

Flexible vectors update

Petr Penzin

Intel

June 20, 2023

Agenda

- ▶ Recap of the proposal
- ▶ Changes to the proposal since the previous update
- ▶ Phase 2?

Brief overview of the proposal

New types and instructions

- ▶ `vec. < type >` – separate vector types for different lane types, size defaults to maximum supported by hardware
 - ▶ `i8`, `i16`, `i32`, `i64` – integer
 - ▶ `f32`, `f64` – floating point
- ▶ `vec. < type > . < op >` – same lane-wise operation as in `simd128 < op >`, applied to vector of `vec. < type > .length`

For example, `vec.f32.mul` is identical to `f32x4.mul` on a 4-lane vector, `vec.i32.add` to `i32x4.add`, and so on

- ▶ `vec. < type > .length` gets number of elements in corresponding vector type
length is fixed during execution and is multiple of 128 bits

Rationale

- ▶ Realizing substantial performance gains from tapping into more advanced SIMD instruction sets
- ▶ Certain algorithms' benefits are limited on 128-bit SIMD
- ▶ Instructions are meant to be direct translated to hardware SIMD, in spirit of existing SIMD standard

Proposal update since previous presentation

New instructions and semantics:

- ▶ Lane index bounds checking
- ▶ Shuffle operations

Lane index bounds checking

Old 'extract lane' and 'replace lane' were taking immediate arguments, which could be out of bounds.

Solution:

- ▶ Change existing 'extract' and 'replace' to take regular integer arguments, but trigger a bounds check at runtime
- ▶ Fallback operations with immediate indices within 128 bits (no bound checking)
- ▶ New 'extract' and 'replace' operations that take their argument modulo the lane count

Shuffle operations

New operations added (already had lane shift):

- ▶ Concat parts of input vectors
- ▶ Select odd lanes from one vector and even - from the other one
- ▶ Reverse lane order
- ▶ Duplicate odd lanes into lower even lanes

Phase 2 entry requirements

- ▶ Precise and complete overview document is available in a forked repo around which a reasonably high level of consensus exists.
- ▶ *Updates to the actual spec document, test suite, and reference interpreter are NOT yet required.*

Current status

- ▶ Added all classes of operations
- ▶ Expected some tweaks based on performance

Poll for Phase 2

Poll?

Thank you