

Flexible vectors use cases

Why go beyond 128-bit SIMD

Petr Penzin

Intel

August 1, 2023

Agenda

- ▶ Performance expectations
- ▶ Targeted hardware
- ▶ Software use cases

Theory

- ▶ Increase in vector size leads to increase in performance
 - ▶ Performing 2x more operations per cycle leads to 2x speedup, 4x to 4x speedup, etc
 - ▶ Realistically that it is an upper bound (Amdahl's law, etc)
- ▶ Actual gain depends on algorithm, but for easily parallelizeable it is going to be substantial
 - ▶ Think element-wide addition of two arrays

Hardware support

Popular hardware extensions:

- ▶ AVX (x86)

supported by virtually all PCs, 128-bit instructions already supported by Wasm runtimes

- ▶ SVE (Arm)

- ▶ AVX512 (x86)

First order focus on AVX and emerging extensions.

Some practical considerations

- ▶ There is a relationship between number of elements in the vector, "complexity" of the algorithm, and resulting performance
- ▶ Relationship between vector size and performance is not linear for some algorithms
- ▶ Faster execution often allows to do "more work"
- ▶ Relatively small gain can be enough to make an algorithm usable

Software

Accelerating existing use cases, GEMM and friends:

- ▶ Machine learning
- ▶ Graphics
- ▶ Physics

Software

Accelerating existing use cases, GEMM and friends:

- ▶ Machine learning
- ▶ Graphics
- ▶ Physics

Expanding use cases:

- ▶ Larger value types (double precision)
- ▶ Text processing / pattern matching
- ▶ Cryptography / off-browser cases

GEMM and basic algebra

Matrix multiplication and basic algebra more broadly cover important existing use cases:

- ▶ Machine Learning ¹
- ▶ Graphics ²
- ▶ Other linear algebra applications, (example: game physics)

Matrix operations are generally susceptible to SIMD parallelization. For 256-bit operations it is safe to assume double digit performance gains, more for wider extensions.

¹Chellapilla et al 2006, Dukhan 2019

²Sobel operator

Higher precision

- ▶ 128-bit SIMD only processes two double precision or 64-bit integer values at a time
- ▶ Only straight-forward kernels would be able to have substantial speedup over scalar
- ▶ 256-bit SIMD instructions process four 64-bit lanes, a lot more to work with

For current 'hard' cases we can expect more than 2x speedup for 256-bit operations.

Expanding into new algorithms

Examples:

- ▶ Pattern matching ³
- ▶ Cryptography, compression
- ▶ Signal processing

³[HyperScan](#)

⁴[RLBox](#), shipping in Firefox since 2020

⁵[WASM as a portable IR for apps](#) (Adnroid NDK)

⁶[MPIWasm](#), [Chadha et al, PPOPP 2023](#)

Expanding into new algorithms

Examples:

- ▶ Pattern matching ³
- ▶ Cryptography, compression
- ▶ Signal processing

Considerations:

- ▶ In general present more steps
- ▶ 128 bit SIMD implementations tend to have low performance gains over scalar
- ▶ Some important non-Web use cases

Examples: sandboxing⁴, portability⁵, virtualization⁶

Need to do a better estimate, rough estimate would be about 2x for 256-bit SIMD.

³[HyperScan](#)

⁴[RLBox](#), shipping in Firefox since 2020

⁵[WASM as a portable IR for apps](#) (Adnroid NDK)

⁶[MPIWasm](#), [Chadha et al, PPOPP 2023](#)

Future work?

See [tracking issue](#).