

01.Nov 2013

Effect Handlers in Low-Level Languages Challenges and Opportunities

Programming Language Team in Edinburgh

www.huawei.com

Using effect handlers from C

■ Why?

- Effect handlers provide: green threads, actors, generators, exceptions
- C: only **modern** language missing **all** of these features
- Therefore: C stands to benefit **the most**!

Using effect handlers from C

■ Why?

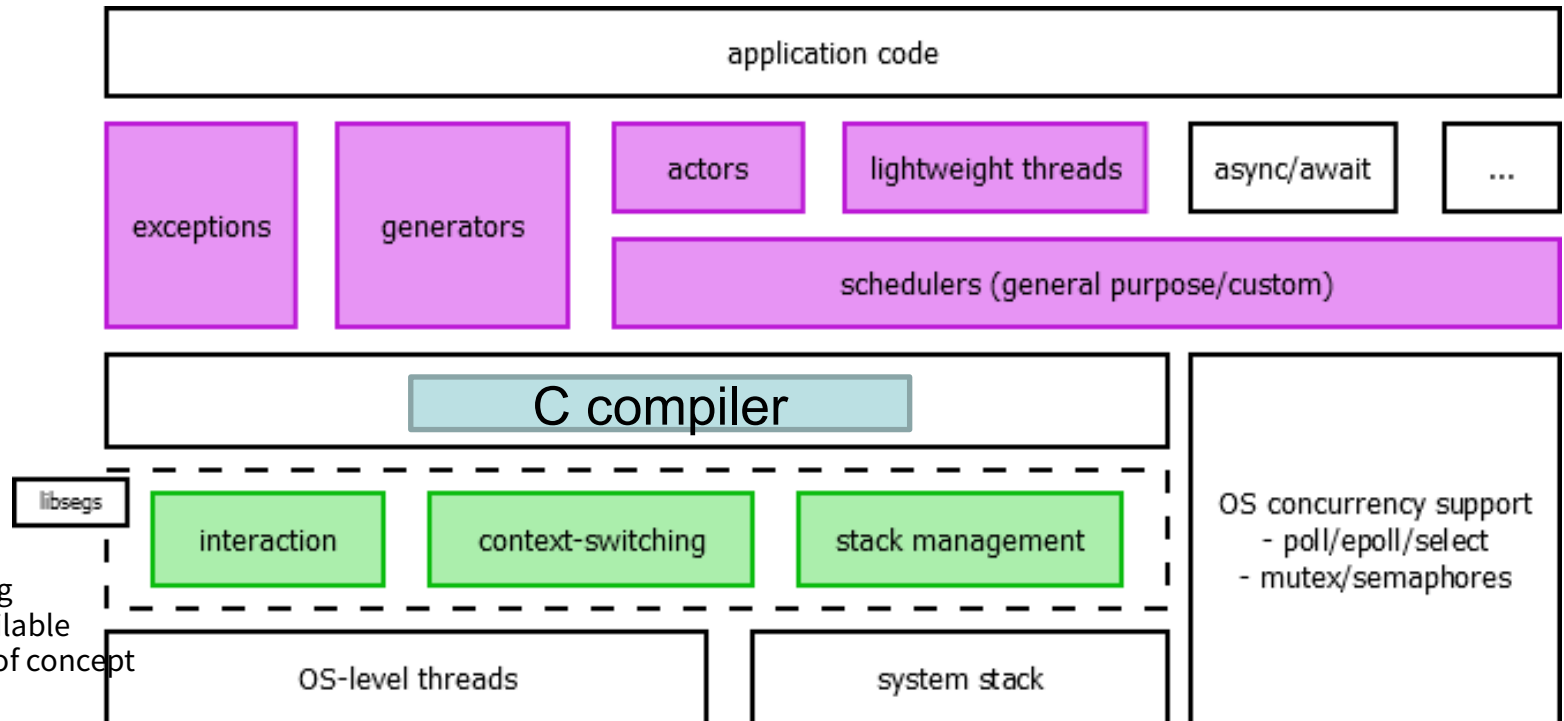
- Effect handlers provide: green threads, actors, generators, exceptions
- C: only **modern** language missing **all** of these features
- Therefore: C stands to benefit **the most**!

■ Ok, but why, really?

- Tons of C code in use at Huawei
- Many projects re-invent concurrency! (Coroutines/actors built on setjmp/longjmp)
- Main goal: use effect handlers to provide **lightweight, modular** concurrency features for C
- Main goal: effect handlers should be **compatible with every C feature** (stack stability)
- Main goal: effect handlers should be **ergonomic** to use **by hand**
- Non-goal: use effect handlers to structure effectful computation
- Non-goal (for now): statically enforce runtime safety

Stackful coroutines in C

- Offer coroutine support through **libseff** library (<https://gitee.com/marioalvarezpicallo/libseff>)
- Prototype implementation in major compilers gcc & clang
- Compiler can provide **extra support, optimizations & better syntax**
- Small asm part needs to be ported to different architectures, **rest is architecture-independent**
- **Effects = stackful coroutines + dynamic binding** (corollary: C programmers are **not scared**)



Effect example – Defining and Performing

```
DEFINE_EFFECT(print, 0, void, { char *str; });

DEFINE_EFFECT(get, 1, int64_t, {});
DEFINE_EFFECT(put, 2, void, { int64_t value; });

void *effectful(int64_t N) {
    for (int64_t i = 0; i < N; i++) {
        int64_t state = PERFORM(get);
        char str[256];
        sprintf(str, "State is %ld", state);
        PERFORM(print, str);
        state = state * state + 1;
        PERFORM(put, state);
    }
    return NULL;
}
```

Effect example – Defining and Performing

```
DEFINE_EFFECT(print, 0, void, { char *str; })  
DEFINE_EFFECT(get, 1, int64_t, {});  
DEFINE_EFFECT(put, 2, void, { int64_t value; });
```

DEFINE_EFFECT macro defines a single operation, a la OCaml

```
void *effectful(int64_t N) {  
    for (int64_t i = 0; i < N; i++) {  
        int64_t state = PERFORM(get);  
        char str[256];  
        sprintf(str, "State is %ld", state);  
        PERFORM(print, str);  
        state = state * state + 1;  
        PERFORM(put, state);  
    }  
    return NULL;  
}
```

Effect example – Defining and Performing

```
DEFINE_EFFECT(print, 0, void, { char *str; });

DEFINE_EFFECT(get, 1, int64_t, {});
DEFINE_EFFECT(put, 2, void, { int64_t value; });

void *effectful(int64_t N) {
    for (int64_t i = 0; i < N; i++) {
        int64_t state = PERFORM(get, {});
        char str[256];
        sprintf(str, "State is %ld", state);
        PERFORM(print, str);
        state = state * state + 1;
        PERFORM(put, state);
    }
    return NULL;
}
```

Return type

Parameter types
(expressed as a struct)

Effect example – Defining and Performing

Programmer must
specify effect ID (0
– 63)

```
DEFINE_EFFECT(print, 0, void, { char *str; });
```

```
DEFINE_EFFECT(get, 1, int64_t, {});
```

```
DEFINE_EFFECT(put, 2, void, { int64_t value; });
```

```
void *effectful(int64_t N) {  
    for (int64_t i = 0; i < N; i++) {  
        int64_t state = PERFORM(get);  
        char str[256];  
        sprintf(str, "State is %ld", state);  
        PERFORM(print, str);  
        state = state * state + 1;  
        PERFORM(put, state);  
    }  
    return NULL;  
}
```


Effect example – Defining and Performing

```
DEFINE_EFFECT(print, 0, void, { char *str; });
```

```
DEFINE_EFFECT(get, 1, int64_t, {});
```

```
DEFINE_EFFECT(put, 2, void, { int64_t value: }):
```

```
void *effectful(int64_t N) {  
    for (int64_t i = 0; i < N; i++) {  
        int64_t state = PERFORM(get);  
        char str[256];  
        sprintf(str, "State is %ld", state);  
        PERFORM(print, str);  
        state = state * state + 1;  
        PERFORM(put, state);  
    }  
    return NULL;  
}
```

Effects are performed via PERFORM macro, this is **relatively** type-safe

Pass pointer to stack-allocated str to effect handler – this is sound because of **stack stability**

Effect example – Handling

```
void *handle_state(closure *clo) {
    seff_coroutine_t *k = seff_coroutine_new(clo->fn, clo->arg);
    effect_set handled = HANDLES(get) | HANDLES(put);
    int64_t state = 0;
    seff_request_t req = seff_handle(k, NULL, handled);
    while (true) {
        switch (req.effect) {
            CASE_EFFECT(req, get, {
                req = seff_handle(k, (void *)state, handled);
                break;
            });
            CASE_EFFECT(req, put, {
                state = payload.value;
                req = seff_handle(k, NULL, handled);
                break;
            });
            CASE_RETURN(req, {
                seff_coroutine_delete(k);
                return payload.result;
            });
        }
    }
}
```

- We often abstract handling to a handle_XYZ function
- Function runs a main loop dispatching on effect
- C lacks closures, we use ad-hoc closure types

Effect example – Handling

```
void *handle_state(closure *clo) {  
    seff_coroutine_t *k = seff_coroutine_new(clo->fn, clo->arg);  
    effect_set handled = HANDLES(get) | HANDLES(put);  
    int64_t state = 0;  
    seff_request_t req = seff_handle(k, NULL, handled);  
    while (true) {  
        switch (req.effect) {  
            CASE_EFFECT(req, get, {  
                req = seff_handle(k, (void *)state, handled);  
                break;  
            });  
            CASE_EFFECT(req, put, {  
                state = payload.value;  
                req = seff_handle(k, NULL, handled);  
                break;  
            });  
            CASE_RETURN(req, {  
                seff_coroutine_delete(k);  
                return payload.result;  
            });  
        }  
    }  
}
```

Effectful functions must be instantiated as **coroutines**: allocate stack space, metadata. Lifetime often managed by handler

Effect example – Handling

```
void *handle_state(closure *clo) {  
    seff_coroutine_t *k = seff_coroutine_new(clo->fn, clo->arg);  
    effect_set handled = HANDLES(get) | HANDLES(put);  
    int64_t state = 0;  
    seff_request_t req = seff_handle(k, NULL, handled);  
    while (true) {  
        switch (req.effect) {  
            CASE_EFFECT(req, get, {  
                req = seff_handle(k, (void *)state, handled);  
                break;  
            });  
            CASE_EFFECT(req, put, {  
                state = payload.value;  
                req = seff_handle(k, NULL, handled);  
                break;  
            });  
            CASE_RETURN(req, {  
                seff_coroutine_delete(k);  
                return payload.result;  
            });  
        }  
    }  
}
```

seff_handle starts/resumes a coroutine, returns a reified request object (can specify return value or performed command)

CASE_EFFECT macro checks effect tag, unpacks request into (typed) payload variable

Coroutine return treated as a special effect with a single result argument

Effect example – Handling

```
void *handle_print(closure *clo) {
    seff_coroutine_t *k = seff_coroutine_new(clo->fn, clo->arg);
    effect_set handled = HANDLES(print);
    seff_request_t req = seff_handle(k, NULL, handled);
    while (true) {
        switch (req.effect) {
            CASE_EFFECT(req, print, {
                puts(payload.str);
                req = seff_handle(k, NULL, handled);
                break;
            });
            CASE_RETURN(req, {
                seff_coroutine_delete(k);
                return payload.result;
            });
        }
    }
}

int main(void) {
    closure closure_1 = (closure){effectful, (void *)10};
    closure closure_2 = (closure){handle_state, (void *)&closure_1};
    handle_print(&closure_2);
}
```

Lack of closures makes composition awkward. Code is equivalent to

```
handle_print() => {
    handle_state() => { effectful(10) }
}
```

Effect example



Coroutine API

- The **coroutine** is the fundamental abstraction of **libseff** (**no resumptions/continuations**)
- The stack frame of any function executing inside the coroutine lives in the coroutine's memory
- Once created, a coroutine may be **resumed**
- Inside a running coroutine, any function may **yield** and provide some information to the context
- Coroutines are **thread-safe** and can be sent between threads to achieve e.g. work-stealing

```
typedef struct { ... } coroutine_t;
```

```
typedef void *seff_start_fun_t(void *);
```

```
seff_coroutine_t *seff_coroutine_new(seff_start_fun_t *, void *);
```

```
void seff_coroutine_delete(seff_coroutine_t *);
```

```
bool seff_coroutine_init(seff_coroutine_t *, fun_t *, void *);
```

```
bool seff_coroutine_release(seff_coroutine_t *);
```

Coroutine API

- The **coroutine** is the fundamental abstraction of **libseff** (**no resumptions/continuations**)
- The stack frame of any function executing inside the coroutine lives in the coroutine's memory
- Once created, a coroutine may be **resumed**
- Inside a running coroutine, any function may **yield** and provide some information to the context
- Coroutines are **thread-safe** and can be sent between threads to achieve e.g. work-stealing

```
typedef struct { ... } coroutine_t;
```

```
typedef void *seff_start_fun_t(void *);
```

```
seff_coroutine_t *seff_coroutine_new(seff_start_fun_t *, void *);
```

```
void seff_coroutine_delete(seff_coroutine_t *);
```

```
bool seff_coroutine_init(seff_coroutine_t *, fun_t *, void *);
```

```
bool seff_coroutine_release(seff_coroutine_t *);
```

- Coroutine & stacklet can be dynamically allocated or programmer can provide memory block
- Implementation is **untyped** (input/return is **void ***)

Coroutine API

- The **coroutine** is the fundamental abstraction of **libseff** (**no resumptions/continuations**)
- The stack frame of any function executing inside the coroutine lives in the coroutine's memory
- Once created, a coroutine may be **resumed**
- Inside a running coroutine, any function may **yield** and provide some information to the context
- Coroutines are **thread-safe** and can be sent between threads to achieve e.g. work-stealing

```
void *seff_yield(coroutine_t *, void *);  
void *seff_resume(coroutine_t *, void *);
```

```
typedef struct {  
    effect_id id;  
    void *payload;  
} seff_request_t;
```

```
void *seff_perform(effect_id, void *);  
seff_request_t seff_handle(coroutine_t *, void *, effect_set);
```

- Handlers are **shallow** (technically **sheep**) – see example later
- Helper macros `DEFINE_EFFECT`, `PERFORM`, `CASE_EFFECT` buy us *some* type-safety/convenience

Coroutine API

- The **coroutine** is the fundamental abstraction of **libseff** (**no resumptions/continuations**)
- The stack frame of any function executing inside the coroutine lives in the coroutine's memory
- Once created, a coroutine may be **resumed**
- Inside a running coroutine, any function may **yield** and provide some information to the context
- Coroutines are **thread-safe** and can be sent between threads to achieve e.g. work-stealing

```
void *seff_yield(coroutine_t *, void *);  
void *seff_resume(coroutine_t *, void *);
```

```
typedef struct {  
    effect_id id;  
    void *payload;  
} seff_request_t;
```

Effects are **reified** as request objects, fit in 2 registers

We use 64-bit bitsets for effects, max 64 definable commands

Coroutine-like API (similar to e.g. libco), can yield to any “parent” coroutine (**not checked**)

Effect-like API: do not yield to specific coroutine, instead search active coroutine stack for installed handler

```
void *seff_perform(effect_id, void *);  
seff_request_t seff_handle(coroutine_t *, void *, effect_set);
```

- Handlers are **shallow** (technically **sheep**) – see example later
- Helper macros `DEFINE_EFFECT`, `PERFORM`, `CASE_EFFECT` buy us *some* type-safety/convenience

Coroutine API

- The **coroutine** is the fundamental abstraction of **libseff** (no resumptions/continuations)
- The stack frame of any function executing inside the coroutine lives in the coroutine's memory
- Once created, a coroutine may be **resumed**
- Inside a running coroutine, any function may **yield** and provide some information to the context
- Coroutines are **thread-safe** and can be sent between threads to achieve e.g. work-stealing

```
void *seff_yield(seff_coroutine_t *, effect_id, void *);  
seff_request_t seff_resume(seff_coroutine_t *, void *);
```

```
typedef struct {  
    effect_id id;  
    void *payload;  
} seff_request_t;
```

Coroutine API actually uses effect request structs, programmer can pass additional data in effect field

```
void *seff_perform(effect_id, void *);  
seff_request_t seff_handle(coroutine_t *, void *, effect_set);
```

- Handlers are **shallow** (technically **sheep**) – see example later
- Helper macros `DEFINE_EFFECT`, `PERFORM`, `CASE_EFFECT` buy us *some* type-safety/convenience

Context Switching

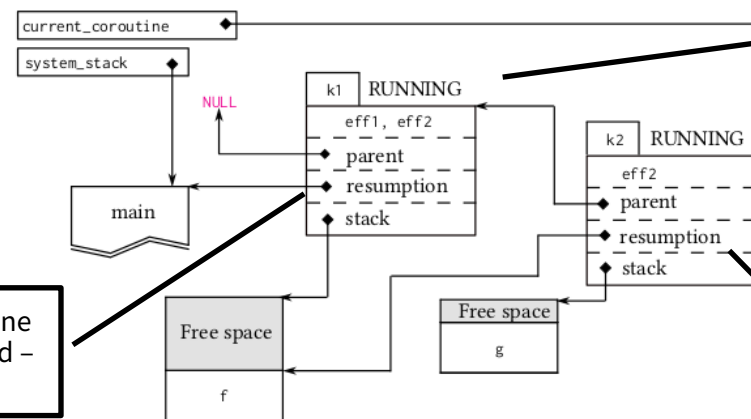
```
1 DEFINE_EFFECT(eff1, void, {});
2 DEFINE_EFFECT(eff2, void, {});
3 void *g(void *arg) { PERFORM(eff1); PERFORM(eff2); }
4 void *f(void *arg) {
5     seff_coroutine_t *k2 = seff_coroutine_new(g, NULL);
6     seff_request_t req1 = seff_handle(k2, NULL, HANDLES(eff2));
7     seff_request_t req2 = seff_handle(k2, NULL, HANDLES(eff2));
8 }
9 void main() {
10    seff_coroutine_t *k1 = seff_coroutine_new(f, NULL);
11    seff_request_t req1 = seff_handle(k1, NULL, HANDLES(eff1) | HANDLES(eff2));
12    seff_request_t req2 = seff_handle(k1, NULL, HANDLES(eff1) | HANDLES(eff2));
13 }
```

A coroutine contains

- A pointer to the **stack frame**
- A pointer to a **parent coroutine**
- A **resumption** (saved register state)
- Bitset of **handled effects**

Thread-local metadata

- Address of **top of system stack**
- Pointer to **currently executing coroutine**



Note that both coroutines are RUNNING since technically k1 never yielded

Resumptions are stored in the coroutine header instead of being passed around – helps avoid allocation

Fig. 1. Before **PERFORM**(eff1)

In a running coroutine the resumption indicates **where to yield**

Context Switching

```
1 DEFINE_EFFECT(eff1, void, {});
2 DEFINE_EFFECT(eff2, void, {});
3 void *g(void *arg) { PERFORM(eff1); PERFORM(eff2); }
4 void *f(void *arg) {
5     seff_coroutine_t *k2 = seff_coroutine_new(g, NULL);
6     seff_request_t req1 = seff_handle(k2, NULL, HANDLES(eff2));
7     seff_request_t req2 = seff_handle(k2, NULL, HANDLES(eff2));
8 }
9 void main() {
10     seff_coroutine_t *k1 = seff_coroutine_new(f, NULL);
11     seff_request_t req1 = seff_handle(k1, NULL, HANDLES(eff1) | HANDLES(eff2));
12     seff_request_t req2 = seff_handle(k1, NULL, HANDLES(eff1) | HANDLES(eff2));
13 }
```

A coroutine contains

- A pointer to the **stack frame**
- A pointer to a **parent coroutine**
- A **resumption** (saved register state)
- Bitset of **handled effects**

Thread-local metadata

- Address of **top of system stack**
- Pointer to **currently executing coroutine**

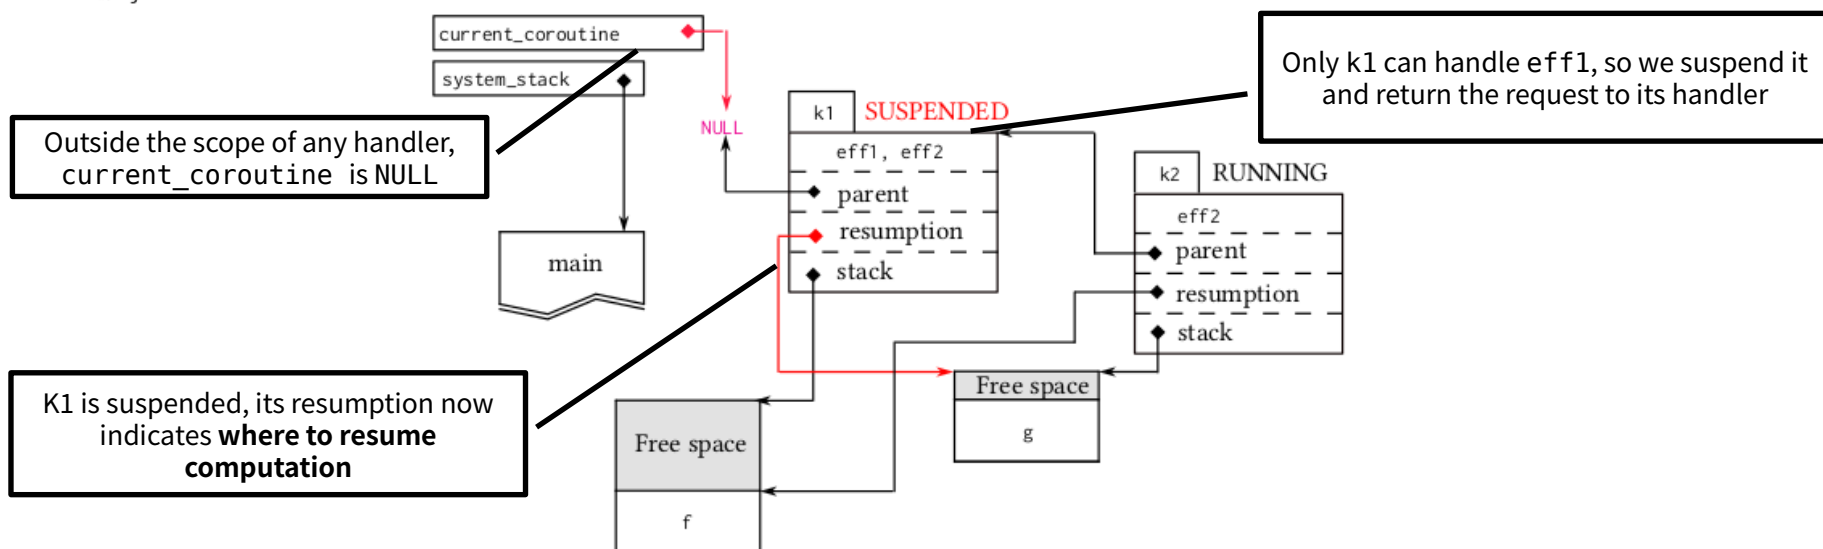


Fig. 2. After `PERFORM(eff1)`

Context Switching

```
1 DEFINE_EFFECT(eff1, 0, void, {});
2 DEFINE_EFFECT(eff2, 1, void, {});
3 void *g(void *arg) { PERFORM(eff1); PERFORM(eff2); }
4 void *f(void *arg) {
5     seff_coroutine_t *k2 = seff_coroutine_new(g, NULL);
6     seff_request_t req1 = seff_handle(k2, NULL, HANDLES(eff2));
7     seff_request_t req2 = seff_handle(k2, NULL, HANDLES(eff2));
8 }
9 void main() {
10     seff_coroutine_t *k1 = seff_coroutine_new(f, NULL);
11     seff_request_t req1 = seff_handle(k1, NULL, HANDLES(eff1) | HANDLES(eff2));
12     seff_request_t req2 = seff_handle(k1, NULL, HANDLES(eff1) | HANDLES(eff2));
13 }
```

A coroutine contains

- A pointer to the **stack frame**
- A pointer to a **parent coroutine**
- A **resumption** (saved register state)
- Bitset of **handled effects**

Thread-local metadata

- Address of **top of system stack**
- Pointer to **currently executing coroutine**

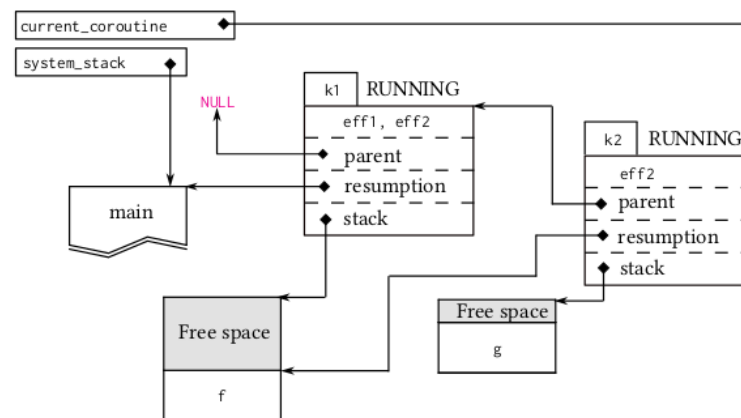


Fig. 3 Before `PERFORM(eff 2)`

Context Switching

```
1 DEFINE_EFFECT(eff1, 0, void, {});  
2 DEFINE_EFFECT(eff2, 1, void, {});  
3 void *g(void *arg) { PERFORM(eff1); PERFORM(eff2); }  
4 void *f(void *arg) {  
5     seff_coroutine_t *k2 = seff_coroutine_new(g, NULL);  
6     seff_request_t req1 = seff_handle(k2, NULL, HANDLES(eff2));  
7     seff_request_t req2 = seff_handle(k2, NULL, HANDLES(eff2));  
8 }  
9 void main() {  
10     seff_coroutine_t *k1 = seff_coroutine_new(f, NULL);  
11     seff_request_t req1 = seff_handle(k1, NULL, HANDLES(eff1) | HANDLES(eff2));  
12     seff_request_t req2 = seff_handle(k1, NULL, HANDLES(eff1) | HANDLES(eff2));  
13 }
```

A coroutine contains

- A pointer to the **stack frame**
- A pointer to a **parent coroutine**
- A **resumption** (saved register state)
- Bitset of **handled effects**

Thread-local metadata

- Address of **top of system stack**
- Pointer to **currently executing coroutine**




Fig. 3 Before `PERFORM(eff 2)`

Context Switching

```
1 DEFINE_EFFECT(eff1, 0, void, {});
2 DEFINE_EFFECT(eff2, 1, void, {});
3 void *g(void *arg) { PERFORM(eff1); PERFORM(eff2); }
4 void *f(void *arg) {
5     seff_coroutine_t *k2 = seff_coroutine_new(g, NULL);
6     seff_request_t req1 = seff_handle(k2, NULL, HANDLES(eff2));
7     seff_request_t req2 = seff_handle(k2, NULL, HANDLES(eff2));
8 }
9 void main() {
10     seff_coroutine_t *k1 = seff_coroutine_new(f, NULL);
11     seff_request_t req1 = seff_handle(k1, NULL, HANDLES(eff1) | HANDLES(eff2));
12     seff_request_t req2 = seff_handle(k1, NULL, HANDLES(eff1) | HANDLES(eff2));
13 }
```

A coroutine contains

- A pointer to the **stack frame**
- A pointer to a **parent coroutine**
- A **resumption** (saved register state)
- Bitset of **handled effects**

Thread-local metadata

- Address of **top of system stack**
- Pointer to **currently executing coroutine**

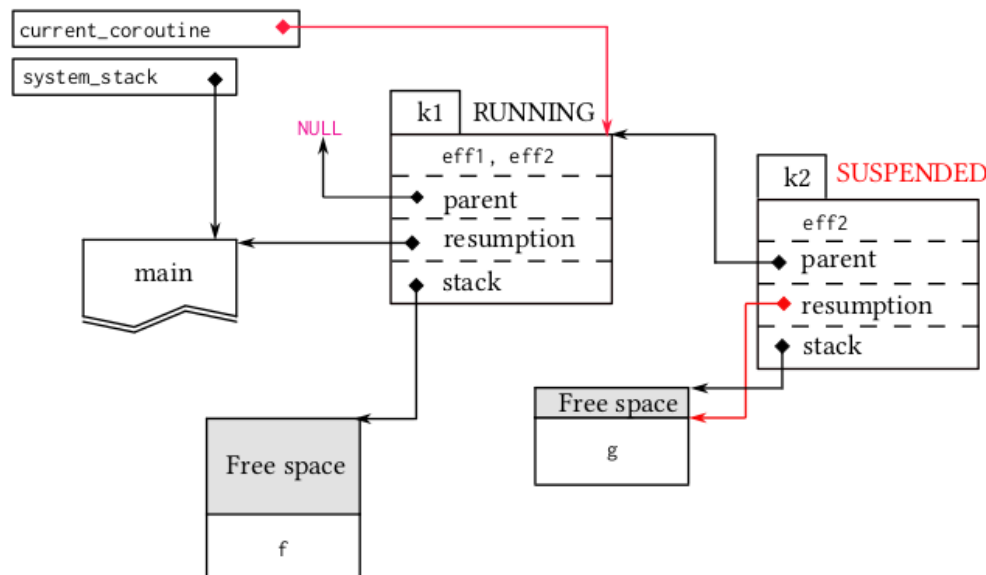


Fig. 3. After `PERFORM(eff2)`

Context Switching

```
1 DEFINE_EFFECT(eff1, 0, void, {});
2 DEFINE_EFFECT(eff2, 1, void, {});
3 void *g(void *arg) { PERFORM(eff1); PERFORM(eff2); }
4 void *f(void *arg) {
5     seff_coroutine_t *k2 = seff_coroutine_new(g, NULL);
6     seff_request_t req1 = seff_handle(k2, NULL, HANDLES(eff2));
7     seff_request_t req2 = seff_handle(k2, NULL, HANDLES(eff2));
8 }
9 void main() {
10     seff_coroutine_t *k1 = seff_coroutine_new(f, NULL);
11     seff_request_t req1 = seff_handle(k1, NULL, HANDLES(eff1) | HANDLES(eff2));
12     seff_request_t req2 = seff_handle(k1, NULL, HANDLES(eff1) | HANDLES(eff2));
13 }
```


A coroutine contains

- A pointer to the **stack frame**
- A pointer to a **parent coroutine**
- A **resumption** (saved register state)
- Bitset of **handled effects**

Thread-local metadata

- Address of **top of system stack**
- Pointer to **currently executing coroutine**

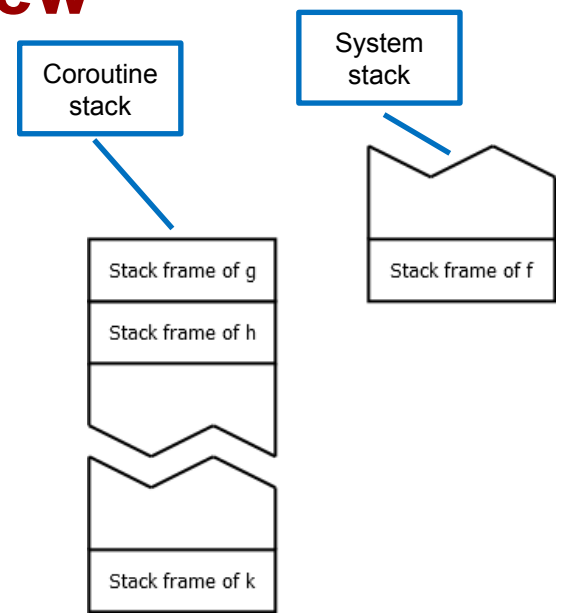
- Single resumption rather than separate yield/resume points
 - Saves space on coroutine metadata at the cost of more time swapping registers – **actually no overhead**
- Effect payload is **allocated on the stack**
 - Perform/resume incur **zero allocations**

PERFORM(put,  put_payload_t payload = (put_payload_t) { 10 });

- Effect tag & payload pointer passed through `seff_perform(put_eff_id, &payload);`

Stackful coroutines: an overview

- Commonly implemented with runtime support
 - Lua (coroutines, **built-in**)
 - Go (goroutines, **built-in**)
 - Java (virtual threads, **built-in since Java 19**)
 - C++ ([Boost::Coroutine](#), **implemented as a library**)
 - Rust ([may](#), **implemented as a library**)
 - Erlang (processes, **built-in**)
- Allocate **entire stack** (not just one frame) for coroutine
 - Stack space can be allocated in heap, global memory, or anywhere
- All calls inside coroutine use coroutine stack
- **Any** function within the coroutine may yield
- Can use **static-sized** stacks or **growable** stacks
 - Growable stacks need more runtime support
- No difference between sync/async functions
 - All functions can call async functions



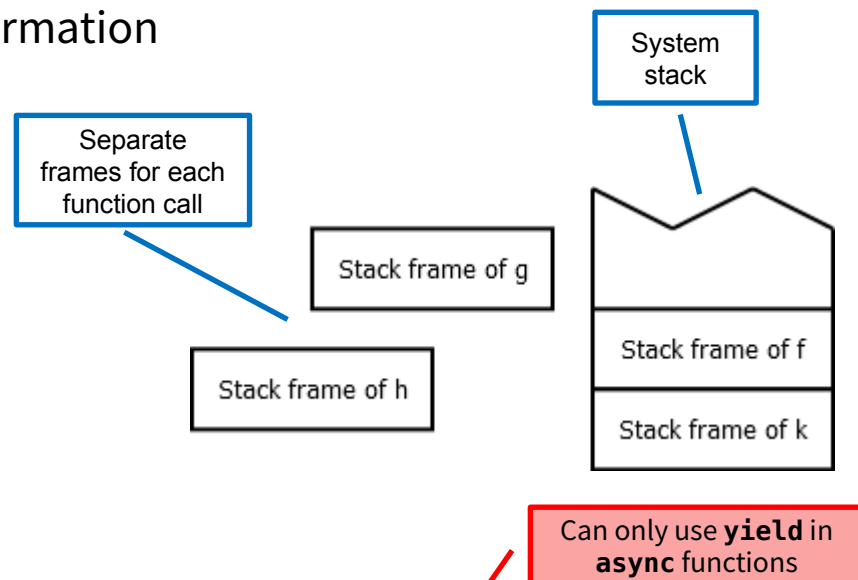
No distinction between sync/async

```
int k() { yield; return 1; }  
int h() { yield; return k(); }  
int g() { h(); }
```

```
void f() {  
    coroutine* coro =  
    create_coroutine(g);  
}
```

Stackless concurrency: an overview

- Commonly implemented via compiler transformation
 - C++ (C++20 coroutines/libcoro)
 - Rust
 - Kotlin
 - Swift
 - Javascript
- Create **single stack frame** for coroutine
 - Frame can be allocated anywhere
 - Function is transformed into state machine
- Calls inside coroutine use system stack
- Can only yield from top-level function
 - Can yield from nested coroutine with special **await** syntax
 - Without complex optimizations, nesting coroutines can be very expensive! (one allocation per coroutine call, chaining yields...)
- Async functions are **special**
 - E.g. cannot be used as function pointers



```
int k() { yield; return 1; }  
int h() async { yield; return k(); }  
int g() async { await h(); }
```

```
void f() {  
    coroutine* coro = g();  
}
```

Hypothetical syntax for stackless coroutines in C


- **yield** for pausing the current coroutine
- **await** for nesting coroutine calls
- **async** for marking coroutine functions

Stackful challenges


- Needs architecture-specific support (portable C library, binds to small platform-dependent asm)
- More complex stack management
 - Resizable stacks
 - Virtual mem } Not suitable for low-level!
 - Stack copying }
 - Segmented stacks — Complex, some runtime overhead
 - Fixed-size stacks — Some memory waste, no recursion
- Cost of context switch — 20~30 μ instructions
- Less efficient use of memory

Stackless:

- Each frame is allocated independently
- More allocations
- But less memory usage!



```
int h() async {  
    yield; return 1;  
}  
int g() async {  
    yield; await  
    h();  
}  
int f() async {  
    await g();  
}
```




Stackful:

- All frames stored in a single memory block
- Some wasted space
- But only 1 allocation!


Stackful challenges

- Needs architecture-specific support (portable C library, binds to small platform-dependent asm)
- More complex stack management
- Cost of context switch
- **Less efficient use of memory**
 - Many optimizations are possible for stackful

```
int h() {  
    yield;  
    return 1;  
}  
int g() {  
    yield;  
    h();  
}  
int f() {  
    g();  
}
```



- When compiler can determine max stack frame size, can allocate exactly what is needed!
- Still some potential for wasted space
- But still only 1 allocation!

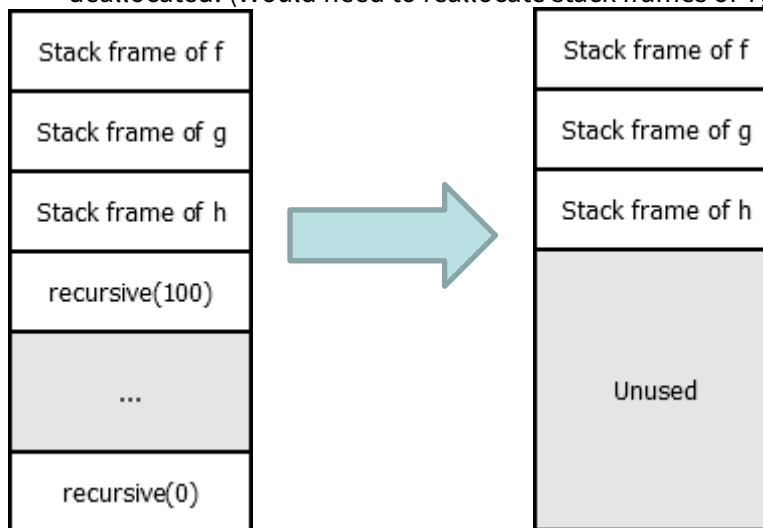


Stackful challenges

- Needs architecture-specific support (portable C library, binds to small platform-dependent asm)
- More complex stack management
- Cost of context switch
- **Less efficient use of memory**
 - Many optimizations are possible for stackful

```
int rec(int n) {  
    ... rec(n-1);  
}  
int h() {  
    yield;  
    recursive(100);  
    yield;  
}  
int g() {  
    yield; h();  
}  
int f() {  
    yield; g();  
}
```

- If compiler can determine max stack frame size, can allocate exactly what is needed!
- **Still some potential for wasted space**
 - After recursive function ends, stack frames are removed but memory cannot be easily deallocated! (Would need to reallocate stack frames of f, g, h)



- Need to allocate stack space for recursive call

- After recursive call, space is not freed, but will not be used!



Stackful challenges

- Needs architecture-specific support (portable C library, binds to small platform-dependent asm)
- More complex stack management
- Cost of context switch
- **Less efficient use of memory**
 - Many optimizations are possible for stackful


```
int recur(int n) {  
    ... recur(n-1);  
}  
int h() {  
    yield;  
    recursive(100);  
    yield;  
}  
int g() {  
    yield; h();  
}  
int f() {  
    yield; g();  
}
```

No **yield**
→ can run
on system
stack

- If compiler can determine max stack frame size, can allocate exactly what is needed!
- **Still some potential for wasted space**
- **Can be mitigated with compiler analysis**



- Run recur on the system stack
- No need for a large coroutine stack!
 - But depends on compiler analysis



- Use segmented stack
- Separate segments for different function calls
 - Deallocate/reuse when finished
 - Some runtime penalty

- Both approaches can be combined
- **Use compiler analysis to decide strategy**
- **If call tree is known at compile time, memory usage can be optimal**

Stack handling

- **libseff** uses **segmented stacks** for stack handling
 - But can easily be adapted to **stack copying** or **virtual memory** if the architecture supports it!
- Coroutines are given an **initial stack** (size can be chosen by the programmer)
- Every function call **checks available stack space vs function stack frame size**
 - If not enough available, **new segment is allocated**
 - The check and allocation are inserted automatically by compiler (clang & gcc -fsplit-stack support)

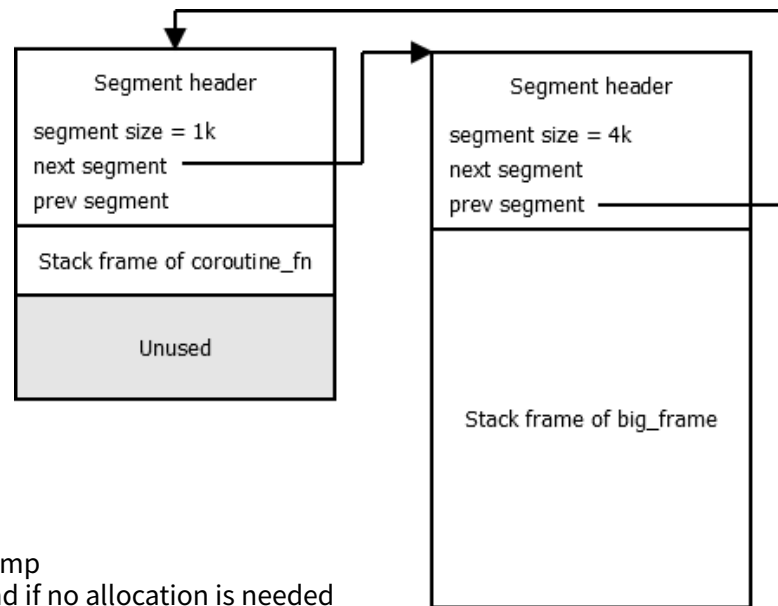
```
int big_frame() {  
    int array[1024];  
    ...  
    return 0;  
}  
int coroutine_fn() {  
    big_frame();  
}
```

- The current stack segment is too small for the big array
- `big_frame` checks stack size in **function prelude**, creates new stack segment in doubly-linked list

```
big_frame():  
    lea    r11, [rsp - 4104]  
    cmp    r11, qword ptr fs:[112]  
    ja     .LBB0_0  
    mov    r10d, 4104  
    mov    r11d, 0  
    call   __morestack  
    ret  
.LBB0_0:  
    ...  
    ret
```

} Prelude: check stack size & allocate

- Prelude is very cheap: load + cmp + jmp
- Branch predictor eliminates overhead if no allocation is needed
- Slow path only taken when stack needs resizing

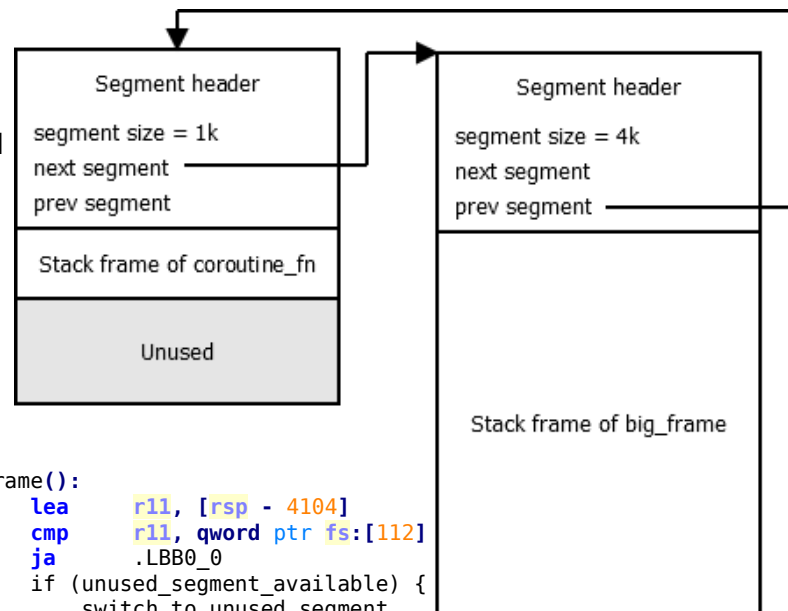


Stack handling

- **libseff** uses **segmented stacks** for stack handling
- Coroutines are given an **initial stack** (size can be chosen by the programmer)
- Every function call **checks available stack space vs function stack frame size**
- Potential performance issue: **hot split problem**

```
int big_frame() {  
    int array[1024];  
    ...  
    return 0;  
}  
int coroutine_fn() {  
    for (int i = 0; i < 1000000; i++) {  
        big_frame();  
    }  
}
```

- `big_frame` allocates new segment, but is deleted at end of function call
- **Allocate and deallocate 1M segments?!**



Can be solved with runtime support!

- Do not deallocate segment upon return, just change pointers
- No need to allocate new segment, just reuse old segment!
- Allocation is replaced by just switching stack ptr
- Change autogenerated function prelude to do check: minimal overhead

Can be solved with compiler analysis!

- Detect big allocation in `big_frame`, lift it to `coroutine_fn`
- Effectively: combine stack frames of `big_frame` and `coroutine_fn`

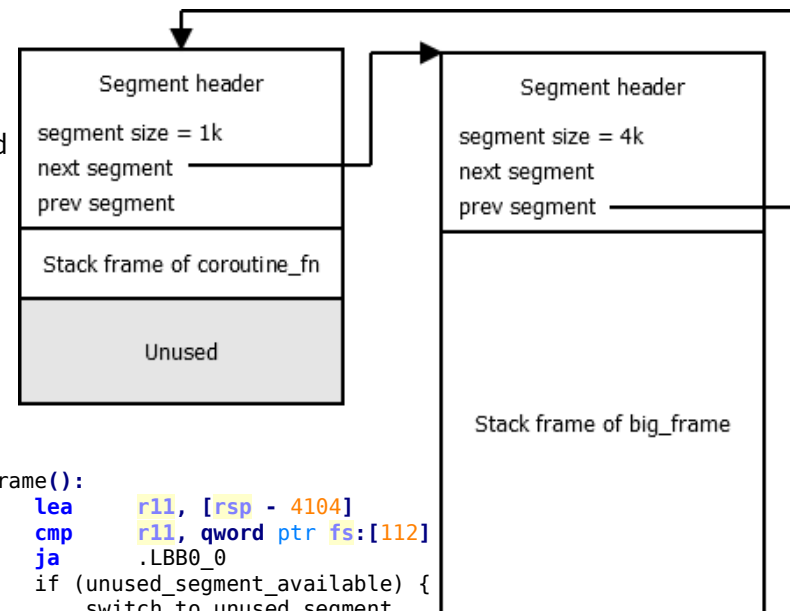
```
big_frame():  
    lea    r11, [rsp - 4104]  
    cmp    r11, qword ptr fs:[112]  
    ja     .LBB0_0  
    if (unused_segment_available) {  
        switch_to_unused_segment  
    }  
    mov    r10d, 4104  
    mov    r11d, 0  
    call   __morestack  
    ret  
.LBB0_0:  
    ...  
    ret
```

Stack handling

- **libseff** uses **segmented stacks** for stack handling
- Coroutines are given an **initial stack** (size can be chosen by the programmer)
- Every function call **checks available stack space vs function stack frame size**
- Potential performance issue: **hot split problem**

```
int big_frame() {  
    int array[1024];  
    ...  
    return 0;  
}  
int coroutine_fn() {  
    for (int i = 0; i < 1000000; i++) {  
        big_frame();  
    }  
}
```

- `big_frame` allocates new segment, but is deleted at end of function call
- **Allocate and deallocate 1M segments?!**



Can be solved with runtime support!

- Do not deallocate segment upon return, just change pointers
- No need to allocate new segment, just reuse old segment!
- Allocation is replaced by just switching stack ptr
- Change autogenerated function prelude to do check: minimal overhead

Can be solved with compiler analysis!

- Detect big allocation in `big_frame`, lift it to `coroutine_fn`
- Effectively: combine stack frames of `big_frame` and `coroutine_fn`

```
big_frame():  
    lea    r11, [rsp - 4104]  
    cmp    r11, qword ptr fs:[112]  
    ja     .LBB0_0  
    if (unused_segment_available) {  
        switch_to_unused_segment  
    }  
    mov    r10d, 4104  
    mov    r11d, 0  
    call   __morestack  
    ret  
.LBB0_0:  
    ...  
    ret
```

Stack handling

- **libseff** uses **segmented stacks** for stack handling
- Coroutines are given an **initial stack** (size can be chosen by the programmer)
- Every function call **checks available stack space vs function stack frame size**
- Potential performance issue: **hot split problem**

```
int8_t *__attribute__((noinline)) bottom() {  
    volatile int8_t arr[BOTTOM_ARR];  
    for (int i = 0; i < MULTS; ++i) {  
        x = x * 3.0;  
    }  
    /* Avoids inlining */  
    __asm__("" : "=o"(v) : "o"(v));  
    return (void *)arr;  
}
```


```
void *__attribute__((noinline)) top(void *_arg) {  
    for (int i = 0; i < REPS; ++i) {  
        int8_t *a = bottom();  
        /* So the result from bottom is actually "used" */  
        __asm__("" : "=r"(a) : "r"(a));  
    }  
}
```

Benchmark a hot-split call

- bottom has enough padding in the stack to require allocating a segment
- Perform variable number of multiplications
- Compare 5 versions
 - Native: no split stack
 - Libseff fixed: fixed-size stacks
 - Libseff baseline: large segments, no hot split
 - Libseff hot split: hot split with segment reuse
 - Libseff dealloc: hot split with eager segment deallocation

Multiplications	0	5	10	15	20
native	1.30	1.00	1.00	1.00	1.00
libseff fixed	1.00	1.00	1.00	1.00	1.00
libseff baseline	1.10	1.00	1.00	1.00	1.00
libseff hot split	9.00	1.83	1.06	1.00	1.00
libseff dealloc	32.22	6.47	3.68	2.50	2.08

(a) Relative execution time of the hot split benchmark



(b)

Fig. 6. Hot Split Results

Conclusions

- Segmented stack prelude does not cause perceptible overhead by itself
- Eager deallocation causes unacceptable performance degradation
- Beyond 8 multiplications, segment recycling completely mitigates hot split
- Below 8 multiplications **bottom should be inlined anyways**
- Stack recycling “happy path” could be faster

Stack handling

- **libseff** uses **segmented stacks** for stack handling
- Coroutines are given an **initial stack** (size can be chosen by the programmer)
- Every function call **checks available stack space vs function stack frame size**
- Potential performance issue: **calling library code**

```
int coroutine_fn() {  
    puts("hello");  
}
```

- printf was compiled without support for segmented stacks
- GNU/Clang segmented stack approach: **conservatively reserve large segment**
- Alternative approach: **switch to system stack**

```
int puts(char*);  
  
int __attribute__((no_split_stack)) puts_syscall_wrapper(char*);  
__asm__(  
    "puts_syscall_wrapper:"  
    "movq %rsp, %fs:_seff_paused_coroutine_stack@TPOFF;"  
    "movq %fs:_seff_system_stack@TPOFF, %rsp;"  
    "movq %fs:0x70, %rax;"  
    "movq %rax, %fs:_seff_paused_coroutine_stack_top@TPOFF;"  
    "movq $0, %fs:0x70;"  
    "callq puts;"  
    "movq %fs:_seff_paused_coroutine_stack@TPOFF, "  
    "%rsp;"  
    "movq %fs:_seff_paused_coroutine_stack_top@TPOFF, %rcx;"  
    "movq %rcx, %fs:0x70;"  
    "retq;"  
)
```

- **libseff** can generate wrapper via macro but:
 - Must be requested manually
 - Programmer must choose which version of the function to call
 - This breaks the promise that stickful concurrency is transparent
- Compiler support can eliminate this need
 - Compiler autogenerates wrappers
 - Or just compile everything with segmented stacks!

Split stack overhead

- **libseff** uses **segmented stacks** for stack handling
- Coroutines are given an **initial stack** (size can be chosen by the programmer)
- Every function call **checks available stack space vs function stack frame size**
- Potential performance issue: **calling library code**

```
int coroutine_fn() {  
    puts("hello");  
}
```

- printf was compiled without support for segmented stacks
- GNU/Clang segmented stack approach: **conservatively reserve large segment**
- Alternative approach: **switch to system stack**

```
int puts(char*);
```

```
int __attribute__((no_split_stack)) puts_syscall_wrapper(char*);  
__asm__(  
    "puts_syscall_wrapper:"  
    "movq %rsp, %fs:_seff_paused_coroutine_stack@TPOFF;"  
    "movq %fs:_seff_system_stack@TPOFF, %rsp;"  
    "movq %fs:0x70, %rax;"  
    "movq %rax, %fs:_seff_paused_coroutine_stack_top@TPOFF;"  
    "movq $0, %fs:0x70;"  
    "callq puts;"  
    "movq %fs:_seff_paused_coroutine_stack@TPOFF, "  
    "%rsp;"  
    "movq %fs:_seff_paused_coroutine_stack_top@TPOFF, %rcx;"  
    "movq %rcx, %fs:0x70;"  
    "retq;"  
)
```

- **libseff** can generate wrapper via macro but:
 - Must be requested manually
 - Programmer must choose which version of the function to call
 - This breaks the promise that stickful concurrency is transparent
- Compiler support can eliminate this need
 - Compiler autogenerates wrappers
 - Or just compile everything with segmented stacks!

Benchmarks: context-switching

- We compare **libseff**, **libco** (Tencent's stackful coroutine library) and C++ coroutines (with **cppcoro**)
 - **libco** is used in **real-world applications** (currently in WeChat backend!)
- All benchmarks running on clang 10.0.0 at optimization level 3
- Context-switching: create a coroutine and resume/yield n times
- We control three different variables: number of yield/resume, depth of the stack, and size of each stack frame

```
void *deep_coroutine(seff_coroutine_t *self, void *arg) {  
    char arr[padding];  
    int64_t depth = (int64_t)arg;  
    if (depth == 0) {  
        volatile bool loop = true;  
        while (loop) {  
            seff_yield(self, nullptr);  
        }  
        return arr;  
    } else {  
        deep_coroutine(self, (void *)(depth - 1));  
        return arr;  
    }  
}
```

Frames padded with
uninitialized data

Infinite loop, volatile to
avoid optimizations

No tail call to avoid
optimizations

libseff version

```
seff_coroutine_t *k1 = seff_coroutine_new(fn, (void *)depth);  
seff_coroutine_t *k2 = seff_coroutine_new(fn, (void *)depth);  
for (size_t i = 0; i < iterations / 2; i++) {  
    seff_resume(k1, nullptr);  
    seff_resume(k2, nullptr);  
}  
seff_coroutine_delete(k1);  
seff_coroutine_delete(k2);
```

Driver code interleaves execution of
2 coroutines iterations/2 times
each

Benchmarks: context-switching

- We compare **libseff**, **libco** (Tencent's stackful coroutine library) and C++ coroutines (with **cppcoro**)
- All benchmarks running on clang 10.0.0 at optimization level 3
- Context-switching: create a coroutine and resume/yield n times
- We control three different variables: number of yield/resume, depth of the stack, and size of each stack frame

```
void *deep_coroutine(void *arg) {  
    char arr[padding];  
    int64_t depth = (int64_t)arg;  
    if (depth == 0) {  
        volatile bool loop = true;  
        while (loop) {  
            co_yield_ct();  
        }  
        return arr;  
    } else {  
        deep_coroutine((void *) (depth - 1));  
        return arr;  
    }  
}
```

Frames padded with
uninitialized data

Infinite loop, volatile to
avoid optimizations

No tail call to avoid
optimizations

Set share_stack on (for resizable
coroutines, otherwise stack size is
fixed)

libco version

API is almost identical to **libseff**

```
stCoRoutine_t *k1;  
stCoRoutine_t *k2;  
stShareStack_t *share_stack  
    = co_alloc_sharestack(1, 1024 * 128);  
stCoRoutineAttr_t attr;  
attr.stack_size = 0;  
attr.share_stack = share_stack;  
co_create(&k1, &attr, fn, (void *)depth);  
co_create(&k2, &attr, fn, (void *)depth);  
for (size_t i = 0; i < iterations / 2; i++) {  
    co_resume(k1);  
    co_resume(k2);  
}  
co_release(k1);  
co_release(k2);
```

Benchmarks: context-switching

- We compare **libseff**, **libco** (Tencent's stackful coroutine library) and C++ coroutines (with **cppcoro**)
- All benchmarks running on clang 10.0.0 at optimization level 3
- Context-switching: create a coroutine and resume/yield n times
- We control three different variables: number of yield/resume, depth of the stack, and size of each stack frame

```
cppcoro::recursive_generator<char*>
deep_coroutine(int64_t depth) {
    char arr[padding];
    if (depth == 0) {
        volatile bool loop = true;
        while (loop) {
            co_yield arr;
        }
    } else {
        co_yield deep_coroutine_rec (depth - 1);
        co_yield arr;
    }
}
```

Frames padded with
uninitialized data

Infinite loop, volatile to
avoid optimizations

No tail call to avoid
optimizations

cppcoro coroutines are heap-
allocated, but RAI1 so there is
no explicit deallocation

cppcoro version

cppcoro api does not allow for an exact comparison. We use recursive_generator here because it is more optimized, but **cppcoro** recursive generators are more limited than coroutines (no async).

```
cppcoro::recursive_generator<char*> k1
= coroutine_fn(depth);
cppcoro::recursive_generator<char*> k2
= coroutine_fn(depth);
cppcoro::recursive_generator<char*>
::iterator k1_iter = k1.begin();
cppcoro::recursive_generator<char*>
::iterator k2_iter = k2.begin();
for (auto i = 0; i < iterations / 2; i++) {
    k1_iter++;
    k2_iter++;
}
```


Benchmarks: context-switching


- We compare **libseff**, **libco** (Tencent's stackful coroutine library) and C++ coroutines (with **cppcoro**)
- All benchmarks running on clang 15.0.0 at optimization level 3
- Context-switching: create a coroutine and resume/yield n times
- We control three different variables: number of yield/resume, depth of the stack, and size of each stack frame

If -O0 we're actually 4.3x
FASTER than cppcoro!

Simple example

- 10,000,000 iterations (resume + yield)
- Call stack has depth 0
- Stack frames have no padding

Framework	Mean time (ms)	Relative
libco	473.4 \pm 33.5	14.00
libseff	60.9 \pm 2.4	1.80
cppcoro	33.8 \pm 3.0	1.00



Conclusions

- **libseff** is much more efficient than **libco**, due to using split stacks instead of stack copying
- **cppcoro** is faster, but less flexible (benchmark code could not be extended with async)

Benchmarks: context-switching

- We compare **libseff**, **libco** (Tencent's stackful coroutine library) and C++ coroutines (with **cppcoro**)
- All benchmarks running on clang 15.0.0 at optimization level 3
- Context-switching: create a coroutine and resume/yield n times
- We control three different variables: number of yield/resume, depth of the stack, and size of each stack frame


Stack size scaling

- 10,000,000 iterations (resume + yield)
- Call stack has depth 0
- Stack frames have 0-5kb of padding

Size = 5000

Framework	Mean time (ms)	Relative
libco	1,428.01	39.63
libseff	60.36	1.67
cppcoro	36.03	1.00

Time (s)



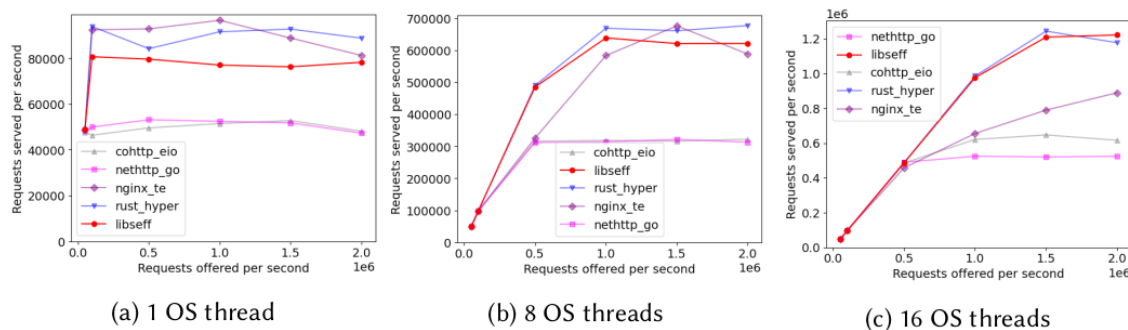
Conclusions

- As expected, **libco** scales linearly with stack size due to stack copying
- Performance of **libseff** and **cppcoro** is independent of stack size

Stack padding (bytes)


Case study

- Goal 1: showcase performance of **libseff** features in “realistic” application
- Goal 2: show how to write applications and schedulers using **libseff**
 1. “Proof of concept” multi-threaded scheduler with async capabilities based on epoll (can easily be adapted to poll/select)
 2. Echo server built on example scheduler, using “listen-accept-fork” approach with coroutines
 3. Benchmark single-threaded performance & multi-threaded scaling
 4. Compare against nginx, hyper (Rust+tokio), OCaml (eio)



■ Toy implementation but performance is competitive

■ Scheduler is extremely naïve, uses simple work-stealing & lock-free datastructures



■ Higher memory consumption than Rust & C due to memory wastage, not significant

Fig. 9. Maximum Memory Consumed (with 1000000 requests offered per second)

Learned lessons

■ Segmented stacks work great, actually!

- “Hot-split” problem not so problematic
- Cost could be reduced further with more ASM

■ Stack allocation + fitting everything into registers

- Biggest advantage over other EH libraries: zero heap allocation. **Needs stack stability!**

■ Group commands into effects

- Solves 64 bit-set problem (though so far not a real problem)
- No need for user to specify effect ID (can use ptrs to per-effect globals)
- Lost flexibility not really an issue

■ Resumption-based API would be possible

- Avoiding heap allocation of resumptions is important, use pointer to coroutine
- Hard to ensure linear usage in presence of race conditions (have to use atomic, cost is prohibitive)
- Some type-safety gains: can type resume properly (C programmers tend not to care)

Conclusions

■ Enormous potential for effects in C

- Can be ergonomic & efficient **without** compiler support!
- But **lots of low-hanging fruit** for compiler support
 - Type-safety, optimizations

■ Major pain point: segmented stacks

- **No real alternative:** virtual memory/stack copying **unworkable**
- Opportunities for optimization
- Gets better with proper effect typing/"purity" tracking!

■ API differences from high-level languages

- No try/handle blocks, continuations not exposed, coroutines as only visible abstraction
- **Session types** obvious candidate for typing coroutines, add extra safety

■ It is worth doing!

- Massive gains in programmer productivity even from a minimal prototype
- Few sharp edges, usable by non-experts!

Thank You

www.huawei.com