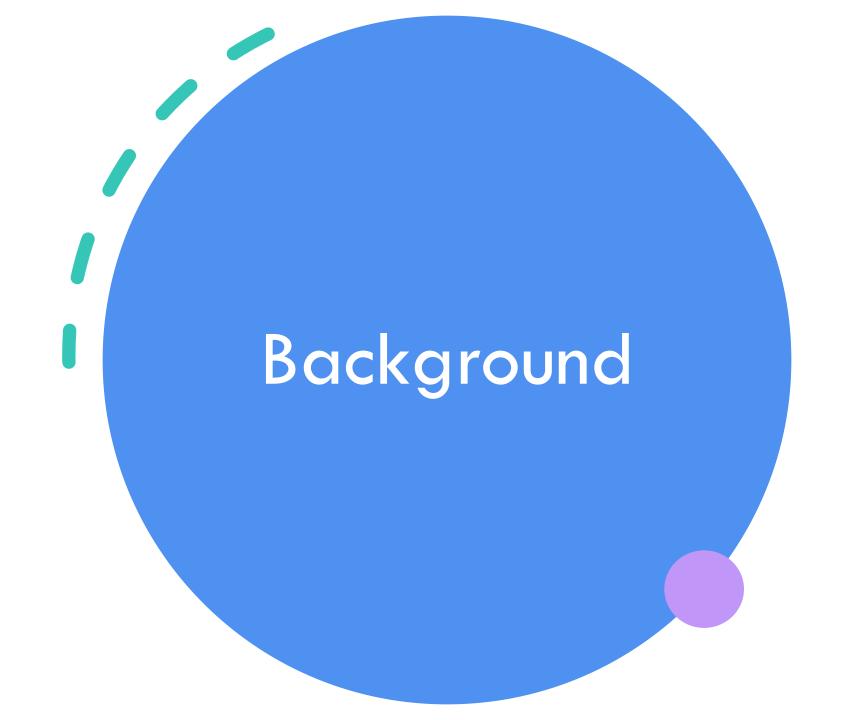


- Some background:
  - Compiling stack switching
  - Passing payloads
- Naive implementation of cont.bind
- Combining cont.bind with more efficient payload passing schemes
- Left-to-right vs. right-to-left binding of cont.bind



#### Use cases (1)

- In general: Giving continuations same type so we can treat them uniformly
- Example: Generator function reading some memory, taking address and length

```
(func $generator (param $addr i32) (param $length i32) ...)
(func $consumer (param i32 $addr)
                                                      ;; $ct2 ~~ cont [i32 i32] -> []
  (local $c (ref $ct0))
                                                      ;; $ct0 ~~ cont [] -> []
                                                      (tag $gen (param i32))
  (local.get $addr) ;; first cont.bind arg
  (i32.const 100) ;; second cont.bind arg
  (cont.new $ct2 (ref.func $generator))
  (cont.bind $ct2 $ct0)
  (local.set $c)
  ($loop
   ;; act on continuation in $c of type $c0:
   ;; either initial value created above or received by handlers on $yield
```

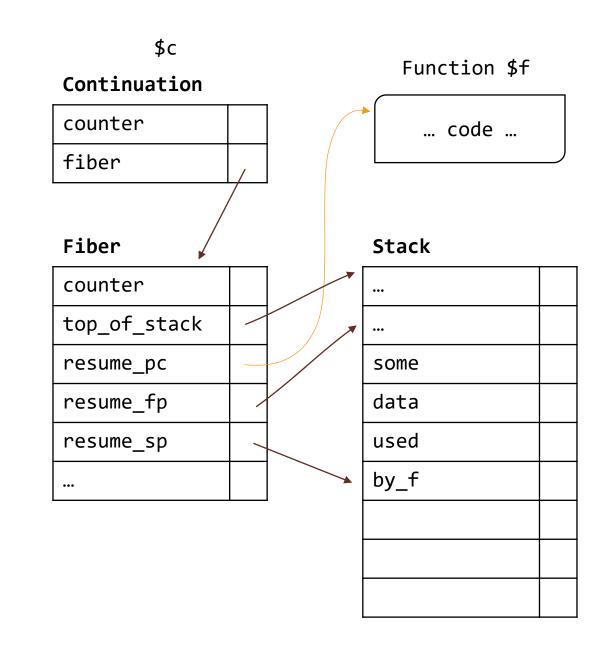
#### Use cases (2)

- In general: Giving continuations same type so we can treat them uniformly
- Example: Task scheduling with two different tag return types
  - (tag \$yield)
     Used by tasks to yield control
  - (tag \$request (result i32))
    Used by tasks to indicate that they need to be assigned next job to work on (identified by i32 id)
  - (table \$task\_queue 0 (ref null \$ct)), with \$ct0 ~~ cont [] -> [] Task queue managed by scheduler
  - Continuations received in \$yield handler:
     Have type \$ct0, immediately re-added to \$task\_queue
  - Continuations received in \$request handler:
     Have type \$ct1 ~~ cont [i32] -> []. Supply work package id using cont.bind \$ct1
     \$ct0, then add resulting continuation to \$task\_queue

## Compiling stack switching

- \$c is currently executing \$f
- What to do on the following? (resume \$ct (local.get \$c))

- Update stack chain
- Save current state (GRPs, SP, FP, IP to resume to)
- Handle resume arguments
- SP := fiber.resume\_sp
- FP := fiber.resume\_fp
- JMP fiber.resume\_pc



## Passing arguments on resume

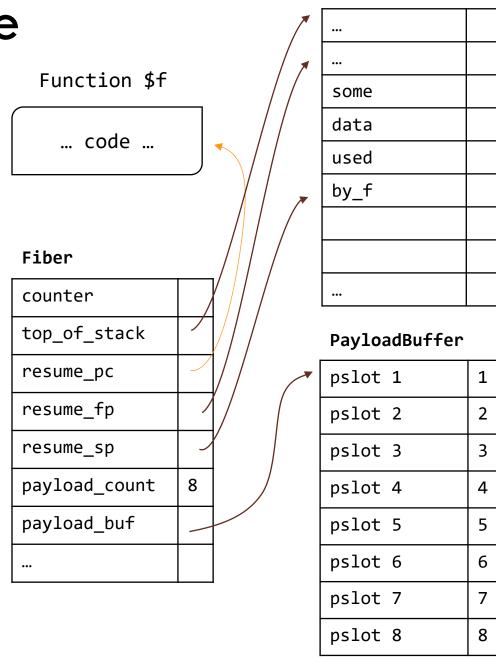
#### Assumption: no cont.bind happened

What to do on the following?

```
;; $ct8~~= cont [i32^8] -> []
(i32.const 1)
(i32.const 2)
...
(i32.const 7)
(i32.const 7)
(i32.const 8)
(resume $ct8 (local.get $c))
```

Option 1: Write to separate buffer, read by other stack after stack switch

 If continuation is "fresh" (at beginning of function \$f), we need extra logic to load arguments into registers as per calling convention



Stack

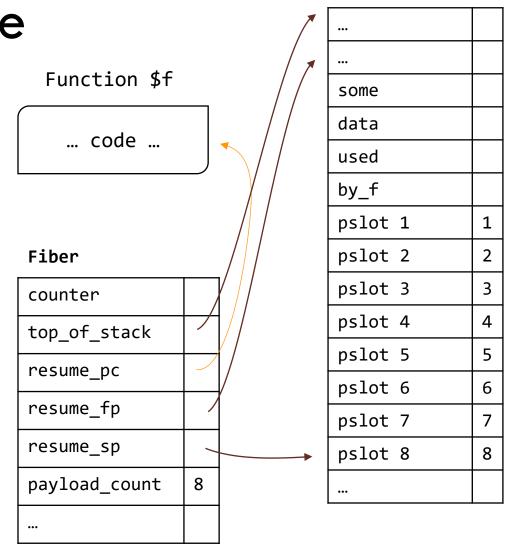
## Passing arguments on resume

Assumption: no cont.bind happened

Situation: resume with i32 args 1 ... 8

Option 2: Same, but use actual stack memory for data

- When c is created (by cont.new, suspend, switch, ...), we know how many slots to reserve)
- Still need to reshuffle if continuation is fresh



Stack

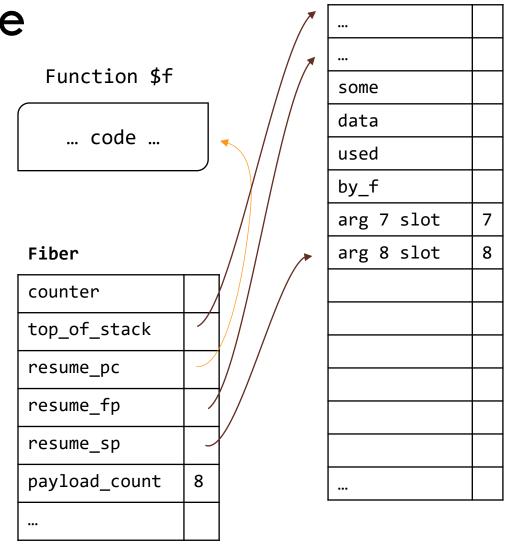
## Passing arguments on resume

Assumption: no cont.bind happened

Situation: resume with i32 args 1 ... 8

Option 3: *Always* use system calling convention

- If passing 8 i32, put first 6 in registers, pass remaining two on stack (assuming System V CC)
- Do this regardless of continuation being fresh or not
- Never need to reshuffle if continuation is fresh, everything is exactly in the right place
- We've ignored cont.bind so far. How to handle that?



Stack

RDI: 1, RSI: 2, RDX: 3, RCX: 4, R8: 5, R9: 6

# cont.bind: Simple implementation

Everything handled through memory

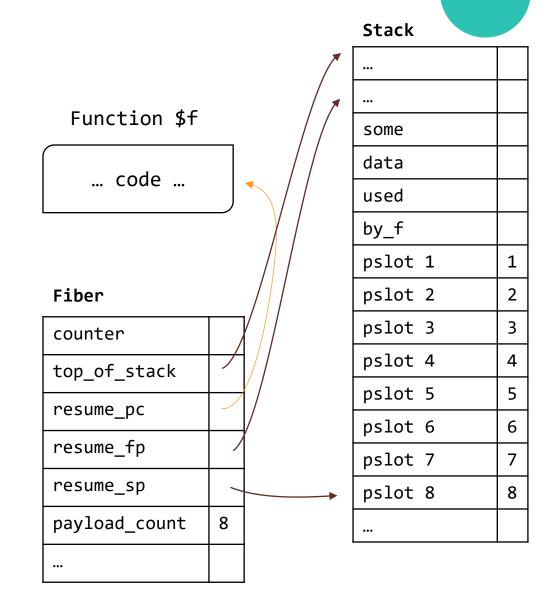
#### Initial observations

 By default, we should assume that every cont.bind argument must be stored somewhere in memory

```
(i32.const 123) ;; cont.bind arg
(cont.new $ct1 (ref.func $f))
(cont.bind $ct1 $ct0)
(local.set $c)

<some 1000 unrelated instructions>
(resume $ct0 (local.get $c))
```

- If all payloads passed using memory ("Option 1 / 2"), we can easily extend that to handle cont.bind
- How to handle more efficient handling scheme? ("Option 3")



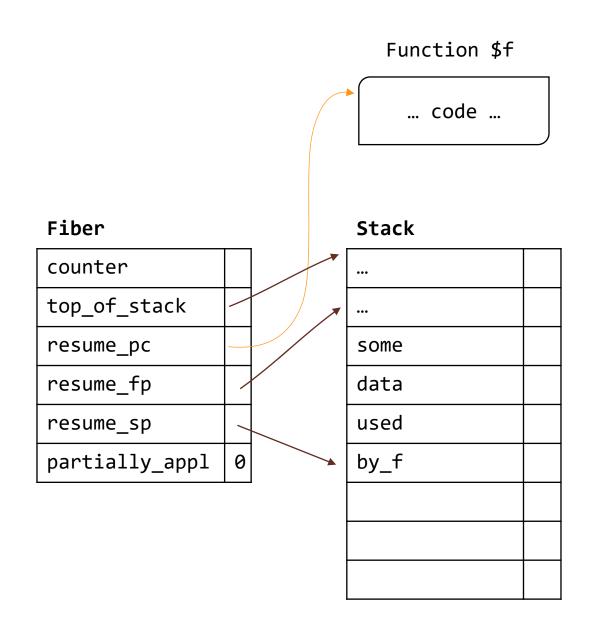
#### The challenge

- Can we achieve the following at the same time?
  - 1. Use some fixed calling convention to pass payloads using registers (& stack, if needed) ("Option 3")
  - 2. Absolutely no overhead on **resume** and **switch** if the continuation was not created by **cont.bind**
- Taking this further: Move all potential overhead into cont.bind code itself. For everything else, must generate and run same code as if cont.bind didn't exist in the proposal
- My claim: Yes! All we need is one boolean flag in each Fiber, never accessed if there was no cont.bind

#### Revisiting resume

We don't want this on resume:

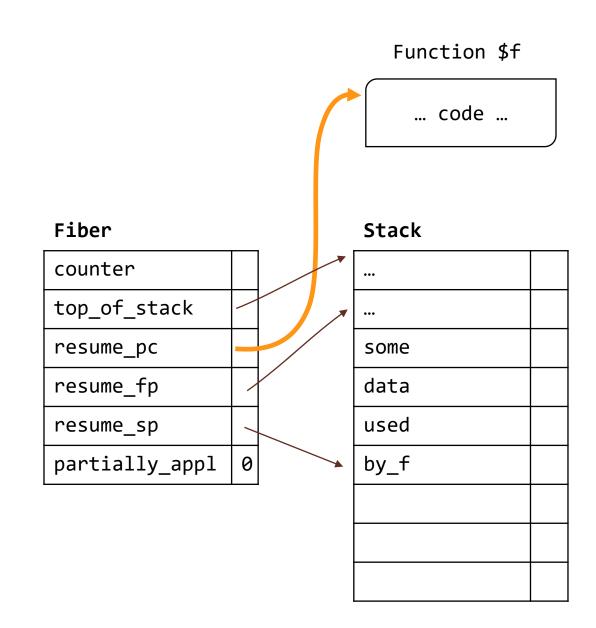
- Update stack chain
- Save current state (GRPs, SP, FP, IP to resume to)
- if fiber.partially\_appl {
   // do something
  } else {
   // do something else
  }
- SP := fiber.resume\_sp
- FP := fiber.resume\_fp
- JMP fiber.resume\_pc



#### Revisiting resume

We don't want this on resume:

- Update stack chain
- Save current state (GRPs, SP, FP, IP to resume to)
- if fiber.partially\_appl {
   // do something
  } else {
   // do something else
  }
- SP := fiber.resume\_sp
- FP := fiber.resume\_fp
- JMP fiber.resume\_pc



# cont.bind: Trampoline-based implementation

#### Idea

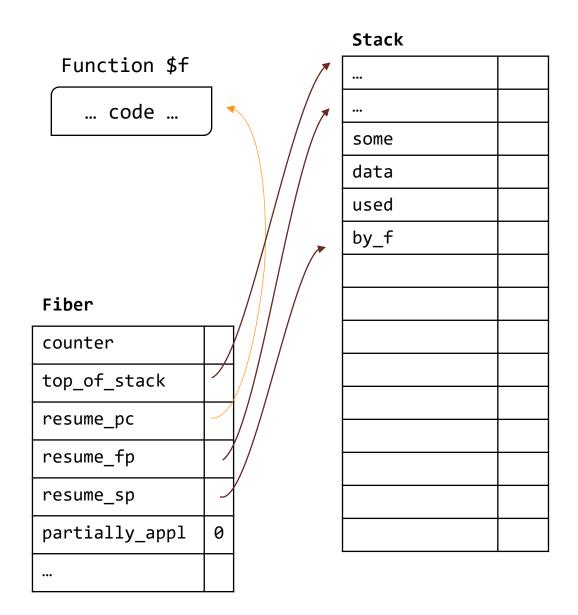
• resume-ing a continuation resulting from *n* cont.bind instructions goes through *n* trampolines before switching to actual continuation

```
;; 5 arguments on Wasm value stack
(local.get $c);; c : $ct5
(cont.bind $ct5 $ct4);; 3rd trampoline to be executed
(cont.bind $ct4 $ct3);; 2nd trampoline to be executed
(cont.bind $ct3 $ct1);; 1st trampoline to be executed
(resume $ct1)
```

- Trampolines accumulate arguments:
  - Receive all previous payloads as parameters
  - Load cont.bind-applied values from memory
  - Pass on combined values to next stage (trampoline or continuation)

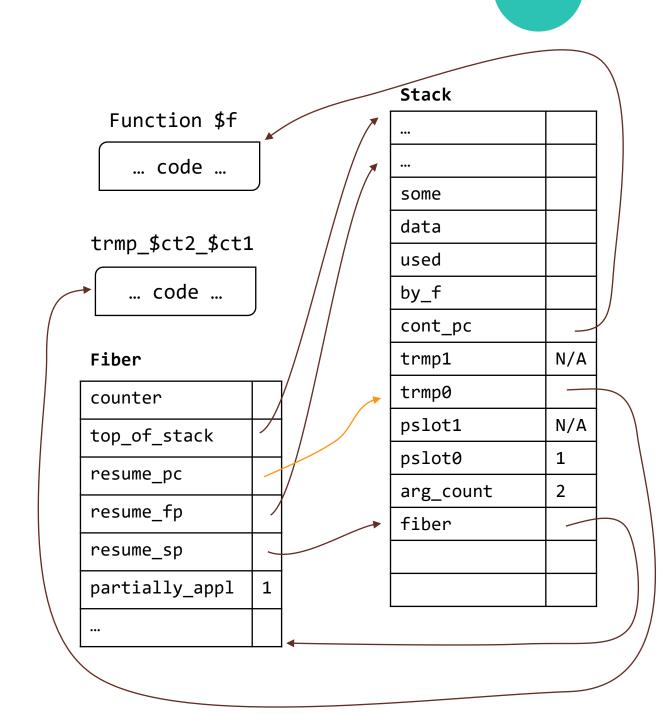
## Example (1)

```
;; $ct2 ~~ cont [i64 i32] -> []
;; $ct1 ~~ cont [ i32] -> []
;; $c2 not created by cont.bind
(cont.bind $ct2 $ct1
  (i64.const 1)
  (local.get $c2))
(local.set $c1)
. . .
(resume $ct1 (i32.const 2) (local.get $c1)
```



# Example (2)

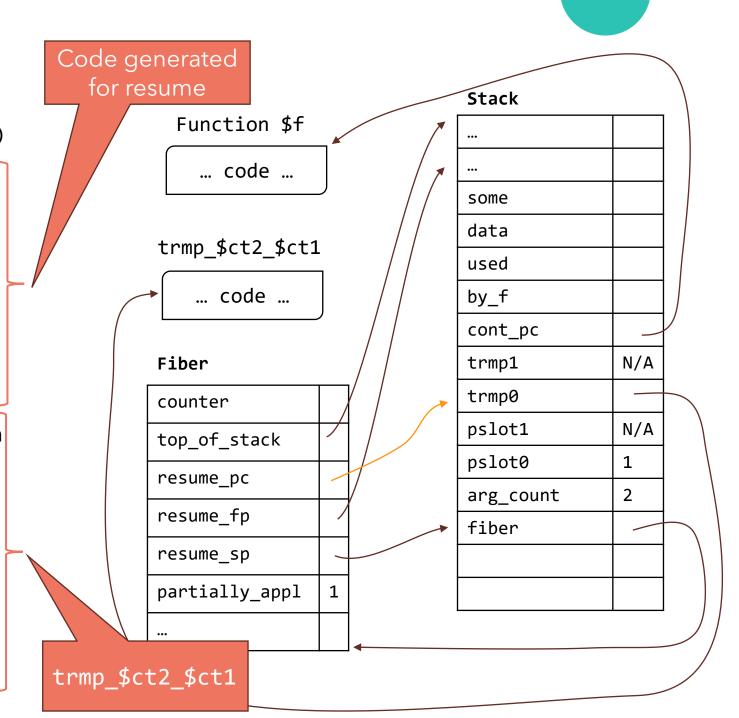
```
;; $ct2 ~~ cont [i64 i32] -> []
;; $ct1 ~~ cont [ i32] -> []
;; $c2 not created by cont.bind
(cont.bind $ct2 $ct1
  (i64.const 1)
  (local.get $c2))
(local.set $c1)
. . .
(resume $ct1 (i32.const 2) (local.get $c1)
```



# Example (3)

(resume \$ct1 (i32.const 2) (local.get \$c1))

- Update stack chain
- Save current state (GRPs, SP, FP, IP to resume to)
- Put resume payload (i32 2) into first argument register (e.g. RDI)
- SP := fiber.resume\_sp
- FP := fiber.resume\_fp
- JMP fiber.resume\_pc
- Receive previous arg (i32 2) as argument in RDI
- load cont.bind argument(i64 1) from pslot0
- Due to no subsequent cont.bind:
  - Remove cont.bind area from stack
  - Set up combined arguments: RDI := 1; RSI := 2
  - JMP cont\_pc



## Trampolines vs cont.bind code

- At cont.bind:
  - Potentially: Do some setup on target stack
  - Write values of types t1\* to memory
- Inside the trampoline (specifically compiled for that combination of types):
  - Receive all previous arguments (types t2\*) as parameters
  - Load cont.bind arguments (types t1\*) from memory
  - Potentially: clean up stack
  - Pass combined sequence (types t1\* t2\*) as arguments on to next stage (continuation or another trampoline) when JMP-ing there

## Handling stack-passed arguments

Not shown before: Trampolines handling arguments that must eventually be passed on in stack slots per CC

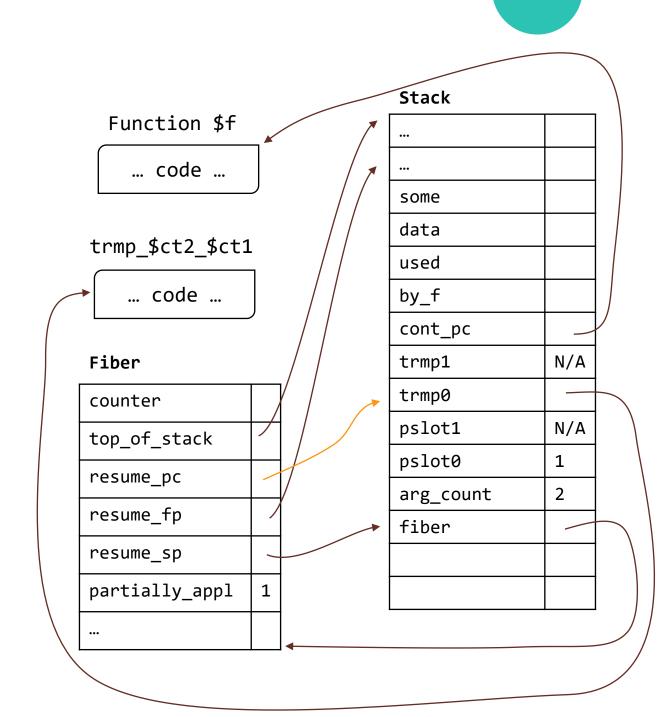
```
Trampoline for (cont.bind $ct_before $ct_after), where
$ct_before = cont [t1* t2*] -> [t3*]
$ct_after = cont [ t2*] -> [t3*]
```

Idea: If t1\* t2\* exceeds what can be passed in registers, start filling stack slots "backwards" starting from end of t1\* t2\*

Just from t1\* t2\*, the trampoline statically knows how many stack slots have already been filled

# What about GC? (1)

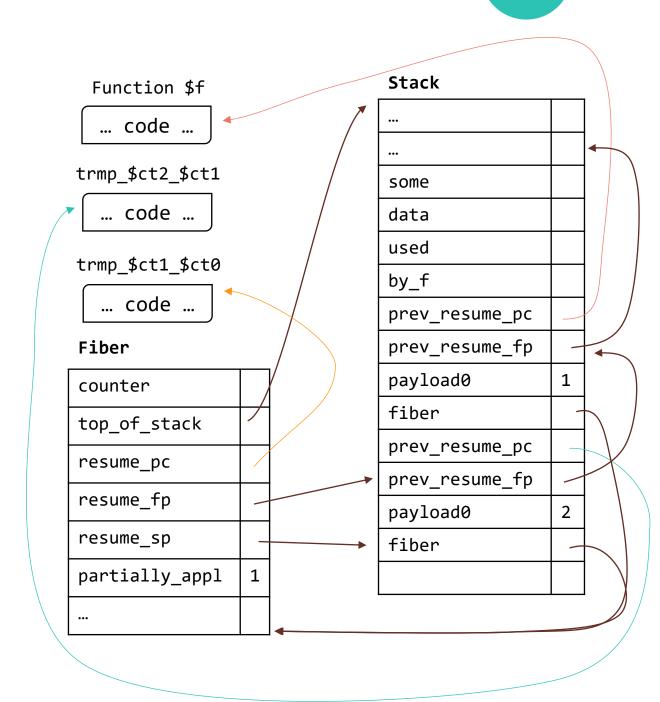
- Observation: When encountering a continuation, GC must scan any cont.bind arguments
- Assumption: Scanning continuations generally works as follows
  - Obtain initial stack map using fiber.resume\_pc, valid stack segment starting at fiber.resume\_sp
  - Then walk continuation's stack and its parents' stacks
- What if fiber.resume\_pc points to a cont.bind trampoline?



# What about GC? (2)

- What if fiber.resume\_pc points to a cont.bind trampoline?
- If cont.bind can only be applied once: Layout similar to what's shown previously should work
- Otherwise: Make each cont.bind add a pseudo stack frame?
  - GC traverses these as always
  - Trampolines are aware of their special nature (e.g., don't RET from their frame)

```
(i32.const 2) ;; second cont.bind
(i32.const 1) ;; first cont.bind
(local.get $c)
(cont.bind $ct2 $ct1)
(cont.bind $ct1 $ct0)
```



## Right-to-left binding (1)

- So far: cont.bind works left-to-right
- What happens if it worked right-to-left?

#### **Left-to-right**

#### ;; \$ct2 ~~ cont [i32 i64] -> [] ;; \$ct1 ~~ cont [ i64] -> [] ;; \$ct0 ~~ cont [ ] -> [] (i64.const 2);; resume arg (i32.const 1) ;; cont.bind arg (local.get c) ;; cont.bind cont (cont.bind \$ct2 \$ct1) (resume \$ct1)

#### **Right-to-left**

```
;; $ct2 ~~ cont [i64 i32] -> []
;; $ct1 ~~ cont [i64 ] -> []
;; $ct0 ~~ cont [ ] -> []
(i64.const 2);; resume arg
(i32.const 1);; cont.bind arg
(local.get c) ;; cont.bind cont
(cont.bind $ct2 $ct1)
(resume $ct1)
```

# Right-to-left binding (2)

- So far: cont.bind works left-to-right
- What happens if it worked right-to-left?

#### **Left-to-right**

#### **Right-to-left**

```
;; $ct2 ~~ cont [i32 i64] -> []
;; $ct1 ~~ cont [ i64] -> []
;; $ct0 ~~ cont [ ] -> []
trmp $ct2 $ct1(prev0 : i64 @ RDI) {
 Load next0 : i32 from memory
 if (is last trampoline) {
                                    Register
   RDI := next0
                                  reshuffle here
   RSI := prev0
   clean up stack, set RSP, RBP
   JMP continuation pc
   else { ... }
```

```
;; $ct2 ~~ cont [i64 i32] -> []
;; $ct1 ~~ cont [i64 ] -> []
;; $ct0 ~~ cont [ ] -> []
trmp $ct2 $ct1(prev0 : i64 @ RDI) {
 Load next0 : i32 from memory
  if (is last trampoline) {
   RDI := prev0
   RSI := next0
   clean up stack, set RSP, RBP
    JMP continuation pc
  } else { ... }
```

# Right-to-left binding (3)

- So far: cont.bind works left-to-right
- What happens if it worked right-to-left?

#### **Left-to-right**

#### **Right-to-left**

#### Filling stack slots:

- Arguments that will eventually be passed in stack slots are passed between trampolines until accumulated arguments exceed what can be passed in registers
- Only then written to corresponding stack slot

#### Filling stack slots:

 Each cont.bind trampoline statically knows whether a given argument will need to go into stack slot end up in a stack slot, puts it there

This only affects situations where we pass many arguments (e.g. > 6 ints on System V) and used **cont.bind** more than once

## Summary (1)

- We want to pass payloads to continuations using same CC as for functions
- Trampoline-based approach can achieve the following:
  - Accumulate payloads previously supplied with cont.bind,
     eventually pass to continuation using chosen CC (in registers & stack slots)
  - Zero overhead on switch/resume if no cont.bind happened!

    All runtime overhead lives in code generated for cont.bind
- Memory cost:
  - In absence of any cont.bind: Potentially a single flag in Fiber
  - Per cont.bind:
    - At runtime: Occupy stack space
    - Code size: Need to compile type-specific trampolines

## Summary (2)

- Two potential approaches towards laying out stack space used by cont.bind
- 1. "One big chunk":
  - Shared by all trampolines
  - Created by first cont.bind instruction
  - Removed by last trampoline before handing over to continuation
  - More difficult to make work with GC if we allow more than one cont.bind to occur
  - Makes it easier to share state between trampolines
- "Pseudo-frames per trampoline"
  - Each cont.bind instruction pushes its own pseudo-frame
  - Each cot.bind trampoline removes its pseudo frame
  - Easier to integrate with GC in presence of multiple cont.bind: The GC just sees a valid sequence of frames with corresponding stack maps.
  - ... haven't thought trough all the details
- Somewhat open question: Which of the two potential memory layouts to use?
  - Make sure that zero overhead introduced into GC

