

Component Model Async Support

WebAssembly CG

May/June, 2022

Outline

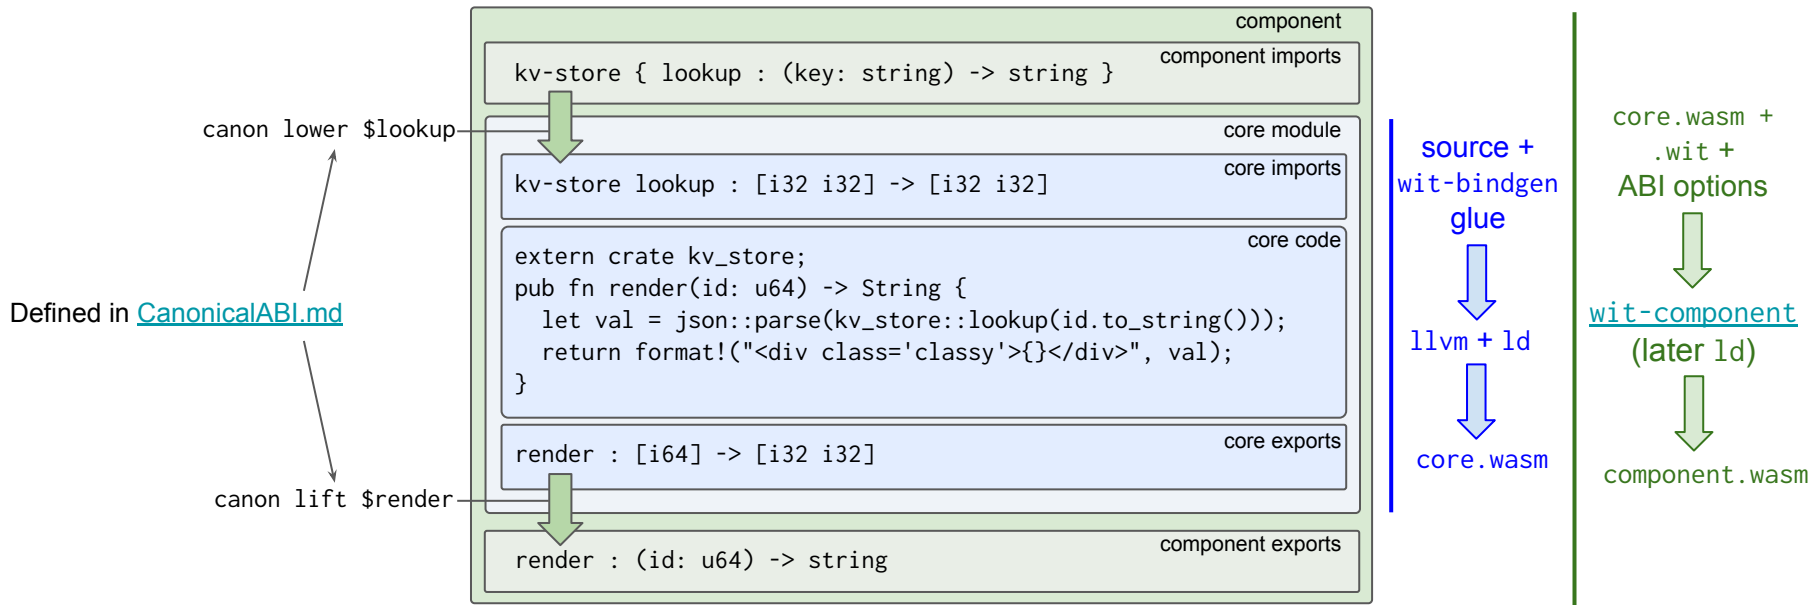
- Motivation
- Background: synchronous Canonical ABI
- Async support
 - future
 - Optimization: callback ABI
 - Optimization: eager return
 - Optimization: stream
 - Optimization: splicing and skipping streams
- Structured concurrency
 - Task tree
 - Task cancellation
 - Task finishing
 - Task scheduling

Caveat: still in flux; feedback welcome

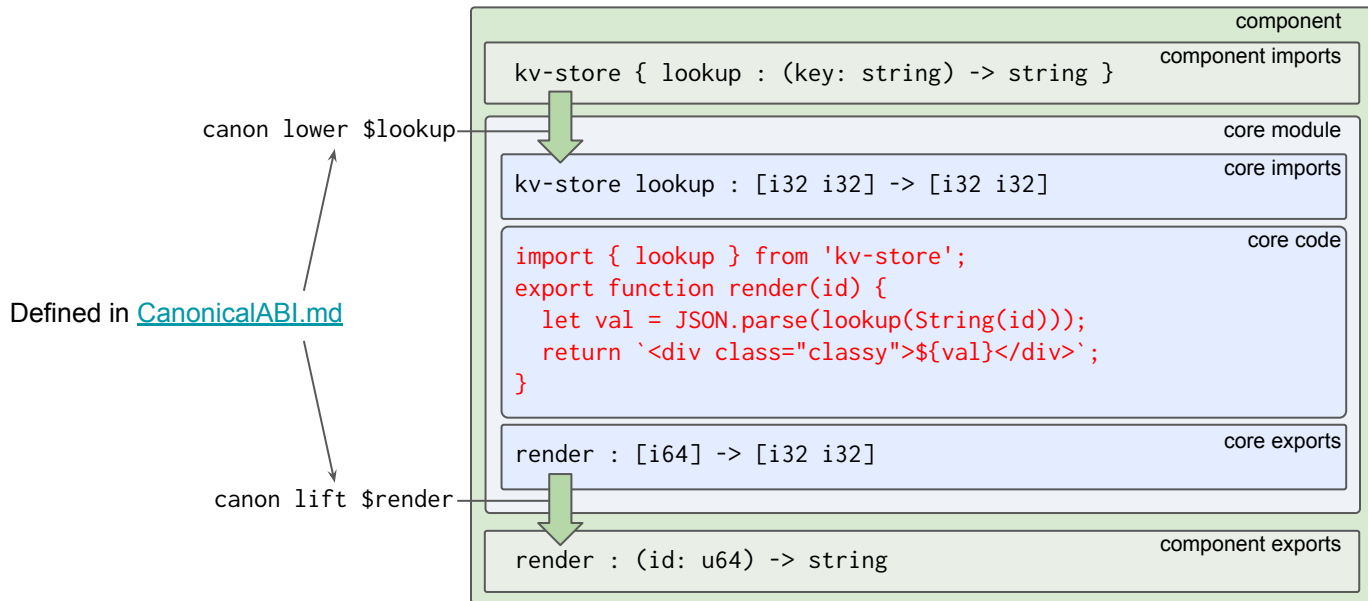
Motivation

- (One slide recap of previously-presented [1,2,3](#))
- How do we specify async/non-blocking operations in WASI and wit?
- Can't we just add first-class functions / callbacks to wit?
 - Cyclic leak problems in non-GC setting (see: Web APIs)
 - Very low-level -- requires manual per-API wrapping to integrate with language concurrency
- Requirements:
 - Virtualizability: async interfaces can be implemented by the host or wasm
 - Efficient I/O implementation when the “other side” is the host (e.g., epoll, io_uring)
 - Ergonomic automatic (wit-bindgen) language bindings
 - Support different styles of language-level concurrency (sync, non-blocking, async, coroutine)
 - Built-in backpressure story (not left as an exercise to the developer)
 - Integrated select / timeout / cancellation across independent interfaces (WASI and host-defined)
 - Ability to keep executing after returning a final value
 - “Just because I want async + modularity doesn't mean I want multi-threading”

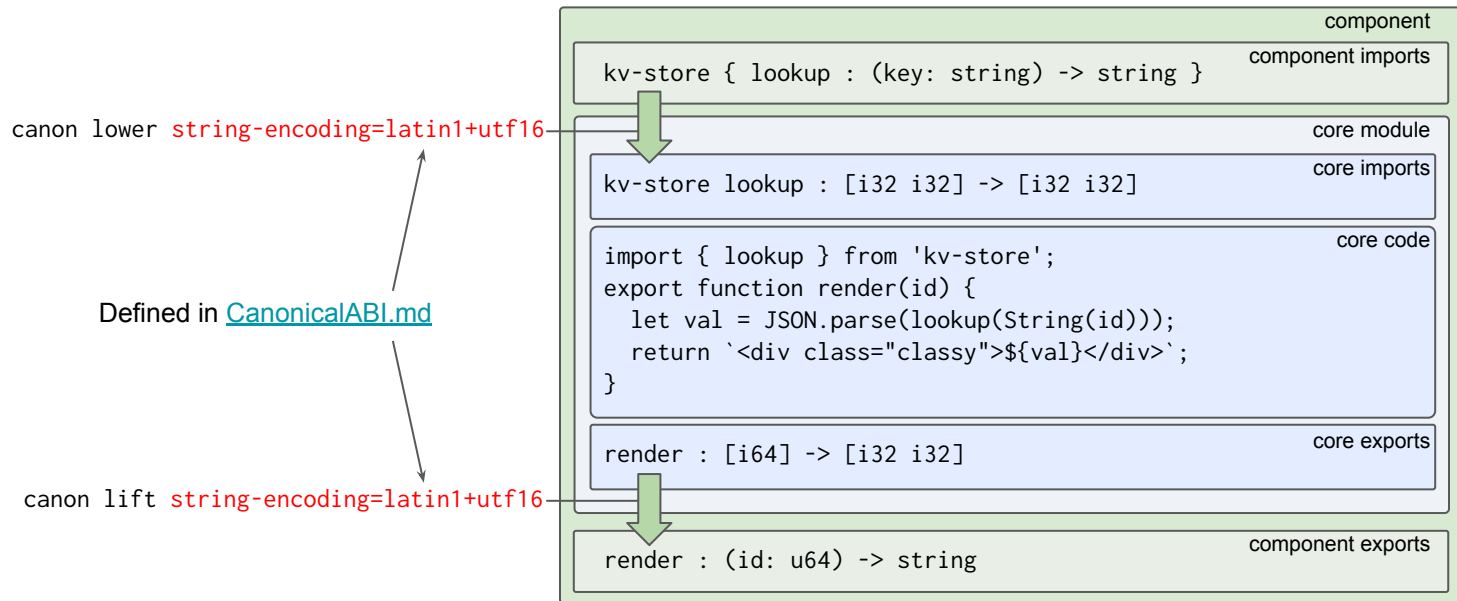
Background: synchronous canonical ABI



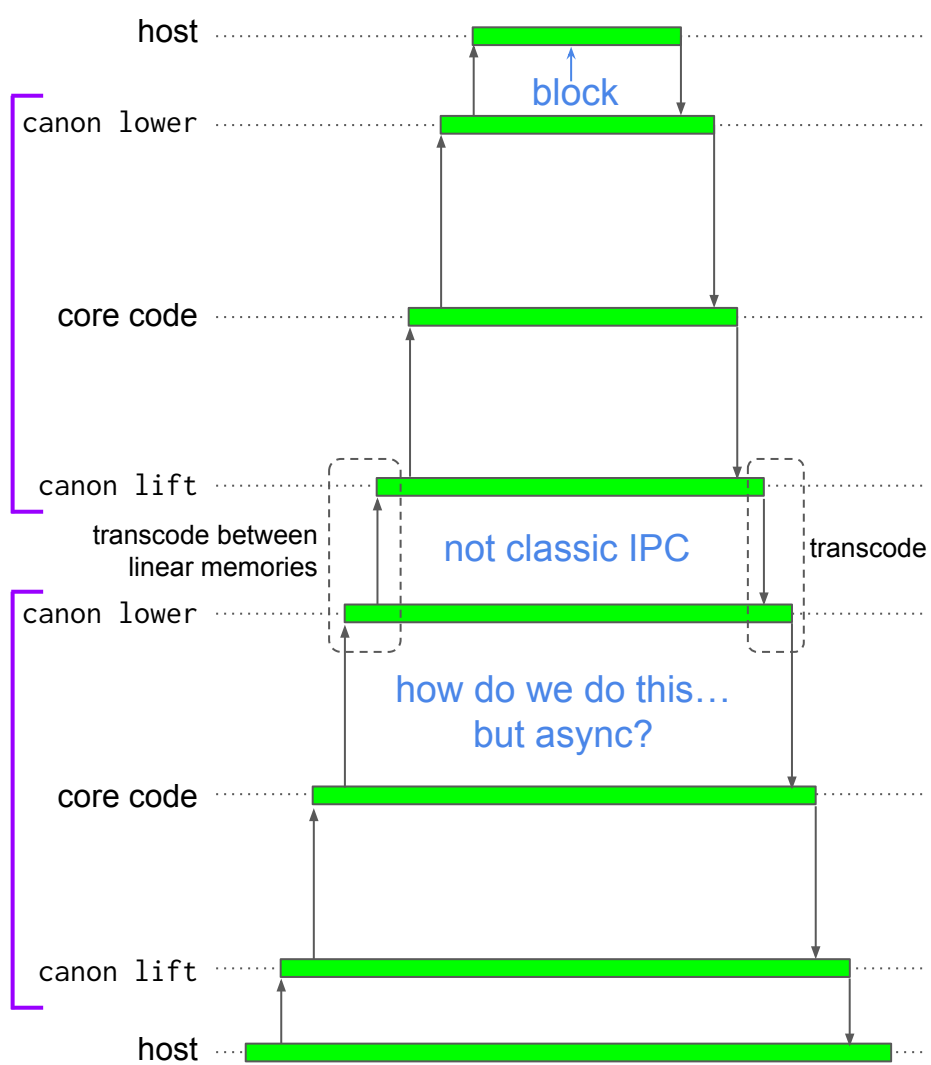
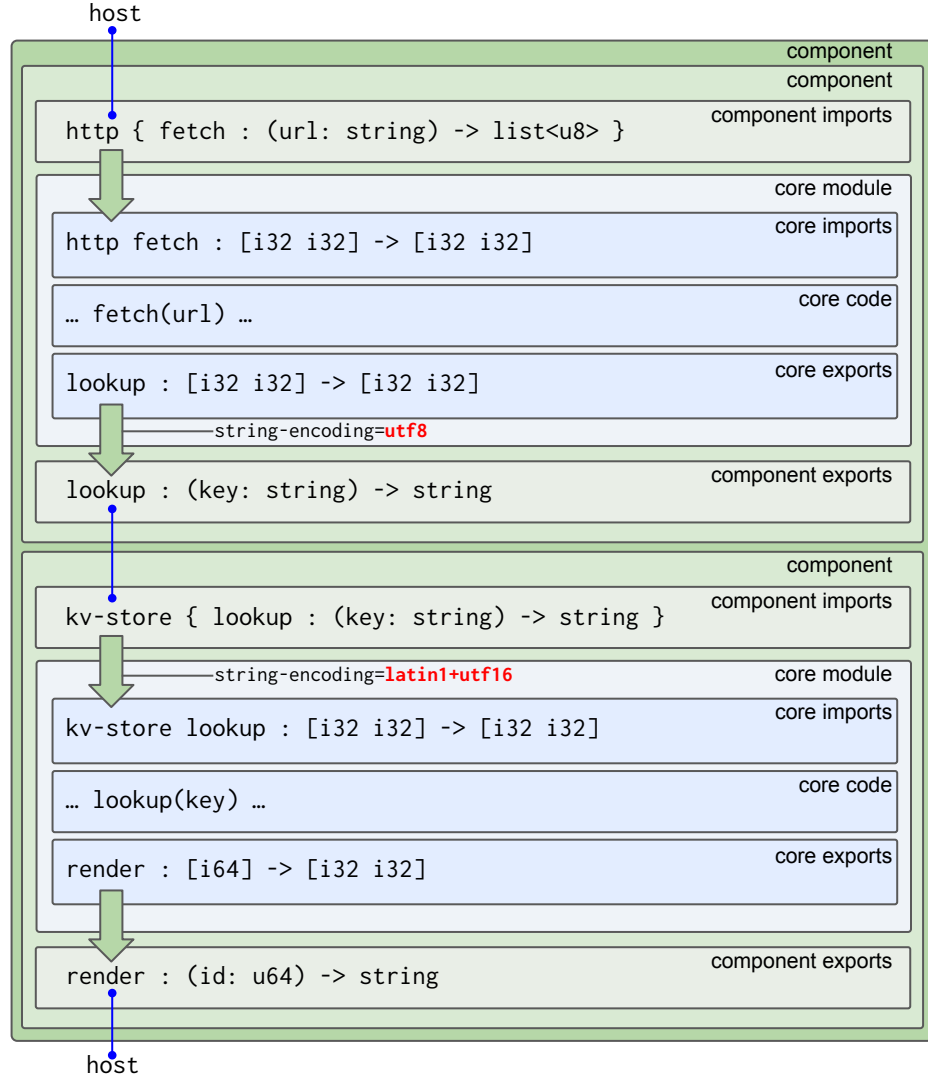
Background: synchronous canonical ABI



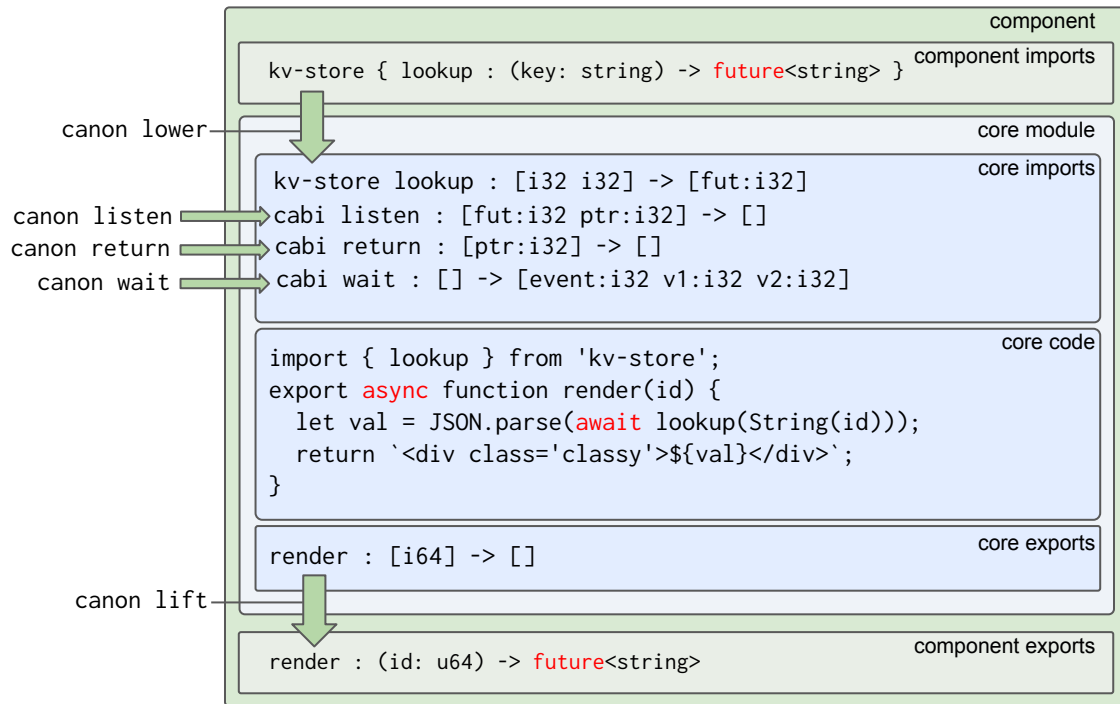
Canonical ABI options



Possible because canon lift and lower bracket all component entry/exit



future



canon lower:

- Returns an index into a component-local future table.
- The future is initially in a “not listening” state.

canon listen:

- Takes a future index returned by lower (et al...)
- Non-blocking: announces *interest* in the future’s value.
- outptr must stay valid until the return event.

canon return:


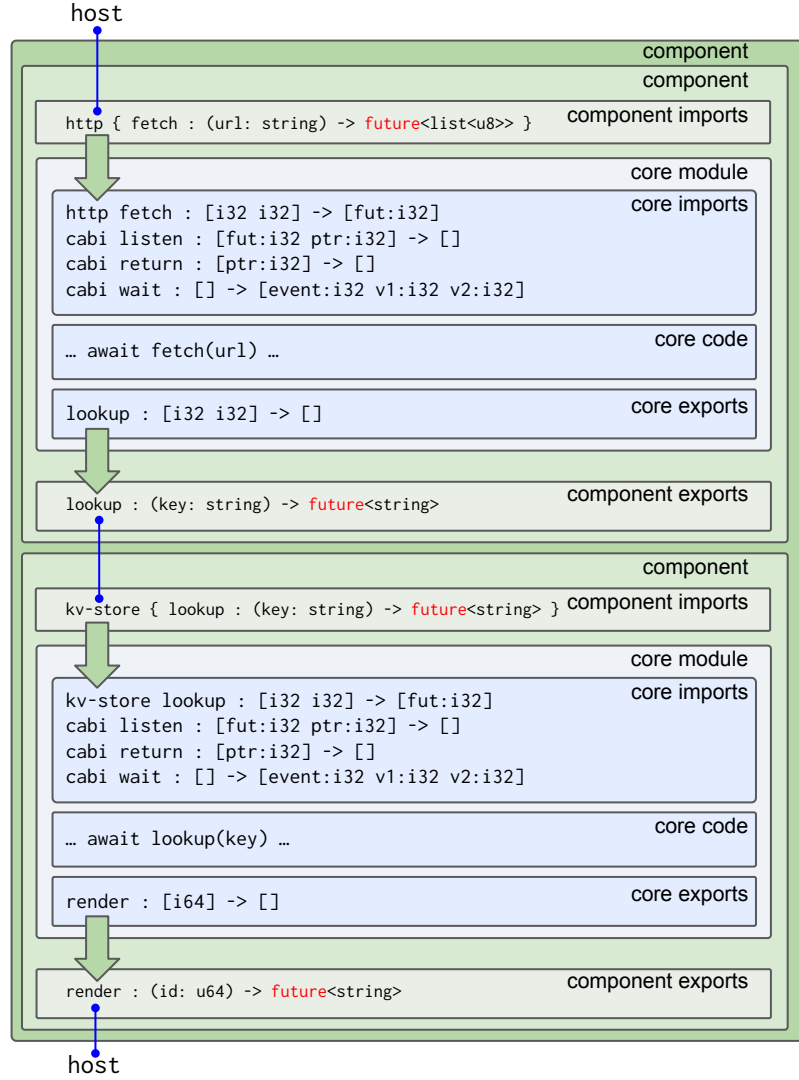
- Non-blocking: *offers* a T return value to the async caller.
- ptr must stay valid until the return-complete event.

canon wait:

- Blocks until *some* event occurs, including:
 - return (v1 is the future index)
 - return-complete

canon lift:

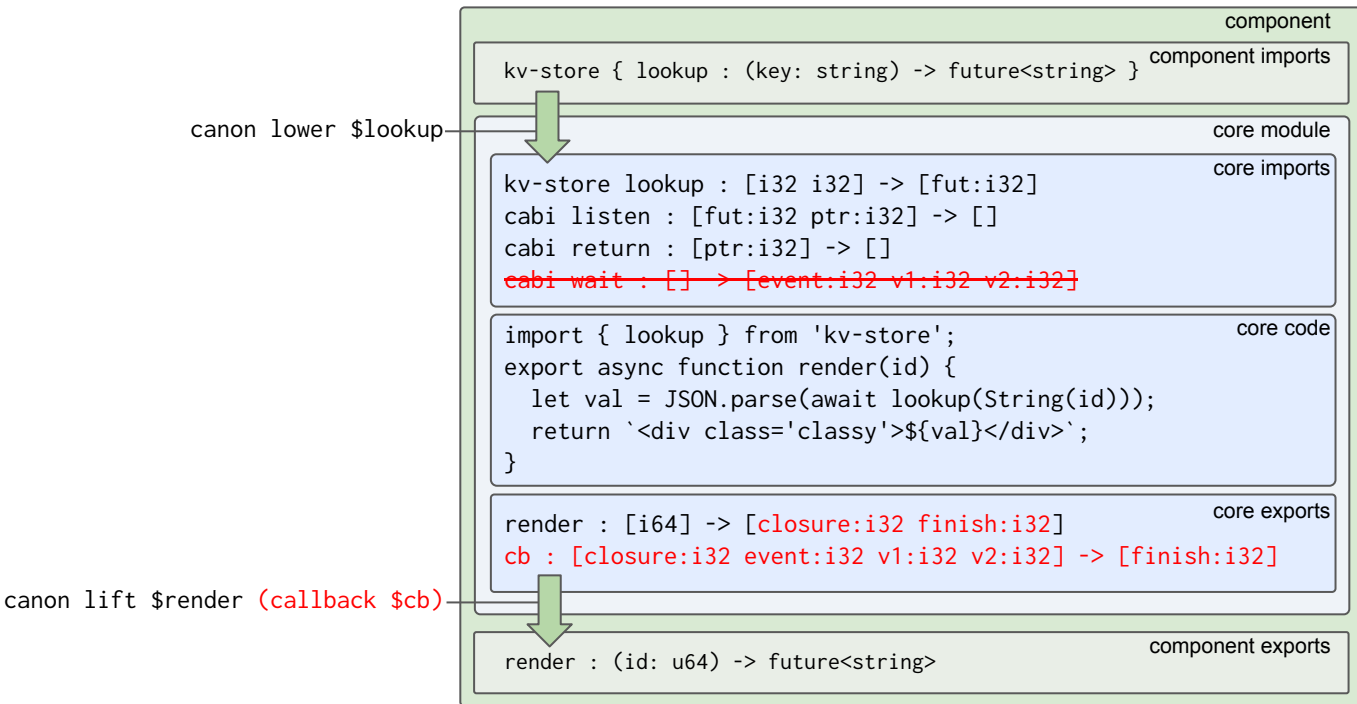
- Traps on return before the return-complete event.



Optimization: callback ABI

- For the future/promise/task/async+await family of languages...
 - viz., .NET, JS, Rust
- ... `wait` will always be performed at the base of the callstack
 - As part of a runtime-owned event loop designed to integrate with OS event system
- In this setting, full stack-switching is overkill
 - The language compiler/runtime already did all the “hard work” of clearing the native stack
- It would be nice to allow producer toolchains to opt out of stack switching
 - Reap the performance benefits paid for by their async model
- Also, some hosts won't support native stack switching for a while (or ever?)
 - Stack switching can be emulated/polyfilled via [asyncify](#), but it's expensive

Optimization: callback ABI



As if:

```
canon-lift (params) {
  let (c, finish) = render(params)
  while (!finish) {
    let (event,v1,v2) = wait()
    finish = cb(c, event, v1, v2)
  }
}
```

Notes:

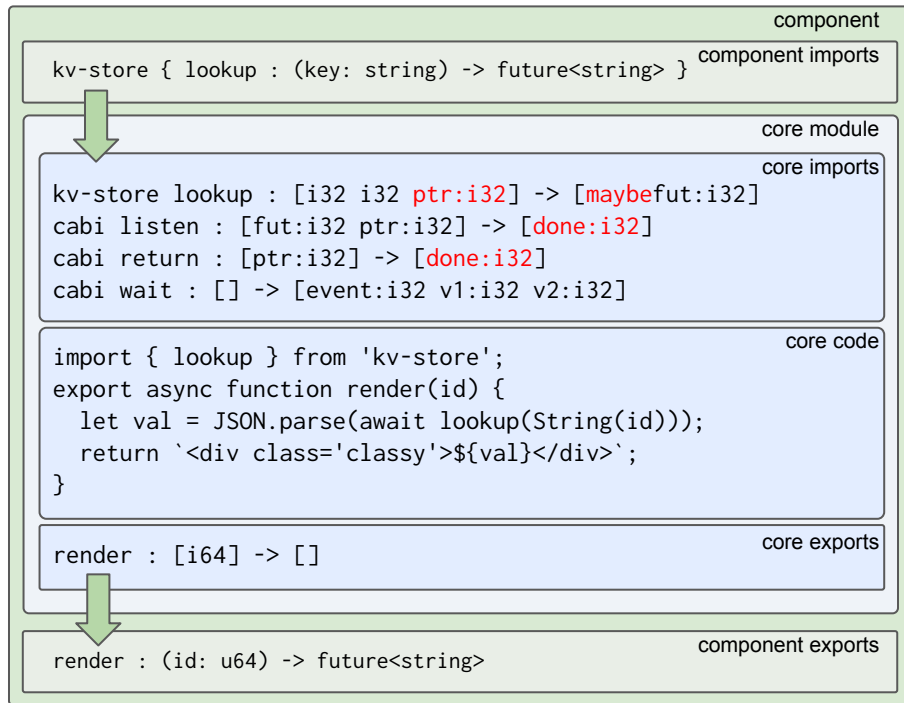
- wait traps under (callback \$cb)
- Encapsulated impl. detail
- Composes with non-callback

The diagram illustrates a hierarchical tree structure, likely representing a data structure or a process flow. The tree is composed of nodes (represented by green rectangles) and edges (represented by arrows). The root node is at the bottom, and the tree branches upwards. The nodes are connected by arrows, indicating the flow of data or the sequence of operations. The tree is divided into two main sections by a vertical dashed line. The left section contains a large tree structure, while the right section contains a smaller tree structure. The root node is labeled 'copy' in a blue oval. The tree structure is composed of nodes (represented by green rectangles) and edges (represented by arrows). The nodes are connected by arrows, indicating the flow of data or the sequence of operations. The tree is divided into two main sections by a vertical dashed line. The left section contains a large tree structure, while the right section contains a smaller tree structure. The root node is labeled 'copy' in a blue oval. The tree structure is composed of nodes (represented by green rectangles) and edges (represented by arrows). The nodes are connected by arrows, indicating the flow of data or the sequence of operations. The tree is divided into two main sections by a vertical dashed line. The left section contains a large tree structure, while the right section contains a smaller tree structure. The root node is labeled 'copy' in a blue oval.


Optimization: eager return

- If the result is already available, future adds overhead
 - Runtime internal allocations, extra listen / wait calls.
- Some languages allow promises/futures to be returned already-resolved
 - Avoiding a trip through the event loop


Optimization: eager return



Before:



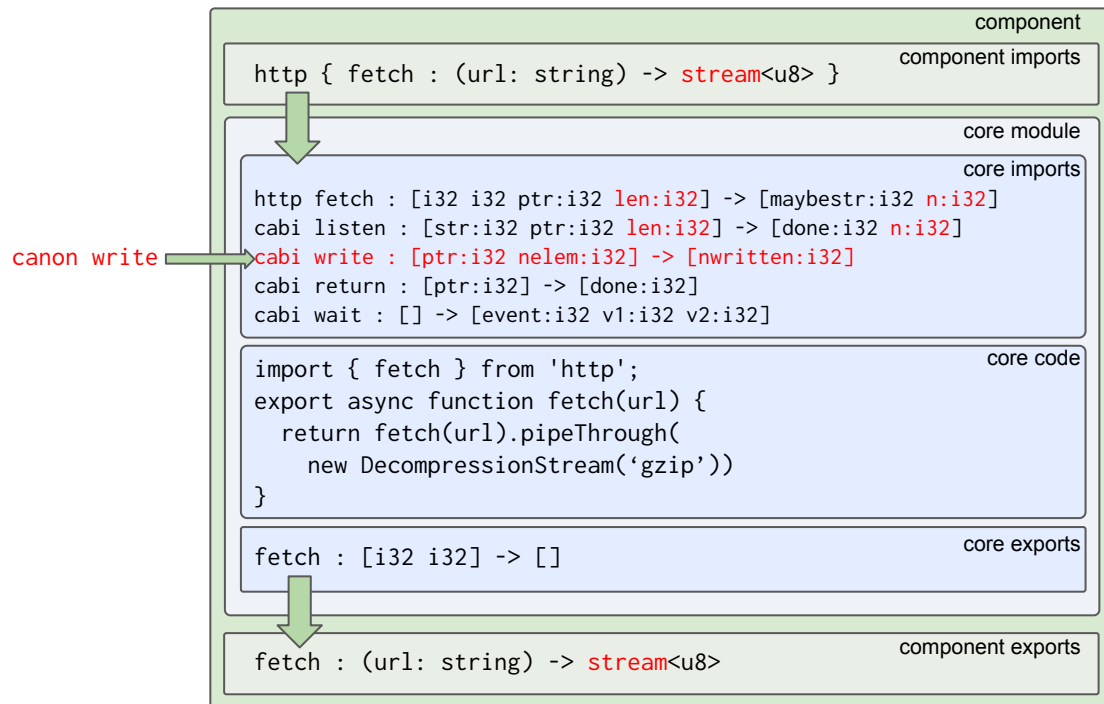
After:



Optimization: stream

- Streams are possible with `stream<T> = future<option<pair<T,stream<T>>>>`
 - (Hand-waving over how we make `stream<T>` recursive...)
- But that's not going to cut it for streams of bytes
 - Need: bulk copies, directly between linear memory (in component-to-component)
- Languages also increasingly have a built-in stream primitives
 - Tightly integrated with the rest of the language (syntax, concurrency model, backpressure, ...)
- So define `stream` as a separate interface type
 - Both as an optimization but also for improved automatic language bindings
- Streams also sometimes have a “closing” value distinct from the elements
 - Effectively: `stream<T,U> = future<either<pair<T,Stream>,U>>`
 - E.g., `main: (argv:list<string>) -> stream<char,s32>`
 - `stream<T> = stream<T,unit>`

Optimization: stream



canon lift (of stream<T,U>):

- Additionally takes the byte-length of ptr
- maybestr=0 means ptr holds $T*U$, $n = |T*|$.
- maybestr≠0 means ptr holds $T*$, $n = |T*|$.

canon listen (of stream<T,U>):

- Additionally takes the byte-length ptr
- ptr receives $T* \text{ xor } U$.
- $n>0$ means “wrote n Ts”, $n=0$ means “returned U”

canon write (of stream<T,U>):

- Offers nelem T values to our async caller
- ptr must stay valid until write-complete

canon return (of stream<T,U>):

- Traps if write in progress.
- Offers a U return value to our async caller
- Closes the stream (no more writes possible)

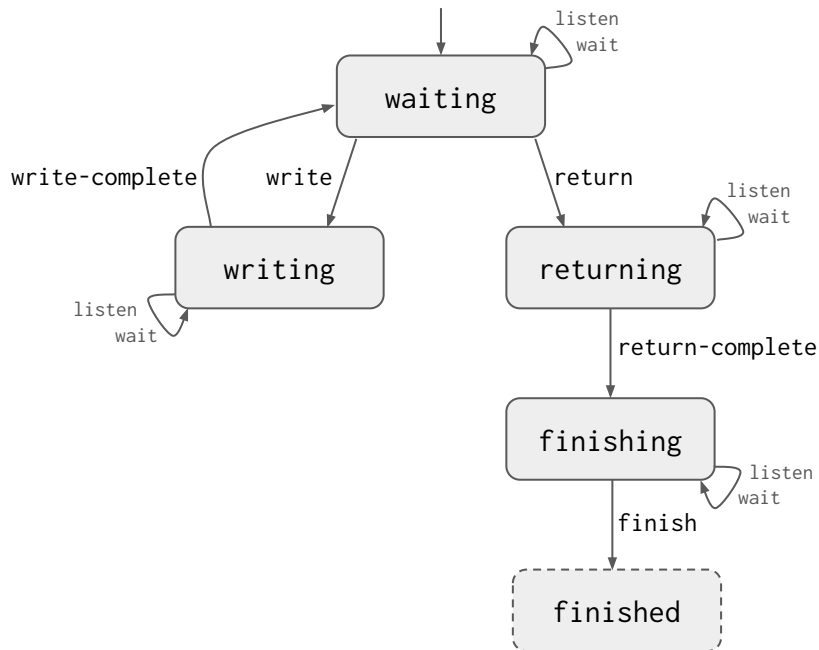
canon wait:

- Additional events:
 - write (v1 is the stream index, v2 is number written)
 - The stream goes back to the “not listening” state
 - Not listening = *backpressure*
 - write-complete (v1 is n written)

Optimization: splicing and skipping streams

- It's very common to copy big chunks from one stream to another
 - Don't want to have to read into linear memory just to immediately write back out.
- `canon splice : [str:i32 nelem:i32] -> [done:i32]`
 - Acts like `listen+write`: if `done=0`, must wait for a write-complete event.
- Sometimes we want to slice out just a subset of a stream
 - Don't want to copy bytes into linear memory just to advance the read offset
- `canon skip : [str:i32 nelem:i32] -> [done:i32]`
 - Acts like `listen`: if `done=0`, must wait for a write event.

State machine




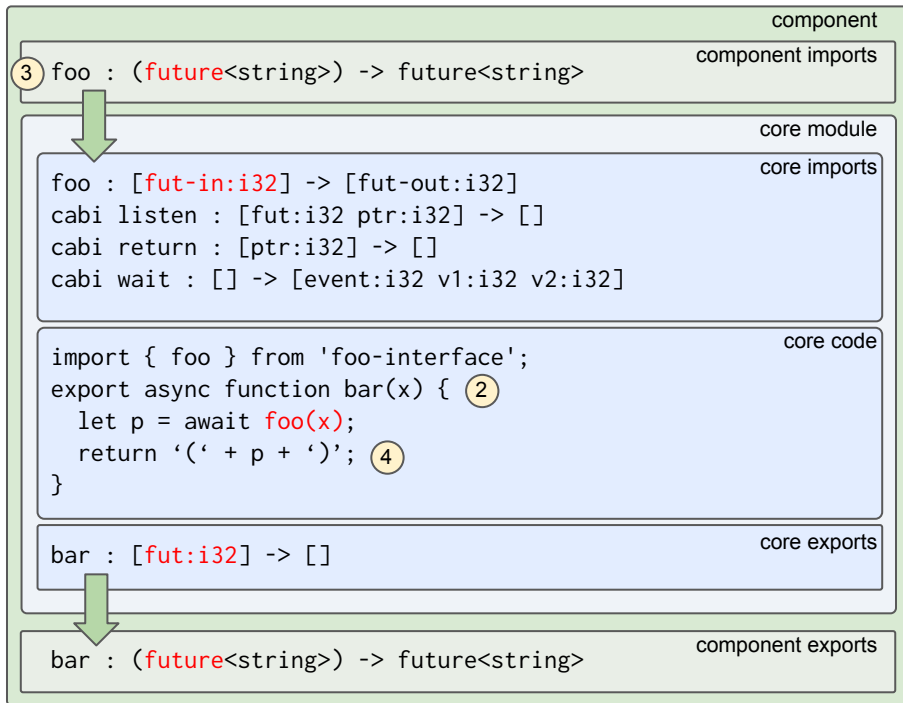
A $\text{future}\langle U \rangle$ is just a $\text{stream}\langle T, U \rangle$ that writes zero T s before returning a U .

We can think of future and stream as two static descriptions of the *dynamic behavior* of a “task”...

Definitions


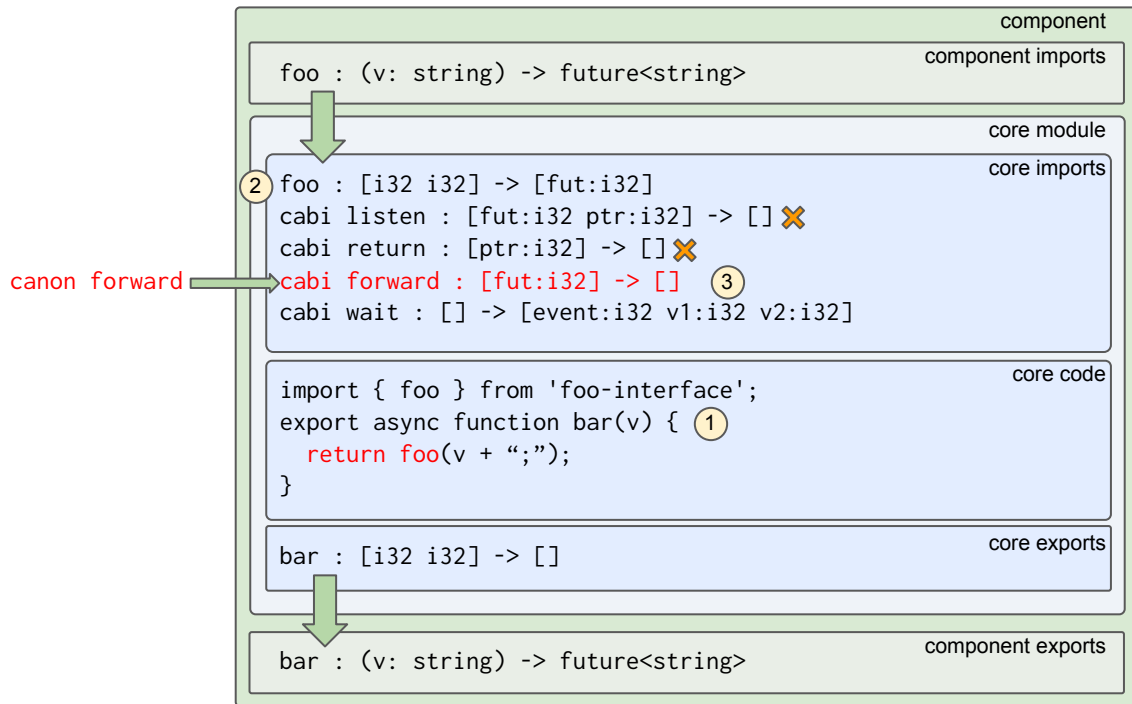
- A **task** is a stateful *resource* representing an asynchronous call
 - ... containing a (live or suspended) stack/continuation
 - ... producing one (future) or multiple (stream) values
 - ... before performing some final execution to finish up and exit.
- Tasks may be **host-implemented** or **wasm-implemented**
 - ... representing both the host async caller and host-implemented async import calls.
- A task also has a **producer state** and a **consumer state**
 - Producer state = { waiting, writing, returning, finishing, finished }
 - Consumer state = { listening, not-listening (= backpressure) }
- Tasks depend on each other via **task handles**
 - The Canonical ABI stores task handles in instance-wide tables, referred to by i32 indices.
- The graph with tasks as nodes and task handles as edges forms a **task tree**.
 - A task handle represents **unique ownership** of a **subtask** by a **supertask**.
 - A supertask **owns** its subtasks, but may **transfer ownership** of them...

Passing tasks as parameters


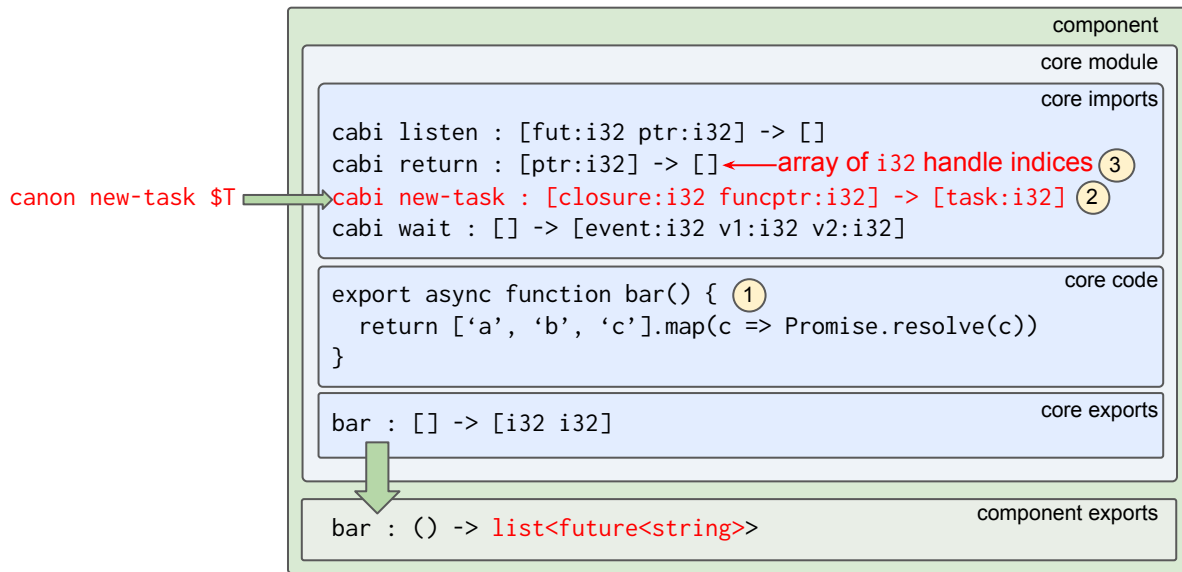


Transfer/move semantics preserves tree-ness

Returning tasks (without copying)



Returning new tasks not derived from imports



In general: `future<T>/stream<T,U>` can be arbitrarily nested in params/results

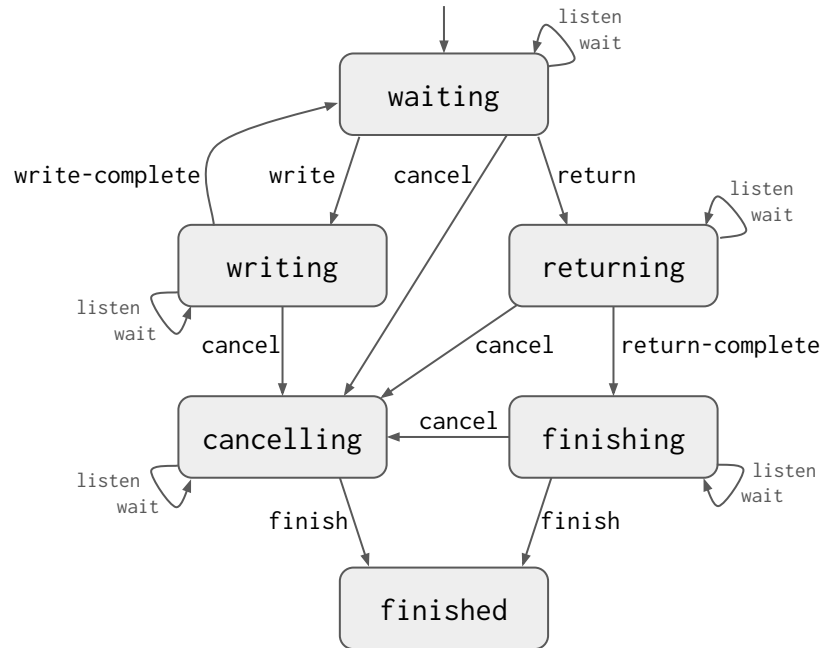
Structured concurrency

- What does “structured concurrency” mean at the component level?
 - **Subtasks can’t outlive their supertasks.**
 - ... although the supertask can change over time (via **explicit** parameter/result passing).
- How is this achieved?
 - `canon lift` traps if a task attempts to finish while it has any subtasks.
 - It’s up to each language toolchain to figure out what this means from a language perspective
 - E.g., maintaining + draining a per-task event loop.
- Why is this useful?
 - Abstractly, it ensures async callees are an encapsulated implementation detail of async callers
 - ... just like with sync calls; all we’re doing is allowing the calls to overlap (= concurrency).
 - Concretely, this enables:
 - Cross-component async callstacks in devtools
 - Automatic “tracing” (in the “observability” sense)
 - A *compositional* recursive cancellation story...

Task cancellation

- What does “cancellation” mean?
 - **The supertask has lost interest**; the result value(s) will not be consumed.
 - Also, the supertask would like the subtask to wrap it up and finish.
- Can we just straight-up *delete* the subtask? No:
 - Its containing instance (and linear memory) lives on, so this might leak / leave in an invalid state
 - Analogous to the usual problem of killing a thread without running destructors
- Can we *force* a subtask to wrap up “promptly”? No:
 - The subtask may legitimately need to perform some async work as part of its cancellation.
 - E.g., rolling back a transaction or posting some logs or metrics
- Thus, cancellation must be *cooperative*:
 - Non-cooperative host/guest scenarios need a new “blast zone” feature *anyhow*
- Canonical ABI additions:
 - `canon cancel : [subtask:i32] -> []`
 - The subtask receives a “cancelled” event from `canon wait`.
 - The supertask later receives a “finish” event from `canon wait` after subtask finishes.

State machine (with cancellation)



Task finishing

- How does a task finish?
 - Normal ABI: return from the export (empty return value)
 - Callback ABI: return the “finish” code from the export or callback
 - Structured concurrency traps if a supertask tries to finish with any remaining subtasks.
- Q: when *precisely* is a subtask taken “off the books”?
 - A: Only when the supertask calls:
 - `canon drop : [subtask:i32] -> []`
 - This lets the toolchain control when the task index may be recycled (like closing an fd).
- Thus: supertasks must explicitly drop all their subtasks before finishing.
- What happens if you drop a subtask that isn’t finished?
 - We can’t just kill it (leak/corruption) or let it keep running (structured concurrency)
 - So: trap
- Due to uniqueness, a dropped task can be eagerly destroyed.

How does cancellation look in the source language?

- JavaScript:
 - If the implementation GCs an unresolved Promise: call `cancel`
 - Eager cancellation via [AbortController](#) signal accepted by JS import bindings.
- Rust:
 - If a Future's destructor is called before the future is resolved: call `cancel`
- Both: the language runtime implicitly waits for all subtasks to finish
 - ... before finishing the Future/Promise returned by the export.
- But what if I want to *explicitly* wait for finish in my source language?
 - Usually I don't care, but I may in advanced scenarios.
 - TBD
 - Maybe the bindings could define a *subclass* of Promise/Future that exposes the finish event?

Task scheduling

- The Component Model defines a scheduler loop executed by the host.
- The task tree serves as the scheduler state.
- Initially, the task tree contains a single root node representing the host.
- On each iteration, one of the following may happen (non-deterministically):
 - The host creates a new task to execute an async export (for whatever host-defined reason).
 - E.g., HTTP call, timer fired, UDF invoked, message arrived, ...
 - An I/O operation completes, transitioning a host-defined leaf task to *writing/returning/finished*.
 - A *writing/returning* subtask copies a value to a listening supertask; both sides receive events.
 - A supertask is notified that one of its subtasks is *finished*.
 - If none of these apply, the loop blocks (waiting for I/O or a new export call to be triggered).
- Thus, we have *two-level* scheduling:
 - *Inter*-component: language-agnostic via the above scheduling loop.
 - *Intra*-component: language-specific as compiled by the language toolchain.
 - Not surprising: this is similar to OS processes, but without the separate threads.

Not covered here, but (maybe) part of the proposal

- Optimization: batched reads and writes
 - Multiple writes can write into a single listen buffer.
 - Multiple listens can read from a single write buffer.
- `notify`
 - How to implement a “tee” and handle the slow-reader/fast-writer case
 - This allows deadlocks :-(
 - BUT, they can be reliably detected by the semantics to produce an error return code.
 - Unavoidable if you can “tee” and “join” (due to resource exhaustion)
- `unlisten/unwrite/unreturn`
 - If the guest code needs to synchronously deallocate a buffer passed to `listen/write/return`
 - ... may conflict with efficient async I/O host implementation?

Summary

- Proposing a common set of high-level **concurrency types**
 - The types prescribe a low-level control-flow *protocol* between the two sides of an async call
 - The runtime mediates and enforces the protocol via dynamic checks and the scheduler.
- Should be “bindable” into different languages’ native concurrency support
 - Ergonomically usable without manual hand-written *per-interface* glue code
 - Turning an $O(N \times M)$ situation into an $O(N+M)$ situation (N = interfaces, M = languages)
 - (Which is the general goal of the Component Model.)
- Interestingly: **not a “process”/“thread” model** (e.g., CSP, π calculus)
 - No (preemptively-scheduled) threads (instead [stack switching](#) / algebraic effects)
 - No channels, pipes, message-boxes (instead direct copy + buffering in the wasm + backpressure)
- ... but could complement a process/thread model
 - Can compile process-style languages to run **inside** a component
 - Can instantiate components **inside** the processes of a process model

Next steps

- Use stream/future in WASI snapshot preview2
 - Just the relevant subset (lower+listen+wait)
 - Using the Canonical ABI to define as a pure Core WebAssembly interface...
 - ... so not dependent on the Component Model
- Write up in a PR to the [component-model](#) repo
 - Extend the explainer (AST), binary format and [CanonicalABI.md](#)
 - Get implementation feedback
- Working with Lucy Amidon and Amal Ahmed on formal semantics
 - Rough idea: define this all in terms of algebraic effects (composable with [stack-switching](#))