

Component Model Async Support

WebAssembly CG

May/June, 2022

Outline

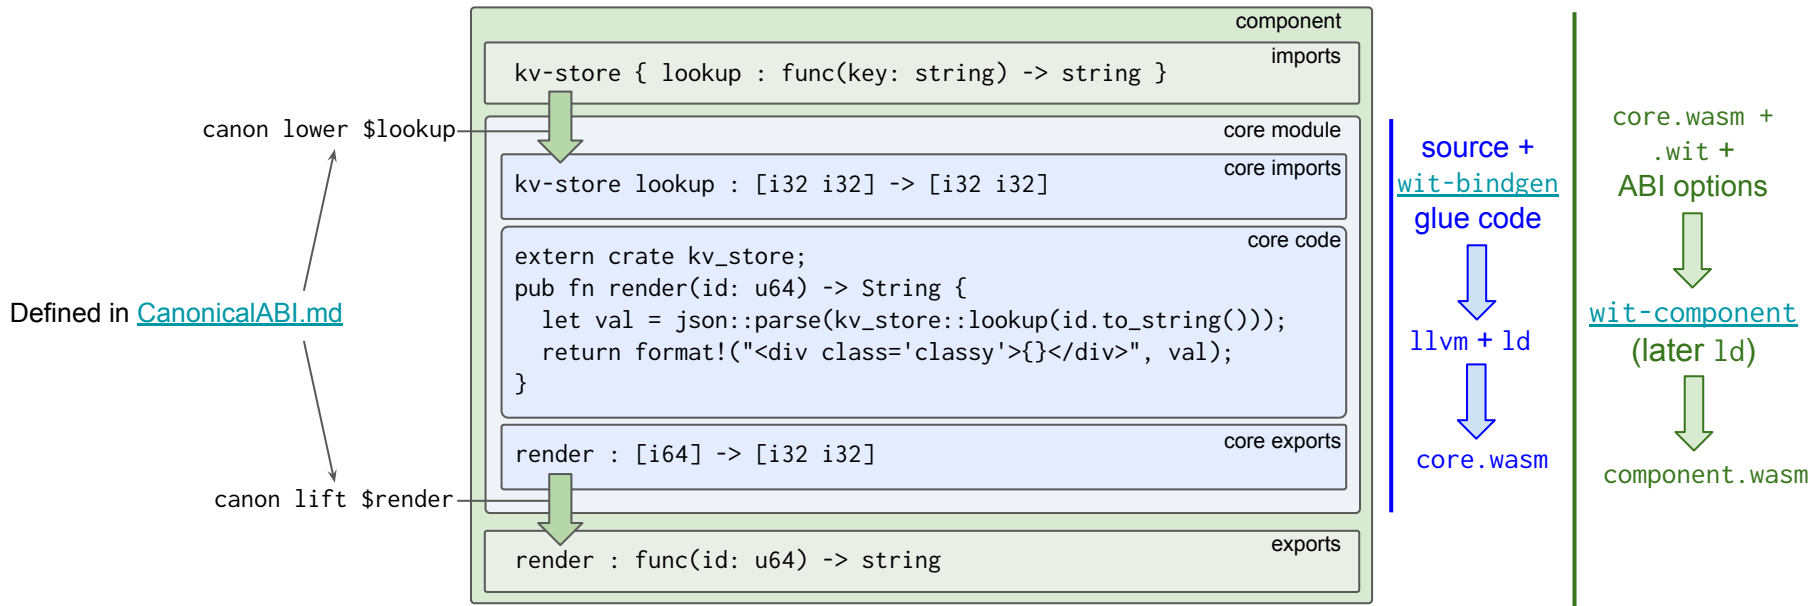
- Motivation
- Background: synchronous Canonical ABI
- Async support
 - future
 - Optimization: callback ABI
 - Optimization: eager return
 - Optimization: stream
 - Optimization: splicing and skipping streams
- Structured concurrency
 - Task
 - Task tree
 - Task cancellation
 - Task scheduling
- Core WebAssembly stack-switching integration

Caveat: still in flux; feedback welcome

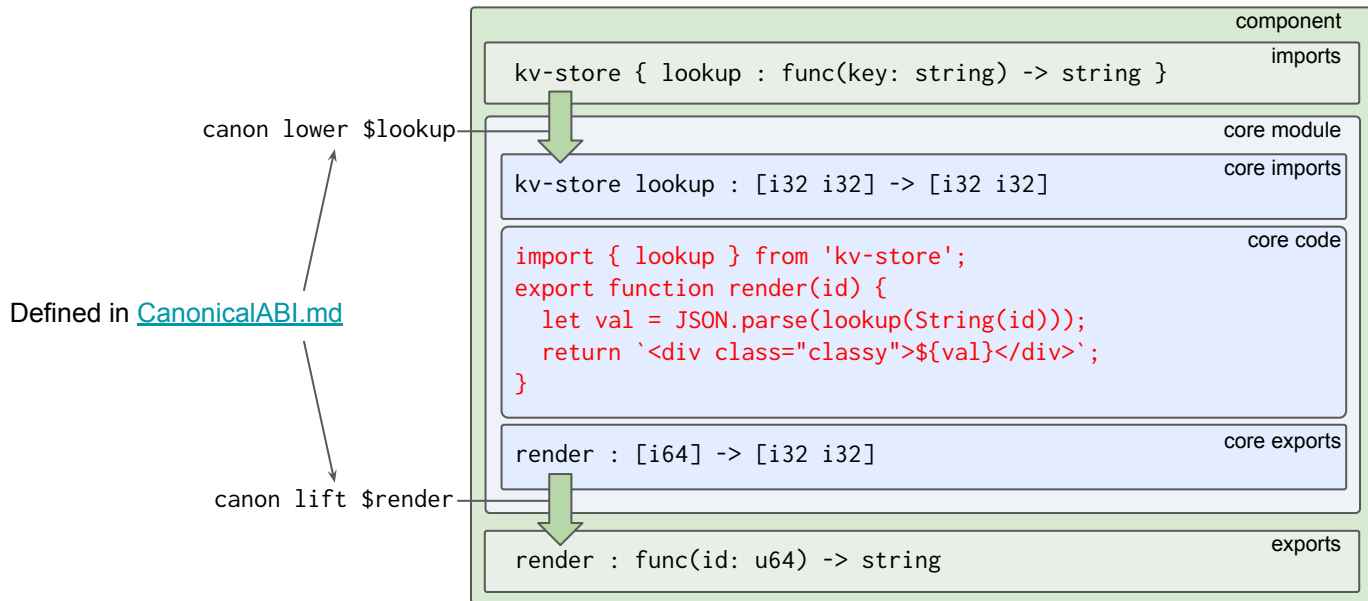
Motivation

- (One slide recap of previously-presented [1,2,3](#))
- How do we specify async/non-blocking operations in WASI and wit?
- Can't we just add first-class functions / callbacks to wit?
 - Cyclic leak problems in non-GC setting (see: Web APIs)
 - Very low-level -- requires manual per-API wrapping to integrate with language concurrency
- Requirements/goals:
 - Virtualizability: async interfaces can be implemented by the host or wasm
 - Efficient I/O implementation when the “other side” is the host (e.g., epoll, io_uring)
 - Ergonomic automatic (wit-bindgen) language bindings
 - Support different styles of language-level concurrency (sync, non-blocking, async, coroutine)
 - Built-in backpressure story (not left as an exercise to the developer)
 - Integrated select / timeout / cancellation across independent interfaces (WASI and host-defined)
 - Ability to keep executing after returning a final value
 - “Just because I want async + modularity doesn't mean I want multi-threading”

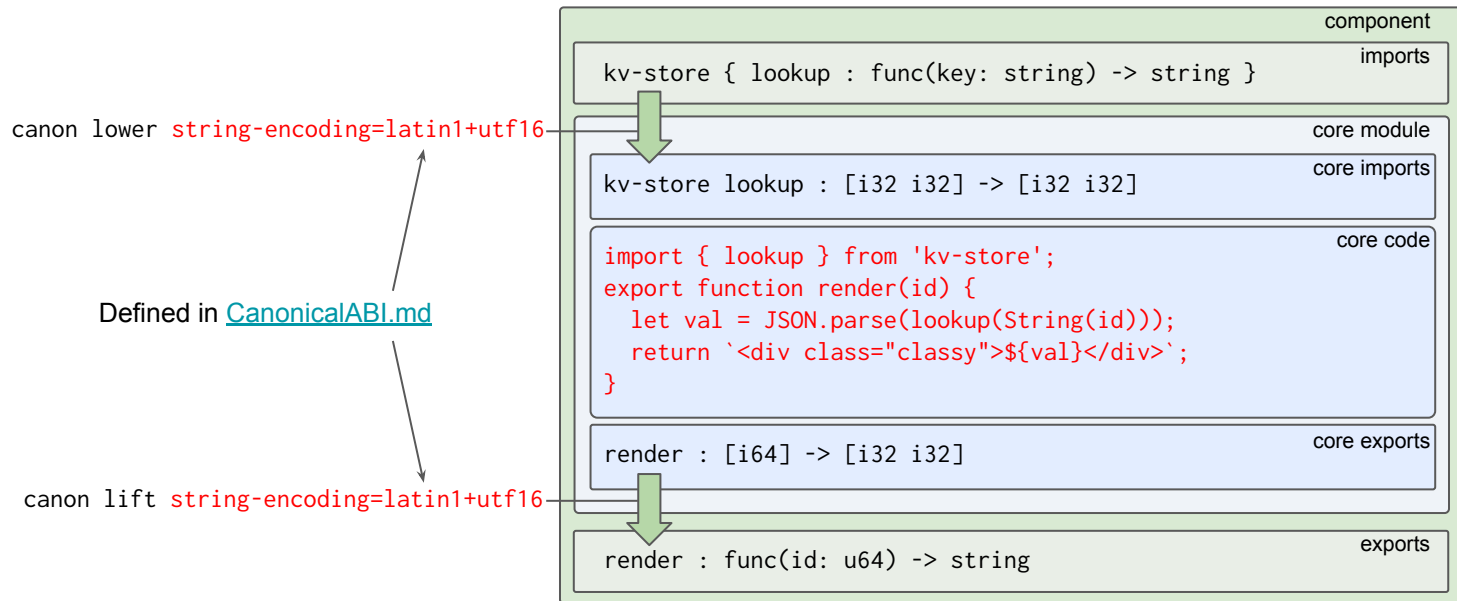
Background: synchronous canonical ABI



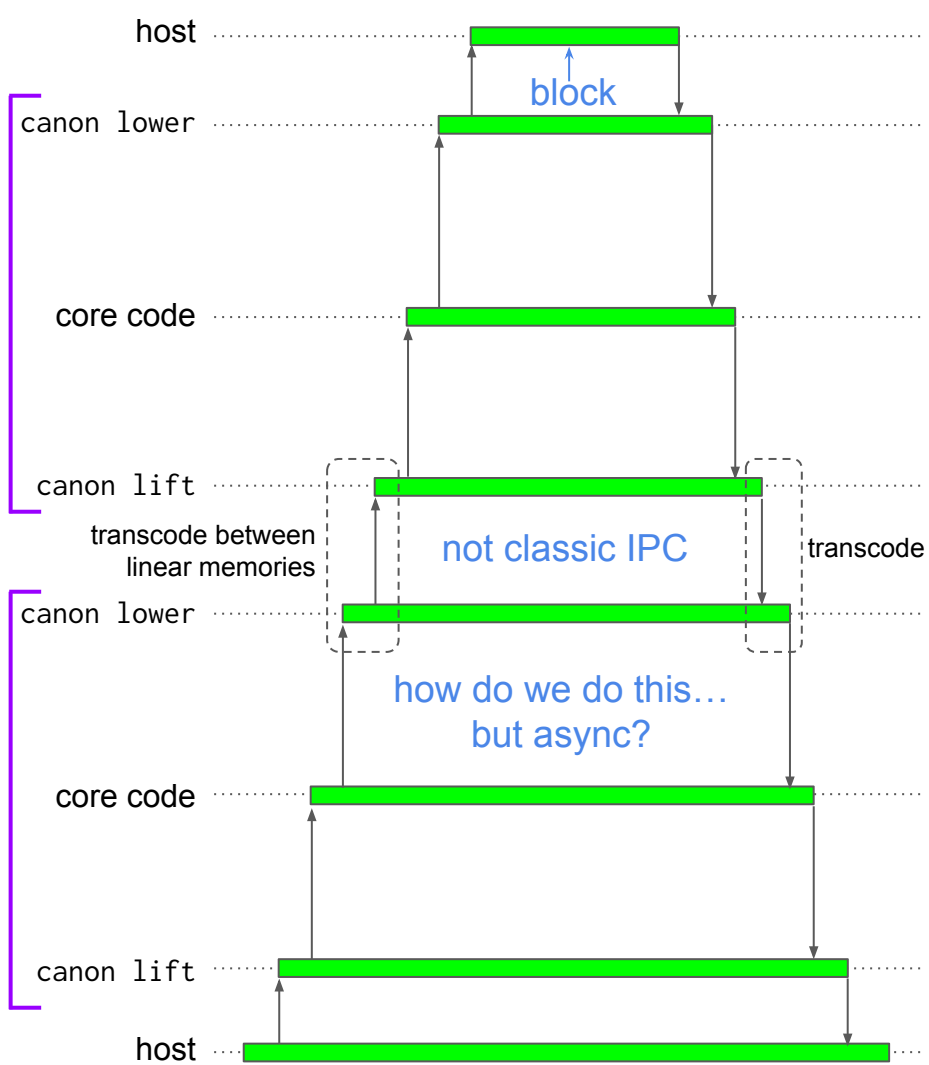
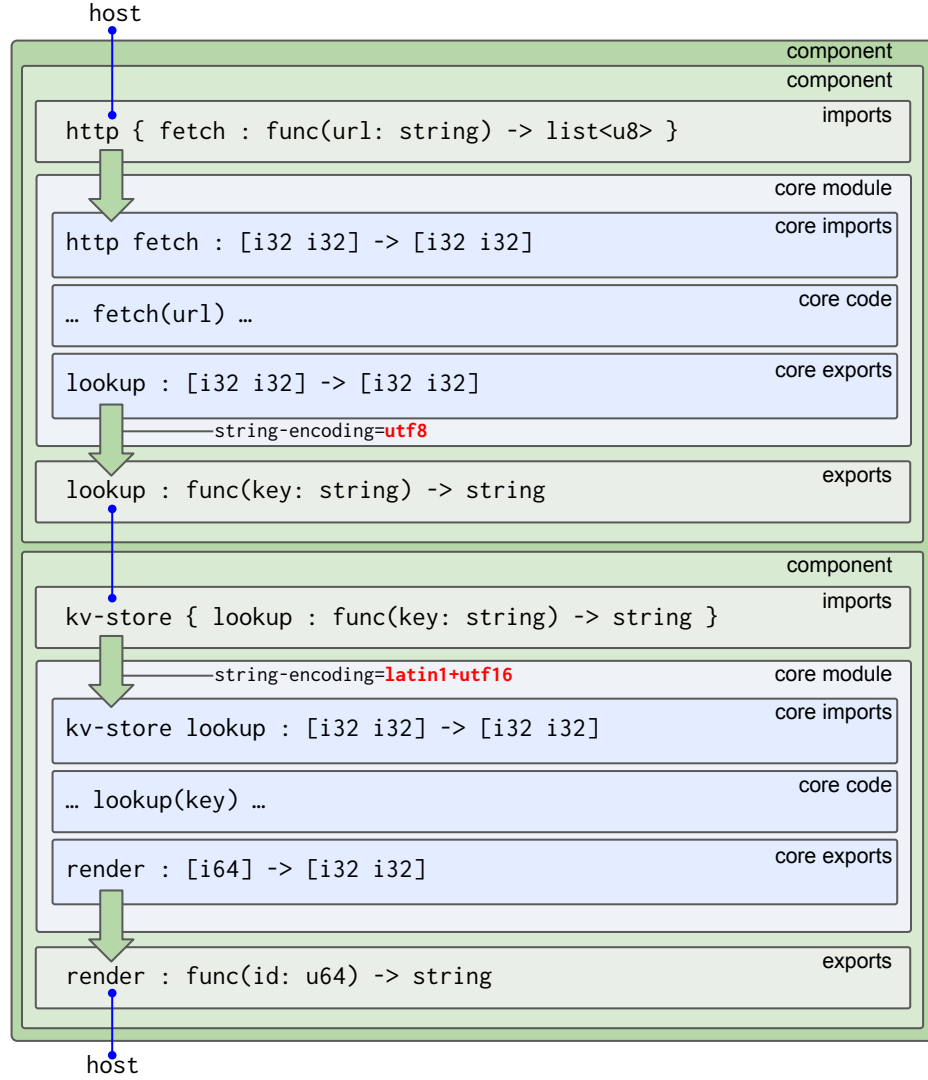
Background: synchronous canonical ABI



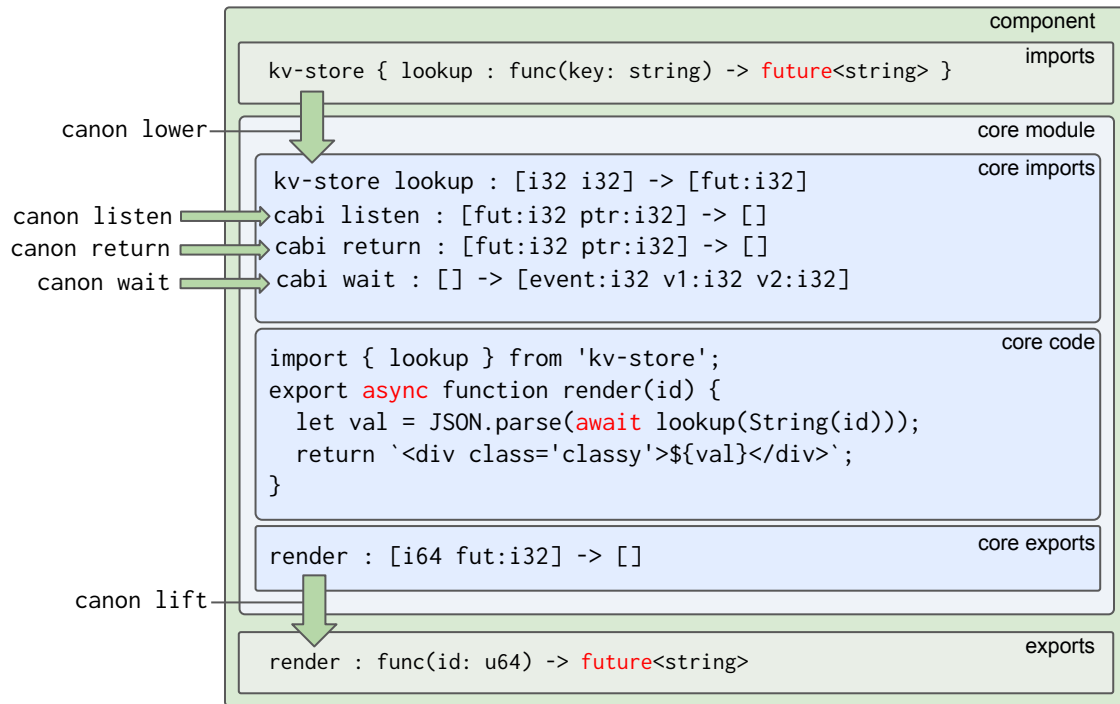
Canonical ABI options



Possible because canon lift and lower bracket all component entry/exit



future



canon lift:

- Passed the index of the future for this export call
- The callee must call return then wait for return-complete

canon return:

- Non-blocking: *offers* a T return value for the given future
- ptr must stay valid until the return-complete event
- Traps if given the result of canon lower

canon wait:

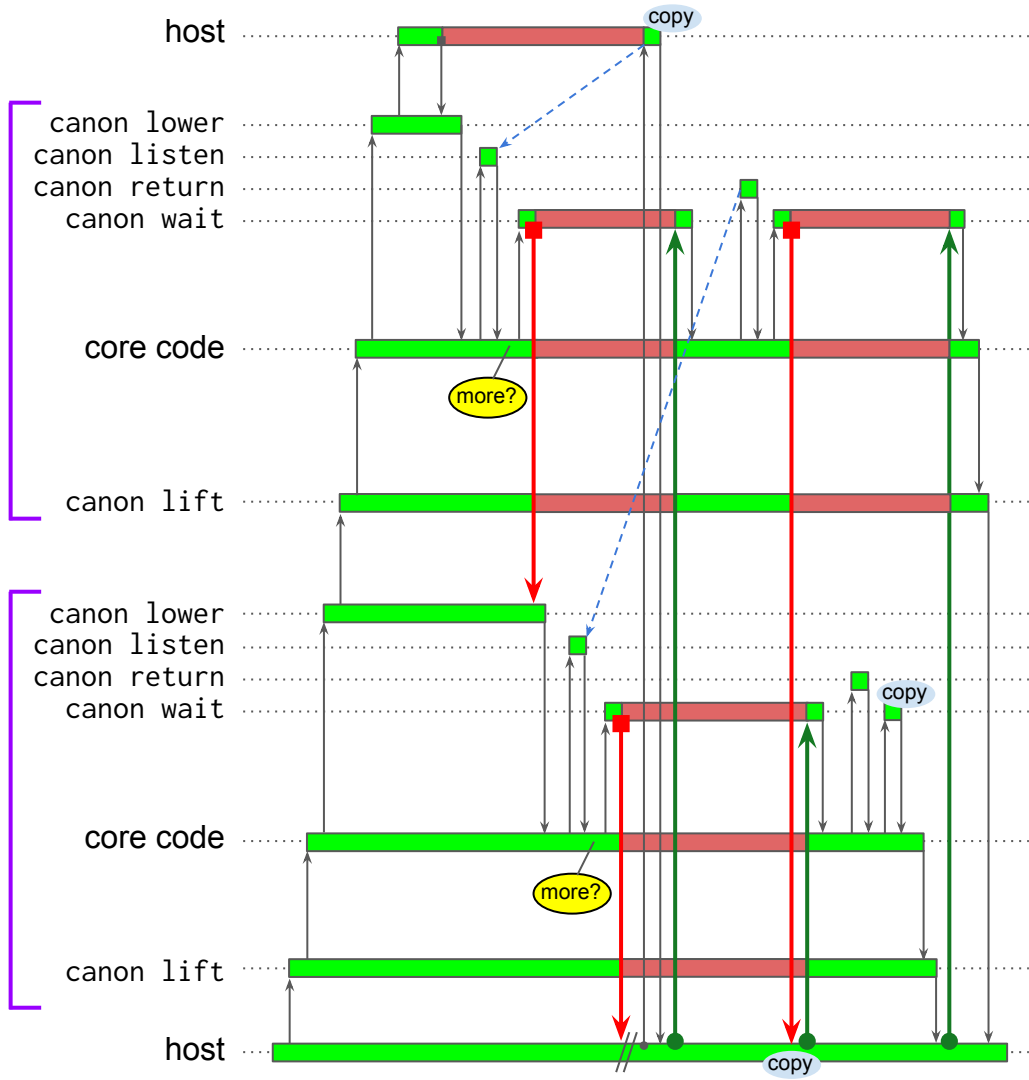
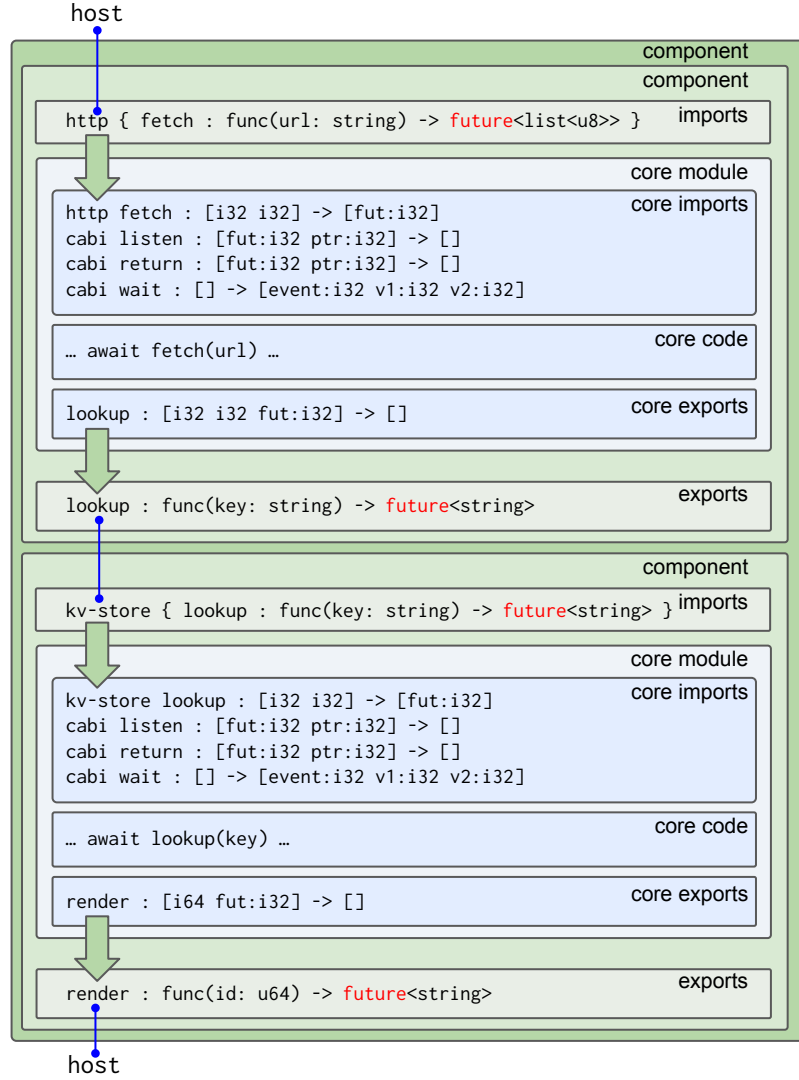
- Blocks until *some* event occurs, including:
 - return-complete (v1 is the future index)
 - returned (v1 is the future index)

canon lower:

- Returns the index of the future for this import call
- The future is initially in a “not listening” state.

canon listen:

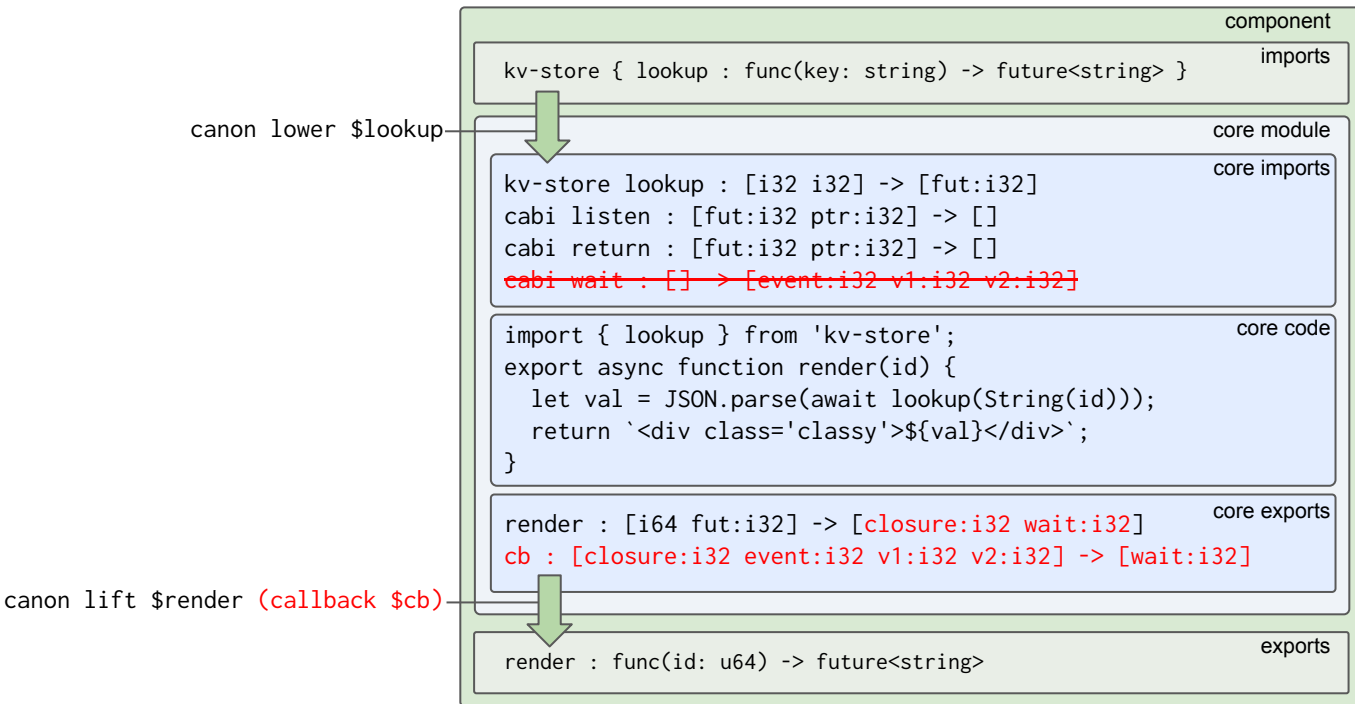
- Non-blocking: *offers* a buffer to receive the future’s value
- ptr must stay valid until the return event.
- Traps if given the parameter from canon lift



Optimization: callback ABI

- For the future/promise/task/async+await family of languages...
 - viz., .NET, JS, Rust
- ... `wait` will always be performed at the base of the callstack
 - As part of a runtime-owned event loop designed to integrate with OS event system
- In this setting, full stack-switching is overkill
 - The language compiler/runtime already did all the “hard work” of clearing the native stack
- It would be nice to allow producer toolchains to opt out of stack switching
 - Reap the performance benefits paid for by their async model
- Also, some hosts won't support native stack switching for a while (or ever?)
 - Stack switching can be emulated/polyfilled via [asyncify](#), but it's expensive

Optimization: callback ABI



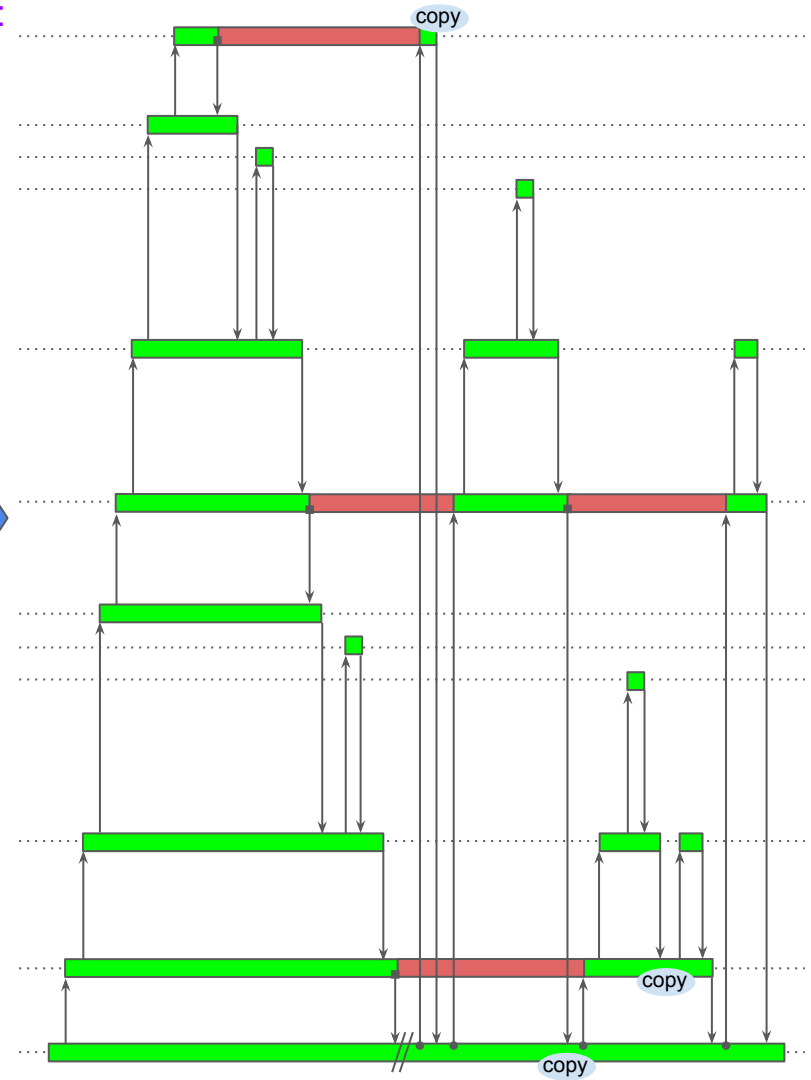
As if:

```
canon-lift (params) {
  let (c, wait) = render(params)
  while (wait) {
    let (event,v1,v2) = wait()
    wait = cb(c, event, v1, v2)
  }
}
```

Notes:

- Encapsulated impl. detail
- Composes with non-callback
- Calling canon wait traps

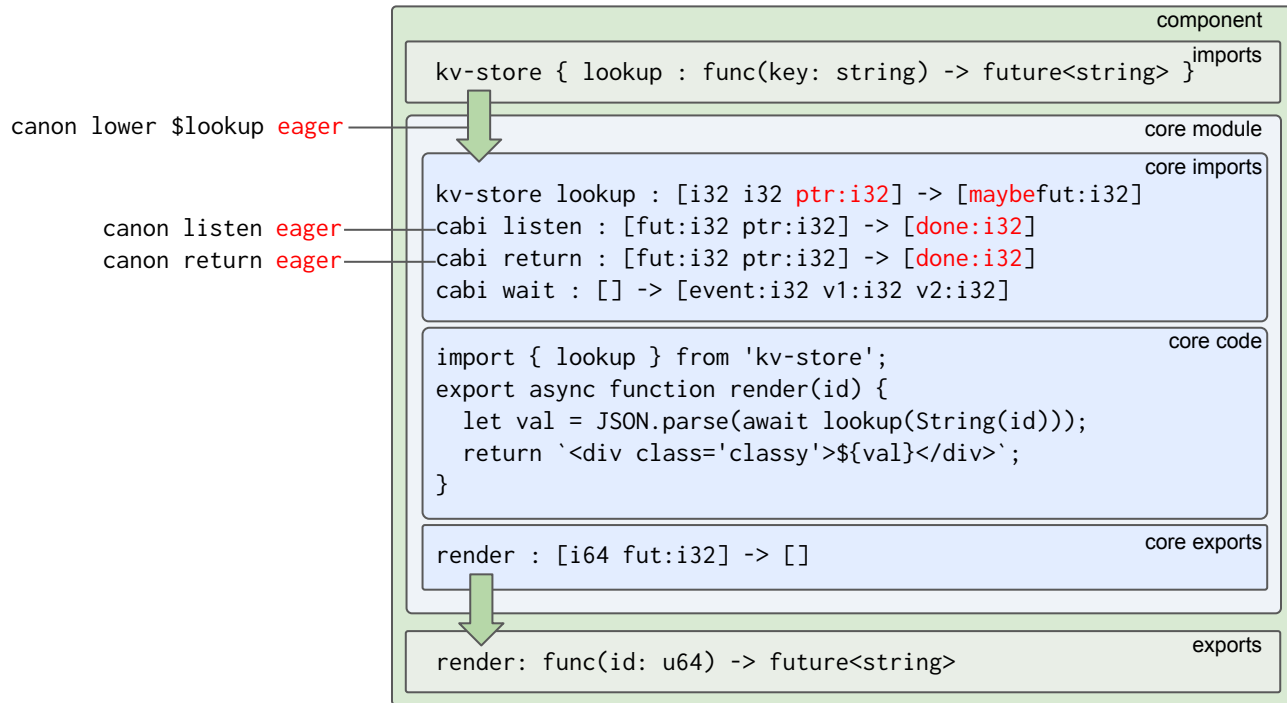
host




Optimization: eager return

- If the result is already available, future adds overhead
 - Runtime internal allocations, extra listen / wait calls.
- Some languages allow promises/futures to be returned already-resolved
 - Avoiding a trip through the event loop


Optimization: eager return



Before:



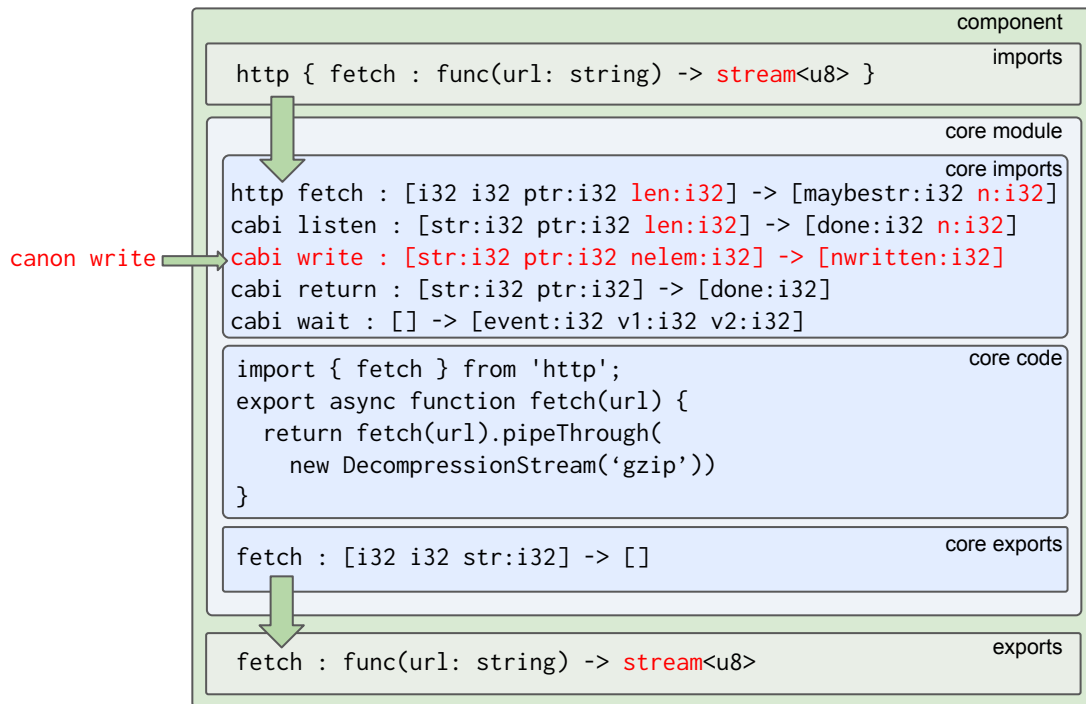
After:



Optimization: stream

- Streams are possible with `stream<T> = future<option<pair<T,stream<T>>>>`
 - (Hand-waving over how we make `stream<T>` recursive...)
- But that's not going to cut it for streams of bytes
 - Need: bulk copies, directly between linear memory (in component-to-component)
 - Don't want to create a completely separate `bytestream` (stream of `vec2` should be fast too!)
- Languages increasingly have a built-in stream primitives
 - Tightly integrated with the rest of the language (syntax, concurrency model, backpressure, ...)
 - Want interface types to automatically bind to these stream language primitives.
- So define `stream<T>` as a new interface type constructor
 - Both as an optimization but also for improved language bindings
- Streams also sometimes have a “closing” value distinct from the elements
 - Effectively: `stream<T,U> = future<either<U,pair<T,stream<T,U>>>>`
 - E.g., `main: (stdin:stream<u8>, argv:list<string>) -> stream<u8,expected<_,_>>`
 - `stream<T> = stream<T,unit>`

Optimization: stream



canon lower (of function returning `stream<T,U>`):

- Additionally takes the byte-length of ptr
- `maybestr=0` means ptr holds `T*U`, `n = |T*|`.
- `maybestr≠0` means ptr holds `T*`, `n = |T*|`.

canon listen (on `stream<T,U>`):

- Requires `len > max(sizeof(T), sizeof(U))`
- ptr receives `T* xor U`; must stay valid until written event.
- `n=0` means "returned U" / `n>0` means "n Ts written"

canon write (on `stream<T,U>`):

- Non-blocking: *offers* `nelem` T values for the given stream
- Requires `nelem > 0` (progress, symmetry with listen).
- ptr must stay valid until write-complete

canon return (on `stream<T,U>`):

- Traps if write in progress.
- Non-blocking: *offers* a U return value for the given stream
- Closes the stream (no more writes possible)

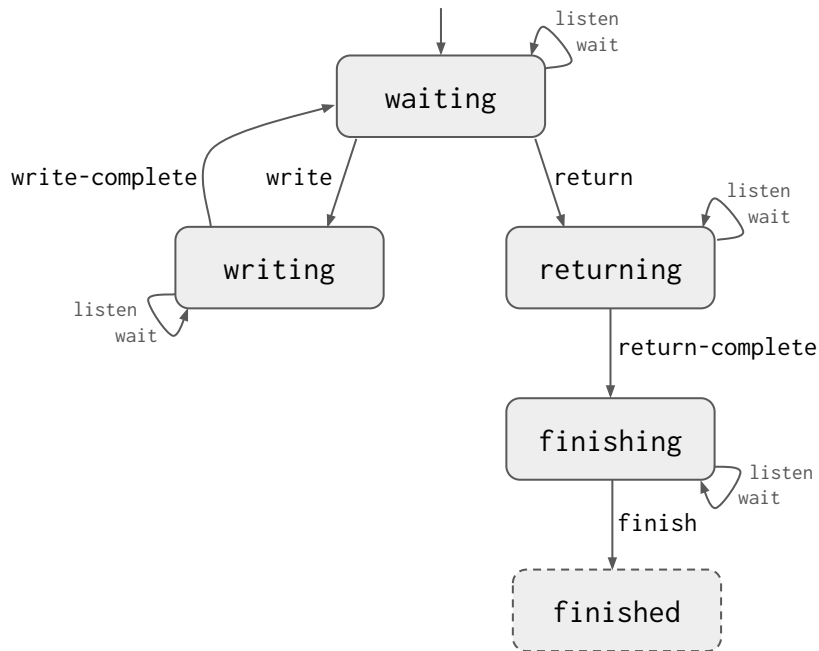
canon wait:

- Additional events:
 - write-complete (`v1` is stream index, `v2` is num written)
 - written (`v1` is stream index, `v2` is num written)
 - The stream goes back to the "not listening" state
 - Not listening = **backpressure**

Optimization: splicing and skipping streams

- It's very common to copy big chunks from one stream to another
 - Don't want to have to read into linear memory just to immediately write back out.
- `canon splice : [src-str:i32 dst-str:i32 nelem:i32] -> [done:i32]`
 - Acts like `listen(src-str, buf) + write(dst-str, buf)`, but without the `buf`.
 - If `done=0`, must wait for a write-complete event (`nwritten <= nelem`).
- `canon forward : [src-str:i32 dst-str:i32] -> []`
 - Like `splice(src-str, dst-str, ∞)`, but caller doesn't have/get to wait on the completion.
 - `dst-str` is immediately removed from table.
 - Also works on futures (analogous to JS rules when a `then()`-function returns a Promise).
- Sometimes we want to ignore a run of elements in a stream
 - Don't want to copy bytes into linear memory just to advance the read offset
- `canon skip : [str:i32 nelem:i32] -> [done:i32]`
 - Acts like `listen`: if `done=0`, must wait for a written event.

State machine



A $\text{future}\langle U \rangle$ is just a $\text{stream}\langle T, U \rangle$ that always writes zero T s before returning a U .

We can think of future and stream as two static descriptions of the *dynamic behavior* of a “task”...

Tasks

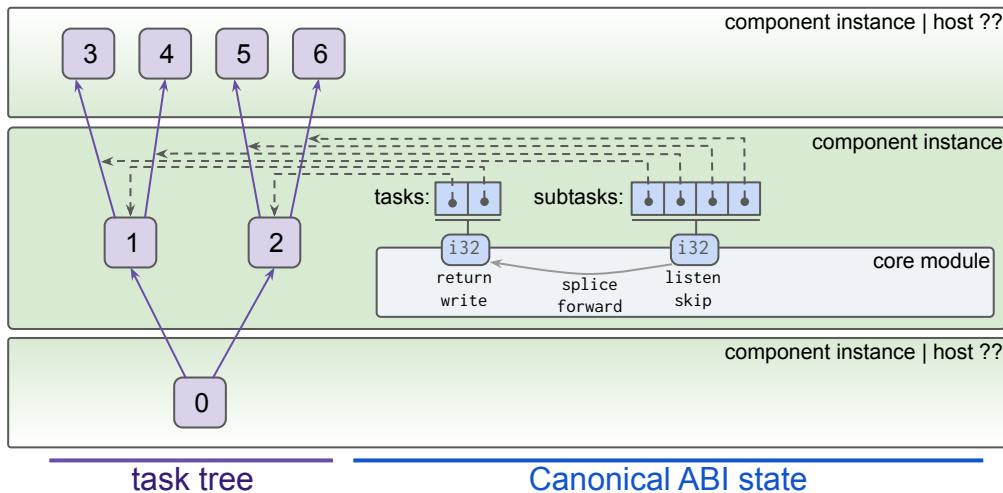
- A **task** is a stateful *resource* representing an asynchronous computation
 - ... producing one value (future) or a sequence of values (stream).
- A task has a **producer state** and a **consumer state**
 - Producer state = { waiting, writing(ptr,n,done?), returning(ptr,done?), finishing, finished }
 - Consumer state = { listening(ptr,len,done?), not-listening }
- canon write/return **consult consumer state before updating producer state**:
 - If consumer state is listening: do the copy, then eager return.
- canon listen **consults producer state before updating the consumer state**:
 - If producer state is writing/returning: do the copy, then eager return.
- The done? boolean sub-state says we still need to *notify* the task
 - ... before transitioning back to waiting/returning/not-listening.
- Batch reads/writes by bumping writing.ptr / listening.ptr
 - ... without setting done?, so that further reads/writes are possible with the same buffer.

Task trees

- Tasks form a **task tree** with edges from **supertasks** to **subtasks**.
 - (The natural “parent”/“child” terminology is already used to describe instance nesting.)
- Because of tree-ness, subtasks are **uniquely owned** by their supertask.
 - However, subtask ownership can be **transferred** (as we’ll see next).
- Tasks can be **host-implemented** or **wasm-implemented**.
 - The task tree’s **root** is an ever-present host-implemented task (calling component exports).
 - All other host-implemented tasks are **leaves** (called by component imports).

Canonical ABI representation of tasks vs. subtasks

- The Canonical ABI defines 2 component-instance-local tables:
 - **Task table:** tasks implemented by the containing instance.
 - **Subtask table:** subtask edges whose sources are tasks in the task table.
- The tables are instance-wide (like linear memory)
 - So any core code can listen/write/return to any (sub)task any time.



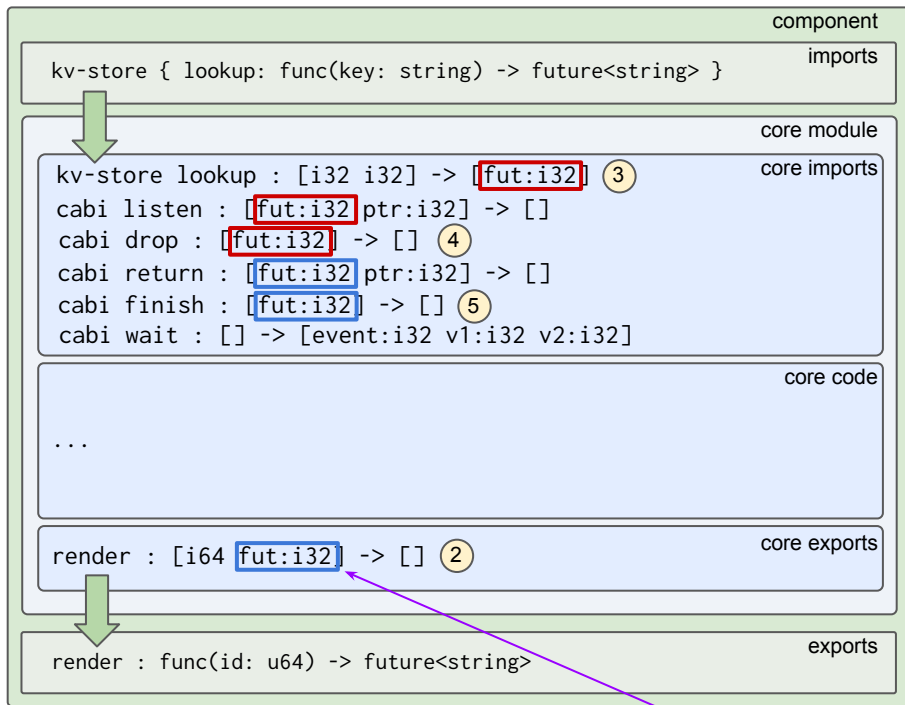
Structured concurrency

- A task tree is effectively the “async version” of a synchronous callstack
 - Asynchronicity means that new stack frames can be created before their sibling frames finish.
 - Hence “tree” not “stack”
 - The same tree structure also shows up when you have lexical closures (callbacks)...
- **Structured concurrency** means making the “async callstack” metaphor hold:
 - Invariant: **subtasks can’t outlive their supertasks**
 - ... although the supertask can change over time via explicit ownership transfer (how next)
- Why is this useful?
 - Abstractly, it ensures async callees are an encapsulated implementation detail of async callers
 - ... just like with sync calls; all we’re doing is allowing the calls to overlap (= concurrency).
 - Concretely, this enables:
 - Devtools / debugging
 - Tracing (in the “observability” sense)
 - A compositional recursive cancellation story (next next)
- How is this invariant achieved?

Structured concurrency


- How exactly does a task “finish”?
 - `canon finish : [task:i32] -> []` (task can be a future or stream)
 - **Traps if task has any remaining subtasks**
 - The supertask of a finished task receives a finished event from `canon wait`.
- When *precisely* is a subtask taken “off the books”?
 - `canon drop : [subtask:i32] -> []` (subtask can be a future or stream)
 - **Trap if subtask isn't finished**
 - Explicit drop lets the toolchain control when the task index may be recycled (like an OS handle).
- Putting these together: before calling `canon finish`, a task must...
 - `canon drop` each subtask, which requires...
 - `canon wait`-ing for finished from each subtask, which requires...
 - Those subtasks to themselves call `canon finish`, which requires ... (recurse)

Original future example *redux*




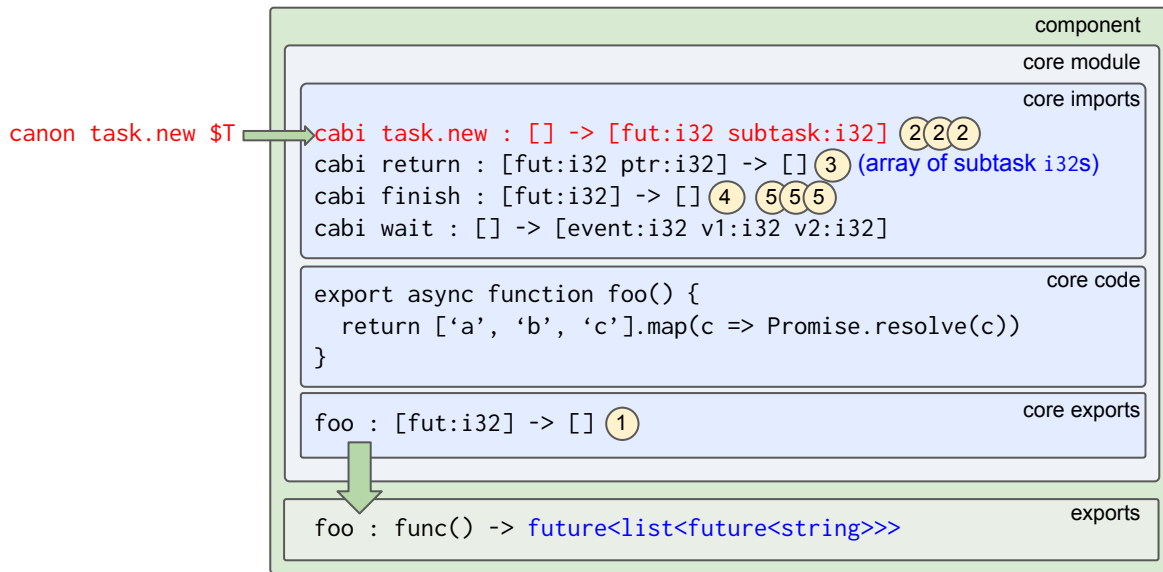
Task table indices

Subtask table indices



Passing the export's returned future as an outparam is a special-case to allow eager return.


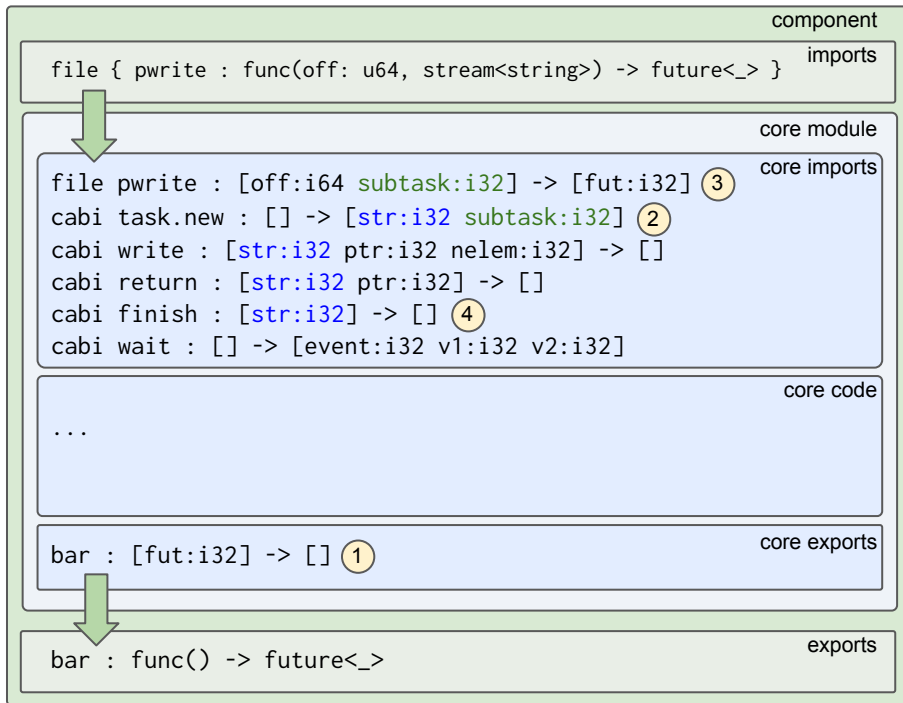
General case: dynamically creating tasks



In general: `future<T>/stream<T,U>` can be *arbitrarily nested* in params *and* results


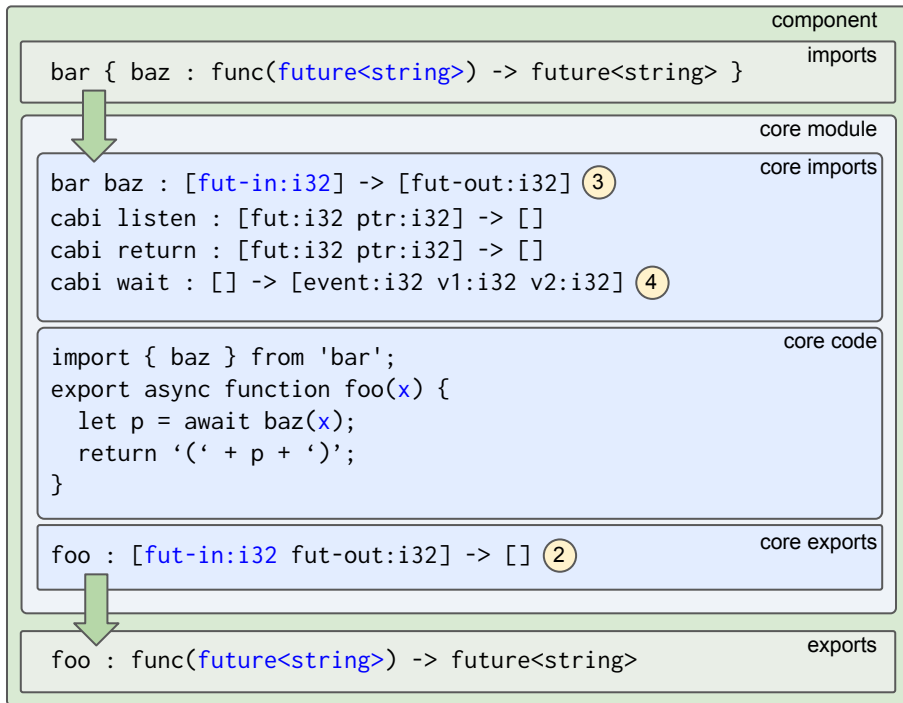
... passed by **ownership transfer** of subtasks from one task to another

Passing tasks as parameters



Note: `bar`→`pwrite`→`<str>` interleaves instances, but task tree ensures acyclicity (no leaks).

Passing tasks as parameters (passthrough)

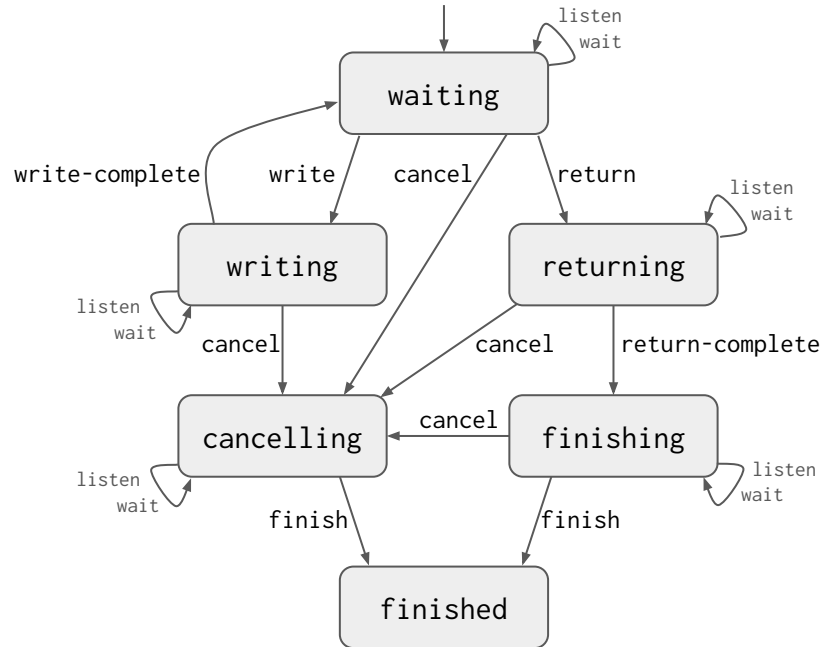


Ownership transfer allows passthrough in all directions

Task cancellation

- What if a supertask starts a subtask but loses interest?
 - E.g., race two network requests, one wins, want to “cancel” the other.
 - `cancel : [subtask:i32] -> []`
- Can `cancel` just straight-up *delete* the subtask? No:
 - Its containing instance (and linear memory) lives on, so this might leak / leave in an invalid state
 - Analogous to the usual problem of killing a thread without running destructors
- Can `cancel` *force* a subtask to wrap up “promptly”? No:
 - The subtask may legitimately need to perform some async work as part of its cancellation.
 - E.g., rolling back a transaction or posting some logs or metrics
- Thus, `cancel` must be *cooperative*:
 - Non-cooperative host/guest scenarios need a new “blast zone” feature *anyhow*
 - So: `cancel` just delivers a cancelled event to the subtask.
 - But the subtask can keep calling imports and waiting before calling `cancel finish`.

State machine (with cancellation)



How does this look in the source language?

- JavaScript:
 - If the implementation GCs an unresolved Promise: call `cancel`
 - Eager cancellation via [AbortController](#) signal accepted by JS import bindings.
- Rust:
 - If a Future's destructor is called before the future is resolved: call `cancel`
- Both: the language runtime implicitly waits for all subtasks to finish
 - ... before finishing the Future/Promise returned by the export.
- But what if I want to *explicitly* wait for finish in my source language?
 - Usually I don't care, but I may in advanced scenarios.
 - TBD
 - Maybe the bindings could define a *subclass* of Promise/Future that exposes the finish event?

Core WebAssembly stack-switching integration

- [TODO]
- What happens when an async export is called while one is in-progress?
 - New stack *and* new task
- Rules ensure every stack implements at least one task
 - Created one first task (on export call)
 - Trap if wait after finish last task
- Same “stacks” as the Core WebAssembly stack-switching proposal
 - Core stack-switching adds: cooperative/green pthreads + coroutines
- Updated definition of “a task”:
 - Resource representing async computation producing one or multiple values
 - Mutable state:
 - Producer and consumer state
 - List of subtasks
 - Stack reference

[TODO: multi-async example]

[TODO: stack.new example]

Task scheduling

- The Component Model defines a scheduler loop executed by the host.
- The task tree serves as the scheduler state.
- Initially, the task tree contains a single root node representing the host.
- On each iteration, one of the following may happen (non-deterministically):
 - The host creates a new task to execute an async export (for whatever host-defined reason).
 - E.g., HTTP call, timer fired, UDF invoked, message arrived, ...
 - An I/O operation completes, transitioning a host-defined leaf task producer state.
 - One or more values are copied from a subtask to a supertask, updating both tasks' state.
 - Notify a task of an event based on its task (producer|consumer) state.
 - “Notify” means resuming the task's stack's `wait` (at which the stack must be suspended).
 - If none of these apply, the loop blocks (waiting for I/O or a new export call to be triggered).
- Thus, we have *two-level* scheduling:
 - *Inter*-component: language-agnostic via the above scheduling loop.
 - *Intra*-component: language-specific as compiled by the language toolchain.
 - Not surprising: this is similar to OS processes, but without the separate threads.

Canonical ABI Summary

[TODO: complete list of canon definitions added/modified]

TODO

- park/unpark
 - How does one stack wait on another stack?
 - This allows deadlocks :-(
 - BUT, they can be reliably detected by the semantics so park returns a failure, not hangs.
 - Unavoidable if you can “tee” and “join” (due to resource exhaustion)
- unlisten/unwrite/unreturn
 - If the guest code needs to synchronously deallocate a buffer passed to listen/write/return
 - ... may conflict with efficient (io_uring) I/O host implementation?
- Async-to-sync adapters
 - A sync import can be implemented by an async function if the caller isn't reentered.
 - Component non-reentrance invariants already enforce this.

Summary

- Proposing a common set of high-level **concurrency types**
 - The types prescribe a low-level control-flow *protocol* between the two sides of an async call
 - The runtime mediates and enforces the protocol via dynamic checks and the scheduler.
- Should be “bindable” into different languages’ native concurrency support
 - Ergonomically usable without manual hand-written *per-interface* glue code
 - Turning an $O(N \times M)$ situation into an $O(N+M)$ situation (N = interfaces, M = languages)
 - (Which is the general goal of the Component Model.)
- Interestingly: **not a “process”/“thread” model** (e.g., CSP, π calculus)
 - No (preemptively-scheduled) threads (instead cooperative scheduling/[stack switching](#))
 - No channels, pipes, message-boxes (instead direct copy + buffering in the wasm + backpressure)
- ... but could complement a process/thread model
 - Can compile process-style languages to run **inside** a component
 - Can instantiate components **inside** the processes of a process model

Next steps

- Use stream/future in WASI snapshot preview2
 - Just the relevant subset (lower+listen+wait)
 - Using the Canonical ABI to define as a pure Core WebAssembly interface...
 - ... so not dependent on the Component Model
- Write up in a PR to the [component-model](#) repo
 - Extend the explainer (AST), binary format and [CanonicalABI.md](#)
 - Get implementation feedback
- Working with Lucy Amidon and Amal Ahmed on formal semantics
 - Rough idea: define this all in terms of algebraic effects (composable with [stack-switching](#))