

# wasi-nn: phase 3 plan

Andrew Brown

August 2024







# History

- wasi-nn initially proposed in 2020
- initial implementation in Wasmtime; subsequent implementations in WasmEdge and WAMR
- interest from VMWare, Microsoft, Fastly, Cosmonic, Fermyon, Midokura, Amazon, Sonos, etc.
- collected feedback in the ML Working Group
- no major changes until:
  - preview1-ABI to preview2-ABI change in late 2023
  - discussions about LLM inferencing

# Implementations



- Wasmtime
- WasmEdge
- WAMR
  - no component model support yet, could target `wasm32-wasip2-module`

# Phase 2 Completed

- Entry requirements:
  -  The portability criteria are documented in the proposal.
  -  Precise and complete overview document is available in a proposal repo around which a reasonably high level of consensus exists.
  -  A WIT description of the API exists.
  -  All dependencies of the WIT description must have reached phase 2.
- During this phase:
  -  One or more implementations proceed on prototyping the API.
  -  A plan is developed for how the portability criteria will be met.

The Bytecode Alliance [ML Working Group](#) has been meeting since October 2023

# Intent

- wasi-nn intends to ask for phase 3 acceptance
  - when: October, November?
- phase 3 entry requirements:
  -  all dependencies of the WIT descriptions must have reached phase 3
  -  portability criteria must be met (or present a plan)
- previous portability concerns:
  - *tricky test suite*: testing internal model logic feels out of scope—too broad
  - *opaque model blobs*: wasi-nn currently has seven graph encodings (`openvino`, `onnx`, `tensorflow`, `ggml`, etc.); likely more to come?

# Plan

1. Implement a test suite based on an “echo” model
2. Replace `load` with `load-by-name`
3. Resolve last few changes (**prompt** interface, **error** shape, etc.)
4. Update implementations (“mostly complete”) in 2 engines




Several WIT snapshots  
implemented in Wasmtime

- Vote on phase 3 in Q4 of 2024

# Plan: Test Suite

- wasi-nn testing is fraught:
  - small floating-point differences between ML frameworks, parameters
  - randomness is inherent in some ML models
  - many frameworks, many models, many values = difficult to check
- Conclusion: **test the wasi-nn API, not the ML model logic** (frameworks already test this!)
- How? An “echo” model...



```
let before = Tensor::new(...)?;  
let graph = nn::load("echo");  
let context = inference::init(graph)?;  
match context.compute([before])? {  
  [after] => assert_eq!(before, after),  
}
```

# Plan: Load-by-name

- A previous concern: “load by bytes”
  - graph encoding list never done
  - encoded bytes are opaque
  - retrieving bytes is too slow
  - models are too big!
- Solution: **replace load with load-by-name**
- ~~Ĵ KǻŸĀ šĝ° žžš/ŭ8Ā/8-\~~

Before:

```
load: func(builder: list<list<u8>>,
         encoding: graph-encoding,
         target: execution-target) ->
    result<graph, error>;
load-by-name: func(name: string) ->
    result<graph, error>;
```

After:

```
load: func(name: string) ->
    result<graph, error>;
```



# Plan: resolve last changes

- See wasi-nn [PRs](#):
  - error shape: **resource** vs **record** vs **string** [#67](#)
  - tensor residence ([#69](#), [#70](#))
  - coalesce tensor parameters ([#77](#))
  - add graph parameters ([#80](#))
  - add prompt interface ([#79](#))

Coalesced compute signature:

```
compute: func(inputs: list<named-tensor>) ->  
    result<list<named-tensor>, error>;
```

New prompt interface:

```
interface prompt {  
    use graph.{graph};  
  
    init: func(graph: graph) ->  
        result<context, string>;  
  
    resource context {  
        compute: func(prompt: string) ->  
            result<string, string>;  
    }  
}
```

# Summary

- Let's move wasi-nn to phase 3!
- A year's worth of design feedback—a little more needed
- Help appreciated on test suite, implementations

# Questions